

The background features a dark teal gradient with a subtle grid pattern. Overlaid on this are several wireframe cubes of varying sizes, some with a slight blue tint. Three thin white arrows point from the text area towards the right side of the slide.

DOCKER FOR ENTERPRISE DEVELOPERS



HOW WE TEACH

- Docker believes in learning by doing, with support.
- The course is lab driven with lecture.
- Work together
- Ask questions at any time



SESSION LOGISTICS

- 2 days duration
- mostly exercises
 - Take time to read the code examples closely
- regular breaks



ASSUMED KNOWLEDGE AND REQUIREMENTS

- Familiarity with using the Linux command line
- A UCP License (download one at <https://store.docker.com/bundles/docker-datacenter/purchase?plan=free-trial>)
- You should know the basics of Docker
 - Run a Docker container
 - Search for and pull images from Docker Store
 - Use Docker for Mac / Windows on your local machine



YOUR LAB ENVIRONMENT

- You have been given several instances for use in exercises.
- Ask instructor for access credentials if you don't have them already.



COURSE LEARNING OBJECTIVES

By the end of this course, learners will be able to:

- Describe the essential patterns used in a highly distributed EE application
- Understand how to configure EE applications for different environments without code change
- Produce and containerize scalable, accessible, and fault-tolerant EE applications
- Apply different debugging and testing techniques to containerized EE applications





DISTRIBUTED APPLICATION ARCHITECTURE



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Describe how familiar design patterns relate to containerized applications
- Enumerate key challenges in containerized application design



DISTRIBUTED APPLICATION ARCHITECTURE

Definition: An application consisting of one or more processes running on one or more nodes.

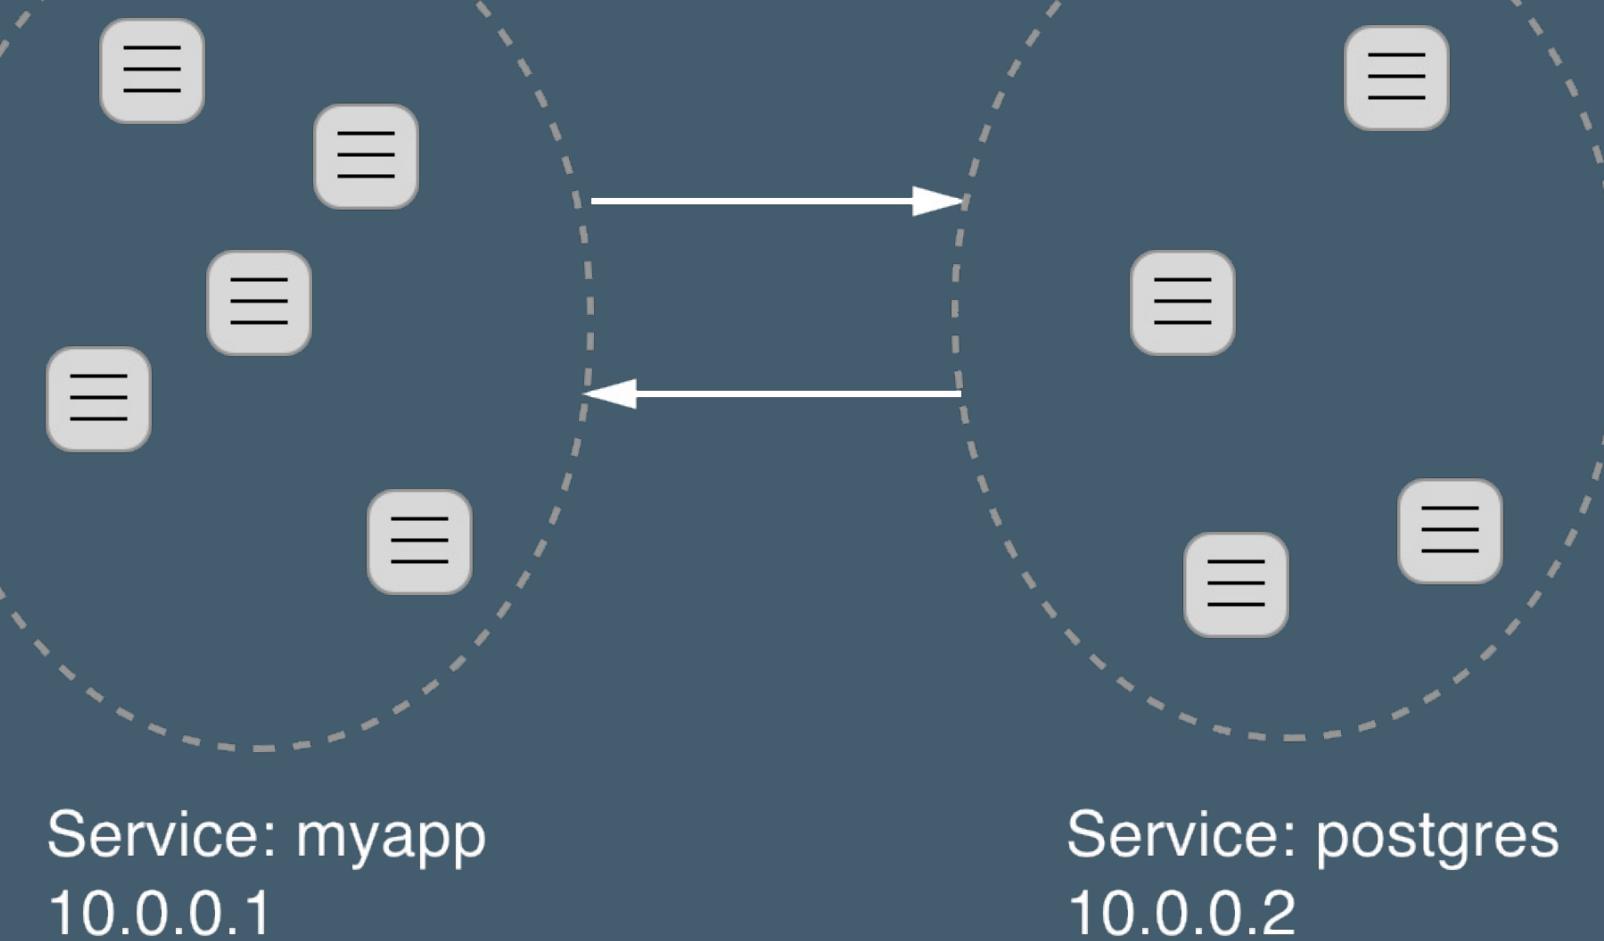


SERVICES OVER PROCESSES

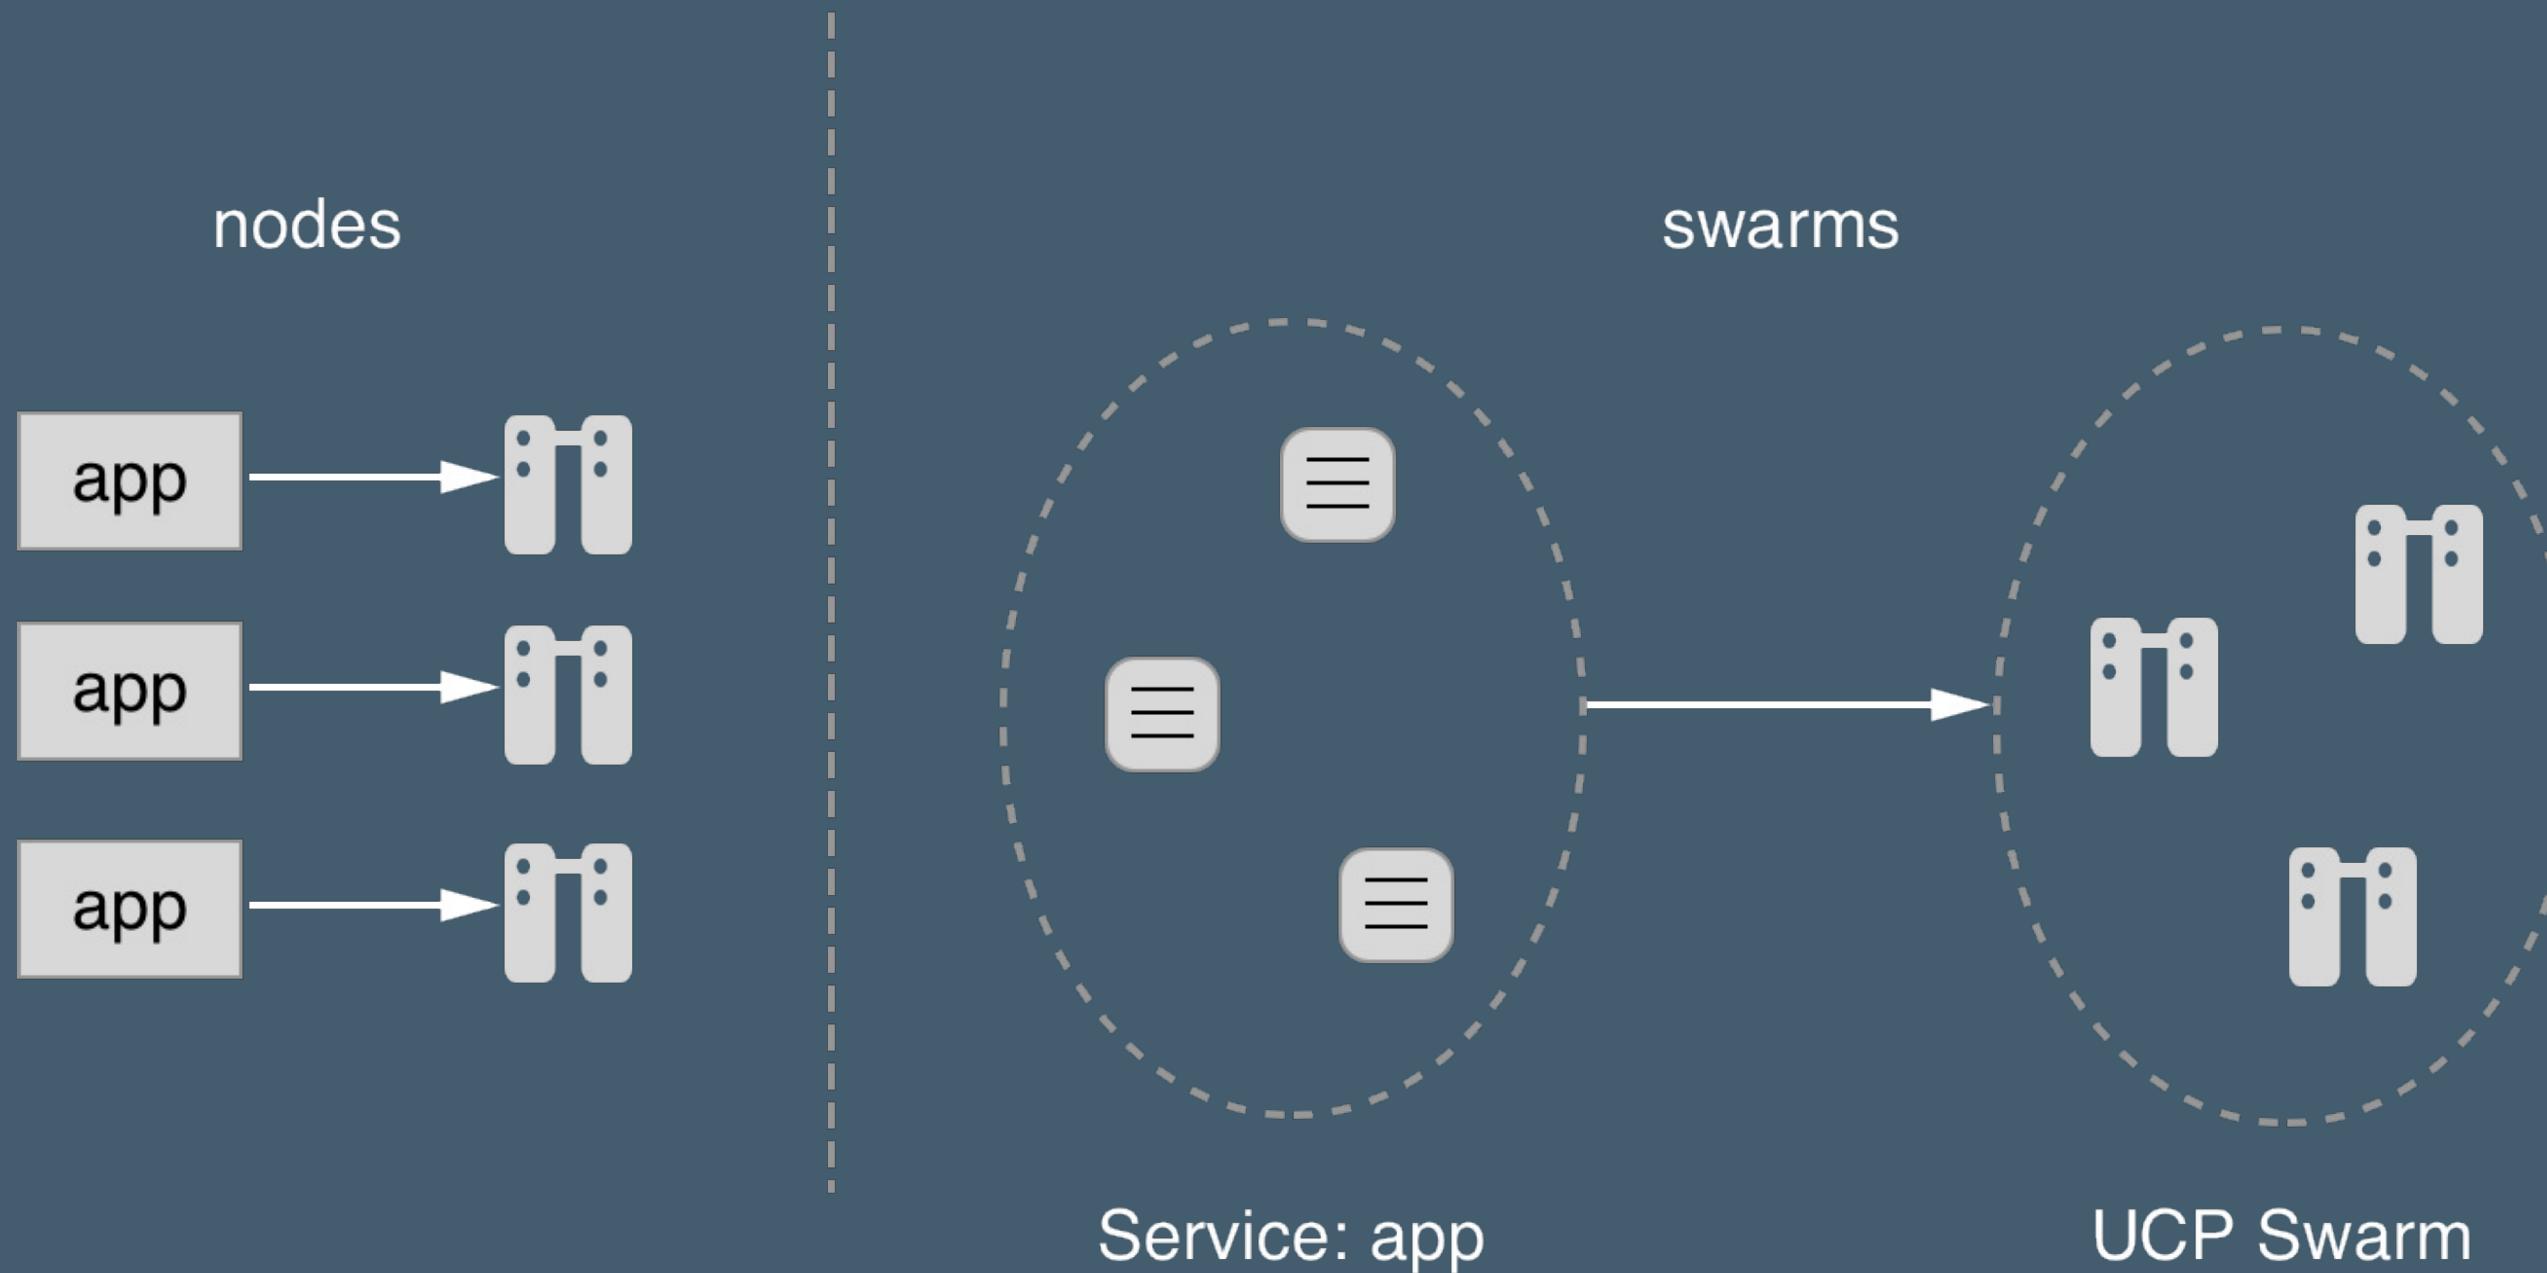
process to process



service to service



CLUSTERS OVER NODES



CHARACTERISTICS & REQUIREMENTS

- Bandwidth and latency
- Ephemeral components
- Stateless versus stateful
- Service discovery
- Load balancing
- Health checking
- Logging and monitoring
- Circuit breakers



BANDWIDTH & LATENCY

Action	Real Time	Scaled Time
1 CPU Cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid state disk I/O	50-150 us	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF -> NYC	40 ms	4 years
Internet: SF -> UK	81 ms	8 years
Internet: SF -> Australia	183 ms	19 years
OS Virtualization Reboot	4 s	423 years
SCSI Command timeout	30 s	3000 years
Hardware virtualization reboot	40 s	4000 years
Physical system reboot	5 min	32 millenia

- Enterprise apps consist of many components
- Remote calls are always expensive
- Keep distances short
- Co-locate where possible
- Use high bandwidth connections



EPHEMERAL CONTAINERS

- Pets versus cattle
- Short lifecycle
- Don't rely on specific containers
- Fast initialization / shutdown



STATELESS VERSUS STATEFUL

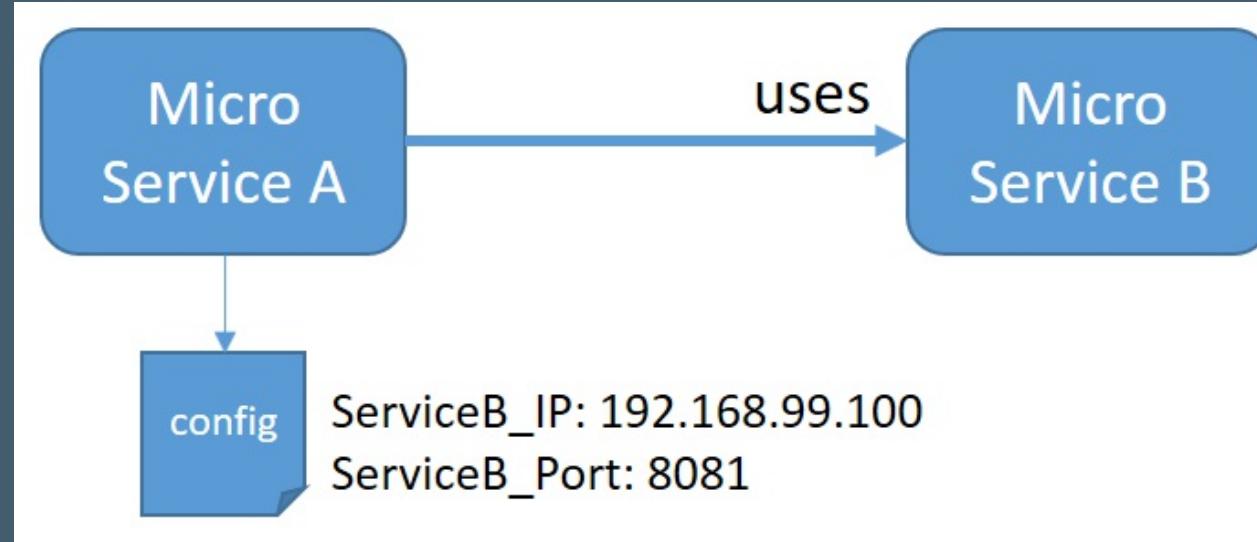
- Only stateless components scale well
- Scaling stateful components is hard
- Stateful: DB, Filesystem, Blob storage, Cache, etc
- Developers: push stateful info out to volumes and databases



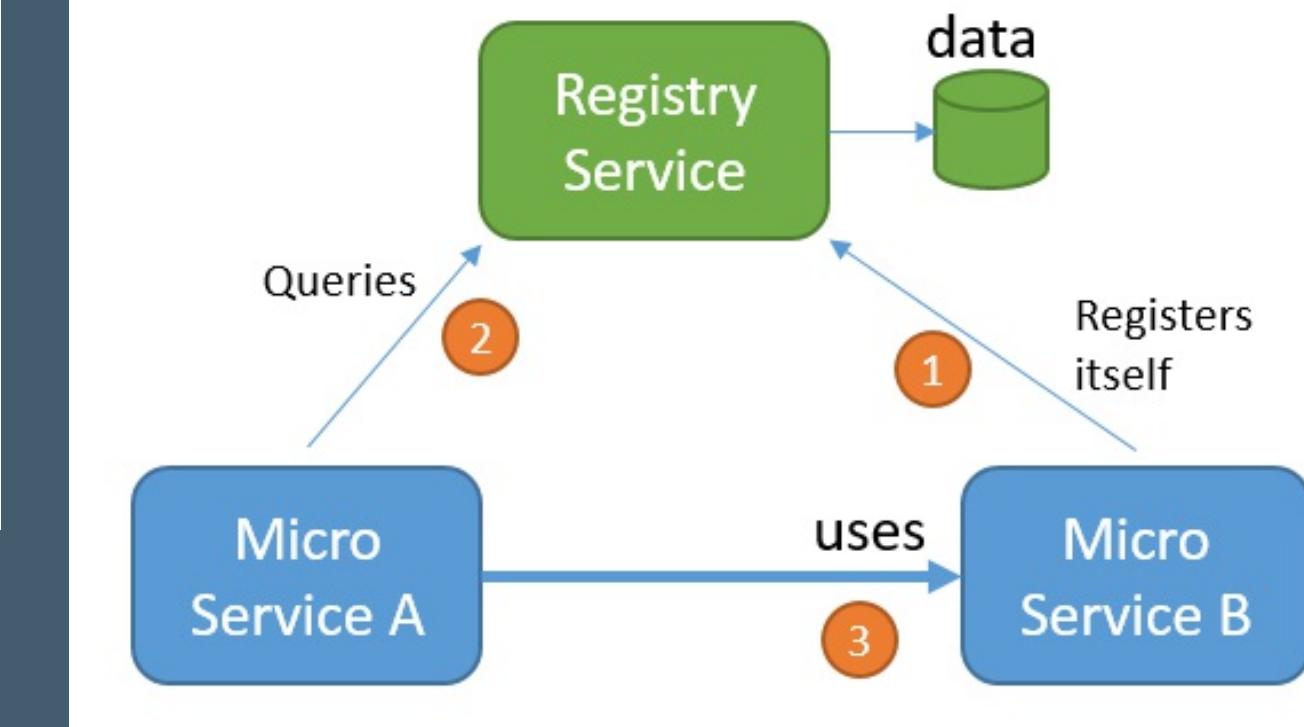
SERVICE DISCOVERY

- Containers are rescheduled frequently
- Hard-coded routes to specific containers will constantly break
- Use Docker's built-in DNS resolver for service names

Bad



Good



LOAD BALANCING

- Service-to-service communication essentially built-in load balancing
- Another reason for statelessness
- Design containers to participate in round-robin LB



HEALTH CHECKING

- Docker EE has no insight into your application logic
- Must provide healthcheck orchestrator can probe for application health
- Unhealthy containers/pods killed and restarted



LOGGING AND MONITORING

- How do we troubleshoot distributed app?
- Centralized logging, log collation & attribution
- Dashboards showing key metrics



CIRCUIT BREAKERS



Avoid cascading failures

- Use cached (stale) data
- Gracefully degrade functionality
- Give failed service time to recover



ARCHITECTURE TAKEAWAYS

Container Developer To-Dos:

- Design for services interacting with other services
- Be radically stateless
- Expect containers to start and stop all the time
- Use Docker EE's service discovery mechanisms
- Provide healthchecks for everything
- Use circuit breaking techniques for damage control
- Design logs for distributed systems



DISCUSSION

- What are the relative strengths and weaknesses of monolithic versus microservice architecture?
- Questions?



FURTHER READING

- View logs for a container or service: <http://dockr.ly/2ezdZdl>
- Docker Reference Architecture: Docker Logging Design and Best Practices: <http://dockr.ly/2gG6ZjG>
- Monitor Docker Trusted Registry: <https://dockr.ly/2HIGTIw>
- High availability architecture and apps with Docker EE: <http://dockr.ly/1sqPrIH>





SAMPLE APPLICATION



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Describe the components of a demonstration three tier web app
- Build and run the sample app using Docker Swarm or Kubernetes



SAMPLE APPLICATION ARCHITECTURE

- Backend
 - Spring Boot Java API
 - Postgres Database
- Frontend
 - Node/Express JS
 - Mustache Templates



SAMPLE APPLICATION - COMPONENTS

FRONTEND

- Node/Express JS
- Dockerfile, inherit from Node

```
FROM node:8-alpine
RUN mkdir /app
WORKDIR /app
COPY package.json /app/
RUN npm install
COPY ./src /app/src
EXPOSE 3000
CMD node src/server.js
```

BACKEND

- Maven, pom.xml
- Dockerfile, inherit from Maven
- Multi-stage build

```
FROM maven:3.5.0-jdk-8-alpine AS appserver
WORKDIR /app
COPY pom.xml .
RUN mvn dependency:resolve
COPY ..
RUN mvn package
```

```
FROM java:8-jdk-alpine
WORKDIR /app
COPY --from=appserver /app/target/pets-api-1.0.0.jar .
EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/app/pets-api-1.0.0.jar"]
```



SAMPLE APPLICATION - APP DEFINITION

Docker Compose

```
version: "3.1"

services:
  database:
    build:
      context: database
    image: ddev_db
    environment:
      POSTGRES_USER: gordonuser
      POSTGRES_DB: ddev
    ports:
      - "5432:5432"
    networks:
      - back-tier
    secrets:
      - postgres_password

  api:
    build:
      context: api
    image: ddev_api
    ports:
      - "8080:8080"
      - "5005:5005"

networks:
  - front-tier
  - back-tier
secrets:
  - postgres_password

ui:
  build:
    context: ui
  image: ddev_ui
  ports:
    - "3000:3000"
  networks:
    - front-tier

secrets:
  postgres_password:
    file: ./devsecrets/postgres_password

networks:
  front-tier:
  back-tier:
```





EXERCISE: WORKSHOP SAMPLE APPLICATION

Work through the 'Workshop Sample Application' exercise in the Docker for Enterprise Developers Exercises book.



FURTHER READING

- Sample applications: <http://dockr.ly/2eU89rn>
- Dockerize a .NET core application: <http://dockr.ly/2jb6GOQ>
- Dockerizing sample applications: <http://dockr.ly/2jb4B5s>





EDIT AND CONTINUE



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Mount a development environment inside a containerized execution environment
- Configure automatic rebuilds and restarts of containerized processes



CONTAINERS ADD FRICTION?

- Use your favorite code editor
- Code lives on host
- Application runs in container

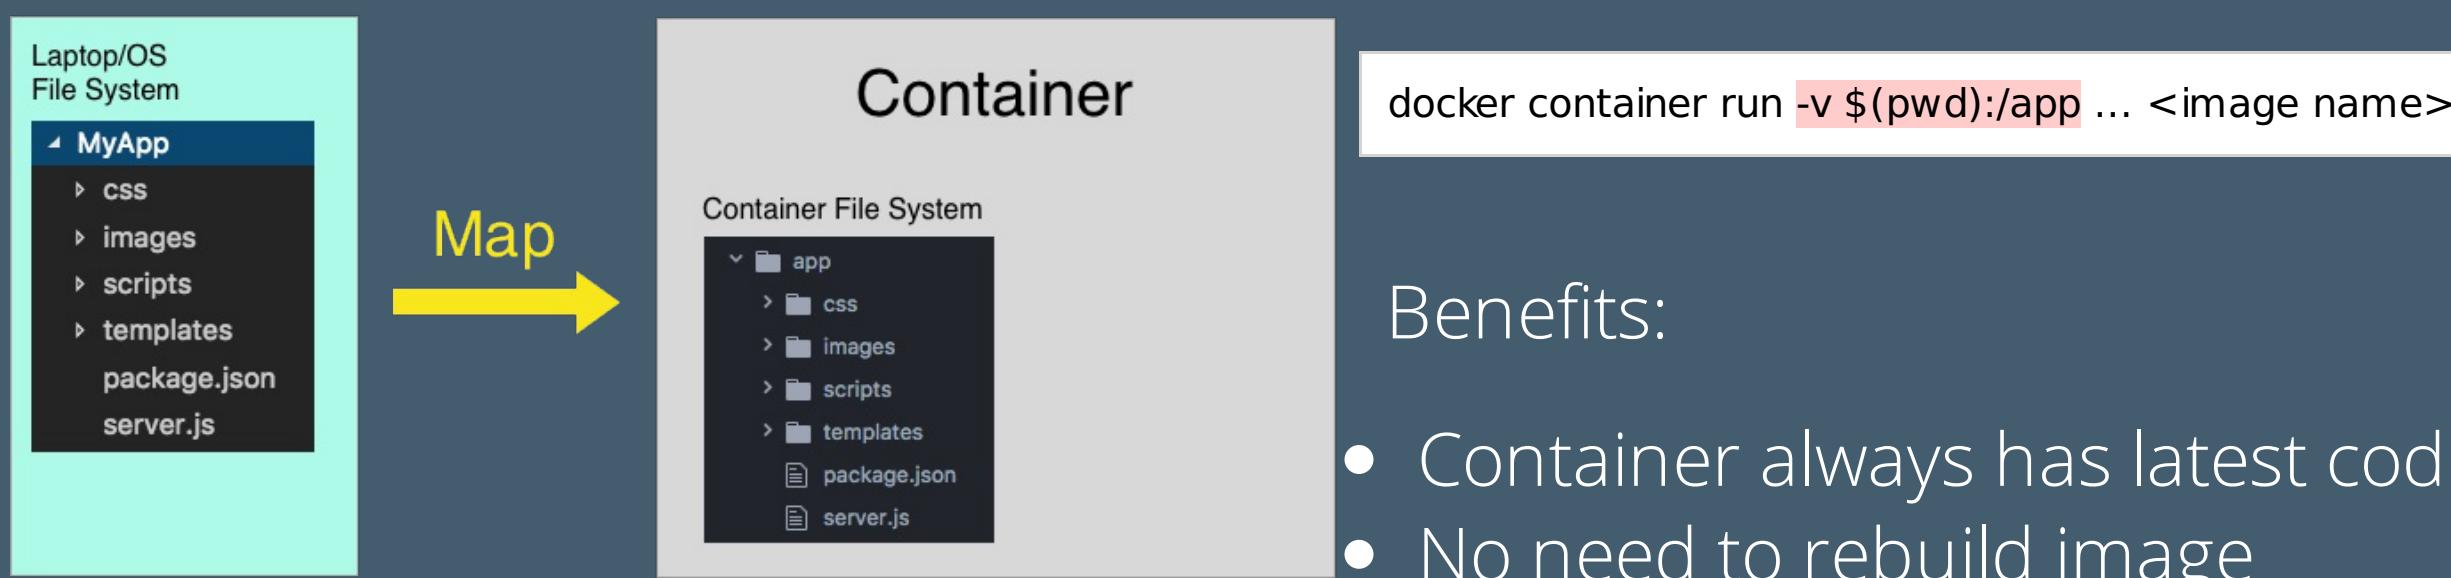
Resulting process:

1. Add or edit code
2. Build the container image
3. Run a container with the new image
4. Test
5. Start over at 1...



MOUNTING CODE

Mounting the source folder into the container:



```
docker container run -v $(pwd):/app ... <image name>
```

Benefits:

- Container always has latest code
- No need to rebuild image

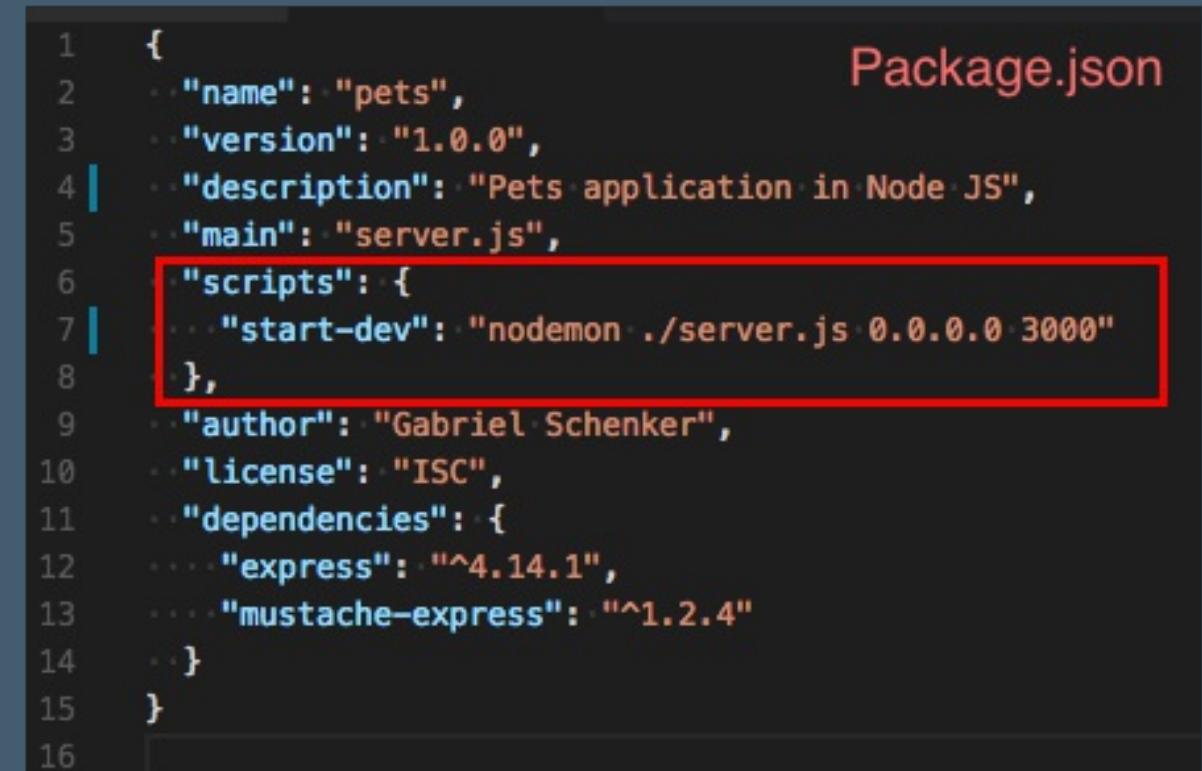


AUTO RESTART APP

Restart web server upon file change...

```
nodemon ./server.js 0.0.0.0 3000
```

Node/Express Application



Package.json

```
1  {
2    "name": "pets",
3    "version": "1.0.0",
4    "description": "Pets application in Node JS",
5    "main": "server.js",
6    "scripts": {
7      "start-dev": "nodemon ./server.js 0.0.0.0 3000"
8    },
9    "author": "Gabriel Schenker",
10   "license": "ISC",
11   "dependencies": {
12     "express": "^4.14.1",
13     "mustache-express": "^1.2.4"
14   }
15 }
16 
```





EXERCISE: EDIT AND CONTINUE

Work through the 'Edit and Continue' exercise in the Docker for Enterprise Developers Exercises book.



FURTHER READING

- Use volumes: <http://dockr.ly/2vRZBDG>
- Use bind mounts: <http://dockr.ly/2wdstvn>





DEBUGGING



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Connect an IDE to a process running in a container.
- Debug containerized processes as if they were remote processes.



DEBUGGING

- We want:
 - Line by line debugging
 - Variable inspection and modification
 - Stack trace
- How to attach an IDE to a containerized process?
- Debugging in container == Remote debugging



DEBUGGING

SUPPORT THROUGH IDE

Language	IDE/Tool	Example
Java	IntelliJ, Eclipse, Netbeans	http://dockr.ly/2rmLhlU , http://bit.ly/2rFoALR
.NET Core	Visual Studio Code	http://bit.ly/2rdVAKX
Node JS	Visual Studio Code	http://dockr.ly/2ac5TVw
	Webstorm	http://bit.ly/2qDcNe1
Ruby	RubyMine	http://bit.ly/2rmPgzd
C++	VisualGDB	http://bit.ly/2rHvcdN
Python/Django	PyCharm	http://bit.ly/2rdFnFg , http://bit.ly/2rn3zDy





EXERCISE: DEBUGGING

Work through the 'Debugging in a Container' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What would you use to keep an eye on the metrics of your container? Why?
- What would you do if you wanted to debug why your container won't start at all (e.g.: initial command or entrypoint that immediately crashes)?
- Questions?



FURTHER READING

- Live debugging with Docker: <http://dockr.ly/2ac5TVw>
- Live debugging Java with Docker: <http://dockr.ly/2rmLhlU>





DOCKER COMPOSE



ee2.0-v2.1 © 2018 Docker, Inc.

LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Describe and deploy a multi-service app using Docker Compose
- Capture configuration differences in environment-specific compose files



DOCKER COMPOSE

Define multiple Docker Compose files

- For development (front- vs. back-end)
- Running unit- and/or integration tests
- Running end-to-end tests
- Build artifacts on CI server
- Run tests on CI server



BACKEND DEVELOPMENT

```
version: '3.1'
services:
  api:
    build:
      context: api
      dockerfile: Dockerfile.dev
    ports:
      - 8080:8080
    volumes:
      - ./api:/usr/src/pets-api
    networks:
      - pets
  watcher:
    build:
      context: api
      dockerfile: Dockerfile.dev
    volumes:
      - ./api:/usr/src/pets-api
    command: mvn fizzed-watcher:run
    networks:
      - pets
networks:
  pets:
```

- Development specific Dockerfile
- Mounting code into container
- Opening port to access via e.g. **curl**
- Watcher container to auto-compile



FRONTEND DEVELOPMENT

```
version: '3.1'
services:
  api:
    image: <username>/pets-api:1.4
    networks:
      - pets
  ui:
    build:
      context: ui
      dockerfile: Dockerfile-dev
    ports:
      - 3000:3000
    volumes:
      - ./ui/src:/app/src
      - ./ui/package.json:/app/package.json
    command: nodemon src/server.js 0.0.0.0 3000
    networks:
      - pets
networks:
  pets:
```

- Use specific version of backend image
- Use development specific Dockerfile
- Mount UI code into container
- Use auto restart in case of change





EXERCISE: DOCKER COMPOSE

Work through the 'Docker Compose' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What are some pros and cons of having different compose files for different environments?
- Questions?



FURTHER READING

- Docker compose reference: <http://dockr.ly/2iHUpex>
- Overview of Docker compose: <http://dockr.ly/2jh9kjv>





TESTING



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Create unit, API, stress and integration tests appropriate for containerized applications.

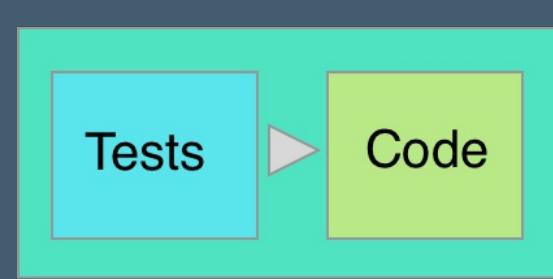


CATEGORIES & CHARACTERISTICS OF TESTS

- Unit Tests
- Integration Tests
- API Tests
- Stress and Load Tests
- End-to-End Tests



UNIT TESTS

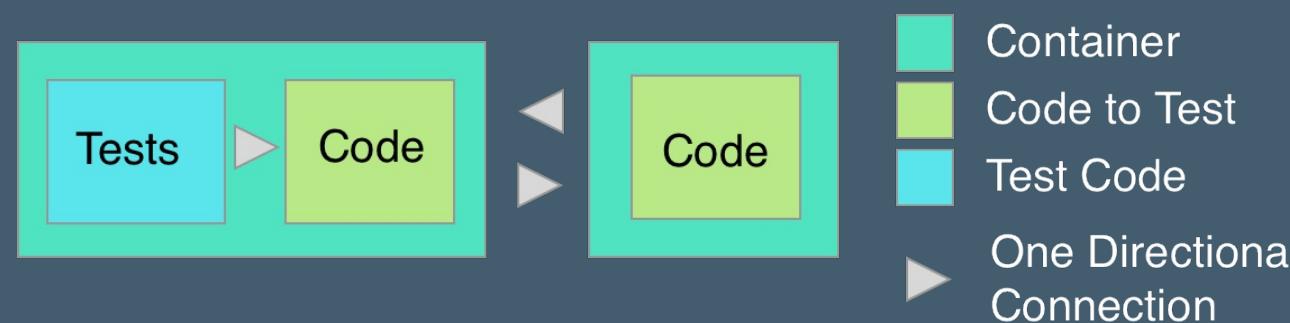


 Container
 Code to Test
 Test Code
► One Directional Connection

- Test a small & isolated piece of code
- All external dependencies are mocked
- Test code is tightly coupled with SUT
- Test code and SUT run in same container



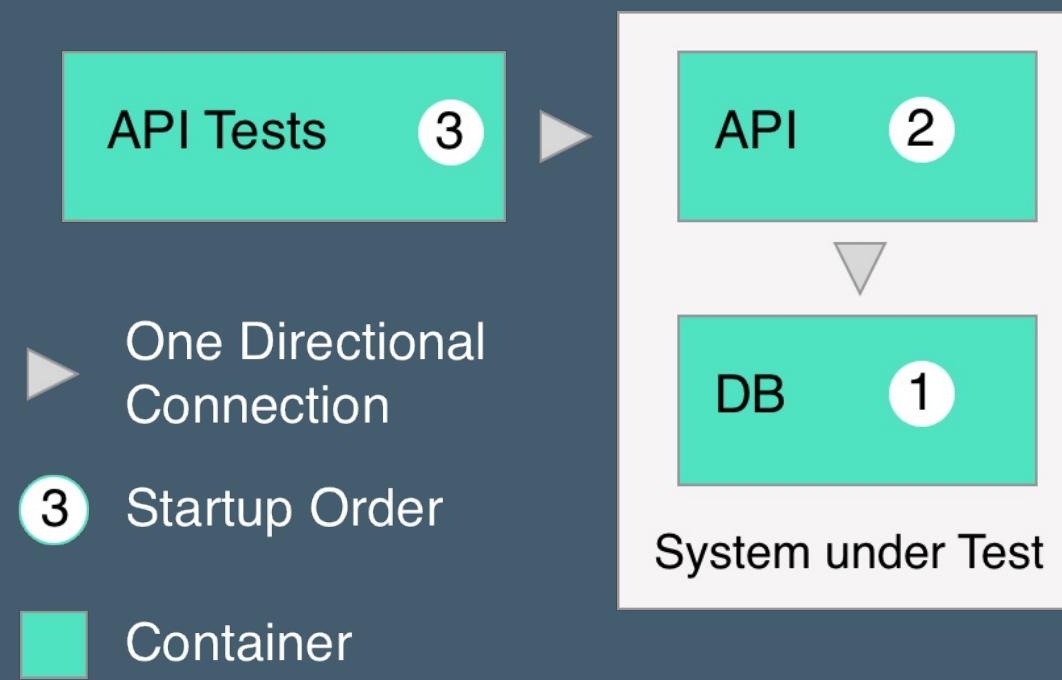
INTEGRATION TESTS



- Test & validate multiple interacting components
- Some external dependencies are mocked
- Test code is tightly coupled with SUT
- Test code and SUT run in same container



API TESTS



- Test & validate service via public API
- Some (expensive) external dependencies can be mocked
- Test code is loosely coupled with SUT
- Test code and SUT run in different containers

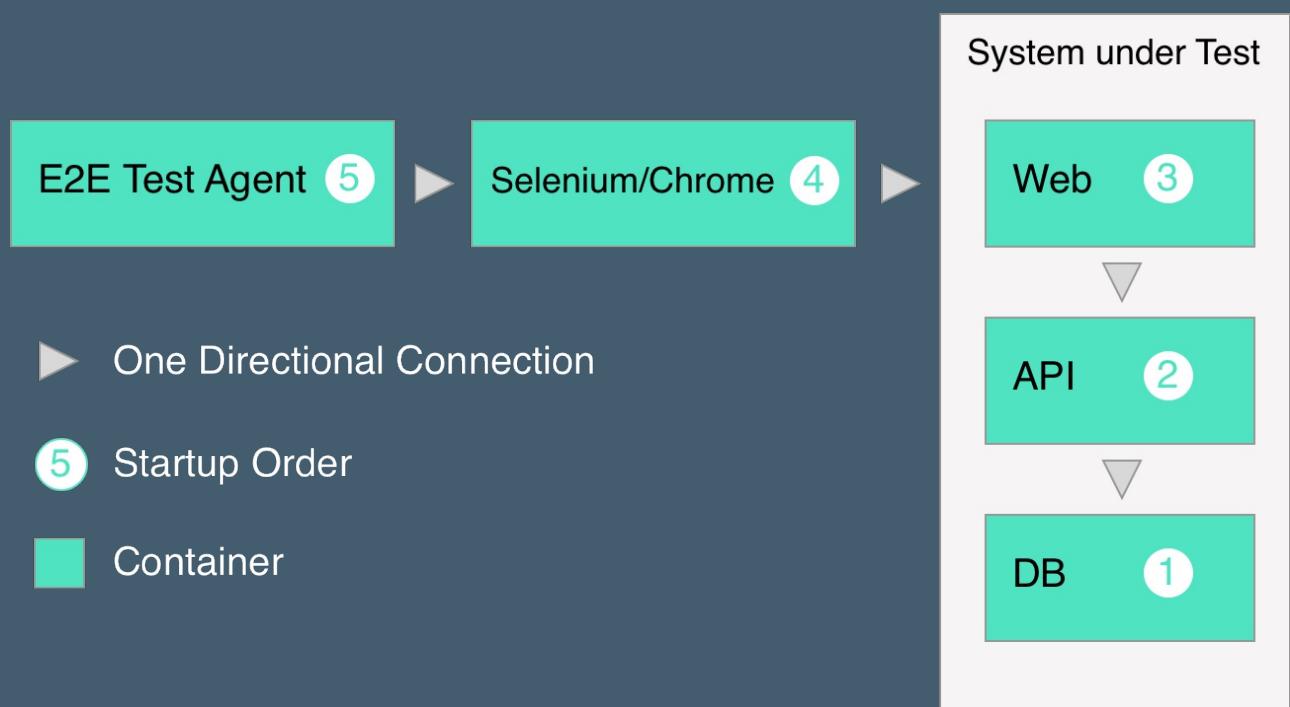


STRESS & LOAD TESTS

- Test code under load
- Use production like environments, data quantities and concurrent connections
- Test code is loosely coupled with SUT
- Test code and SUT run in different containers



END-TO-END TESTS



- Test whole application end-2-end
- Use production like environments, data quantities and concurrent connections
- Test code automates UI
- Test code and SUT run in different containers





EXERCISE: UNIT, API & E2E TESTS

Work through the exercises

- Unit Tests
- API Tests
- End-to-End Tests

in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What new tests become more important after refactoring a monolithic application into multiple containerized services?
- Questions?



FURTHER READING

- Continuous integration testing with Docker: <http://dockr.ly/2xSILXM>
- Scale testing Docker swarm to 30000 containers: <http://dockr.ly/2smWgvf>





SERVICE DISCOVERY



LEARNING OBJECTIVES

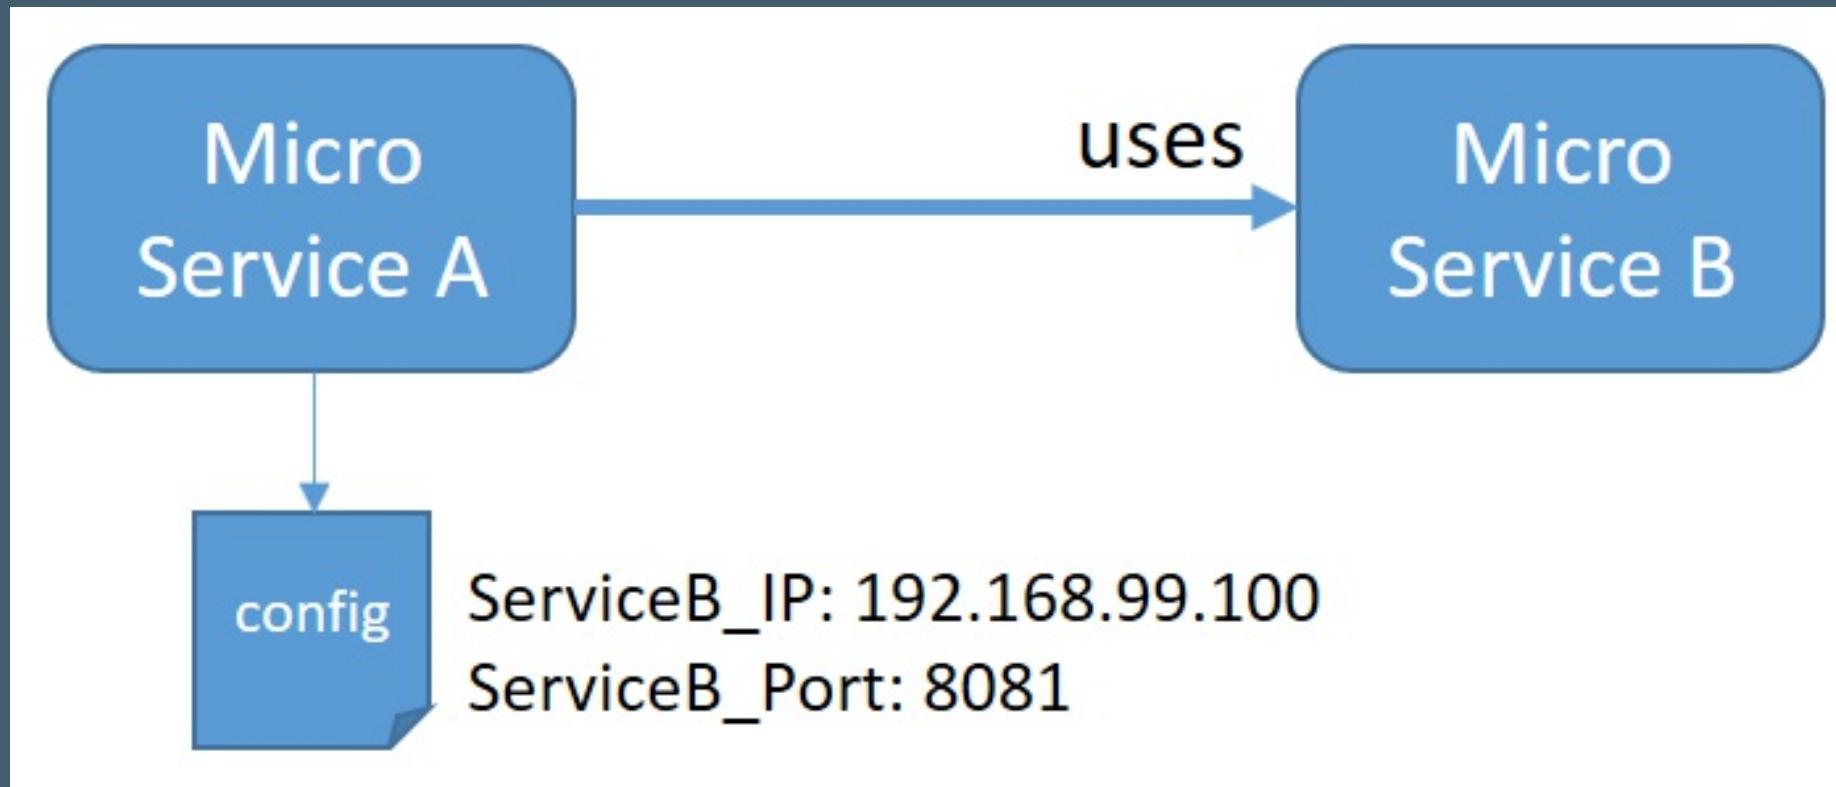
By the end of this module, learners will be able to:

- Route traffic between containers using orchestrator-driven service discovery
- Compare and contrast service discovery in Swarm and Kubernetes



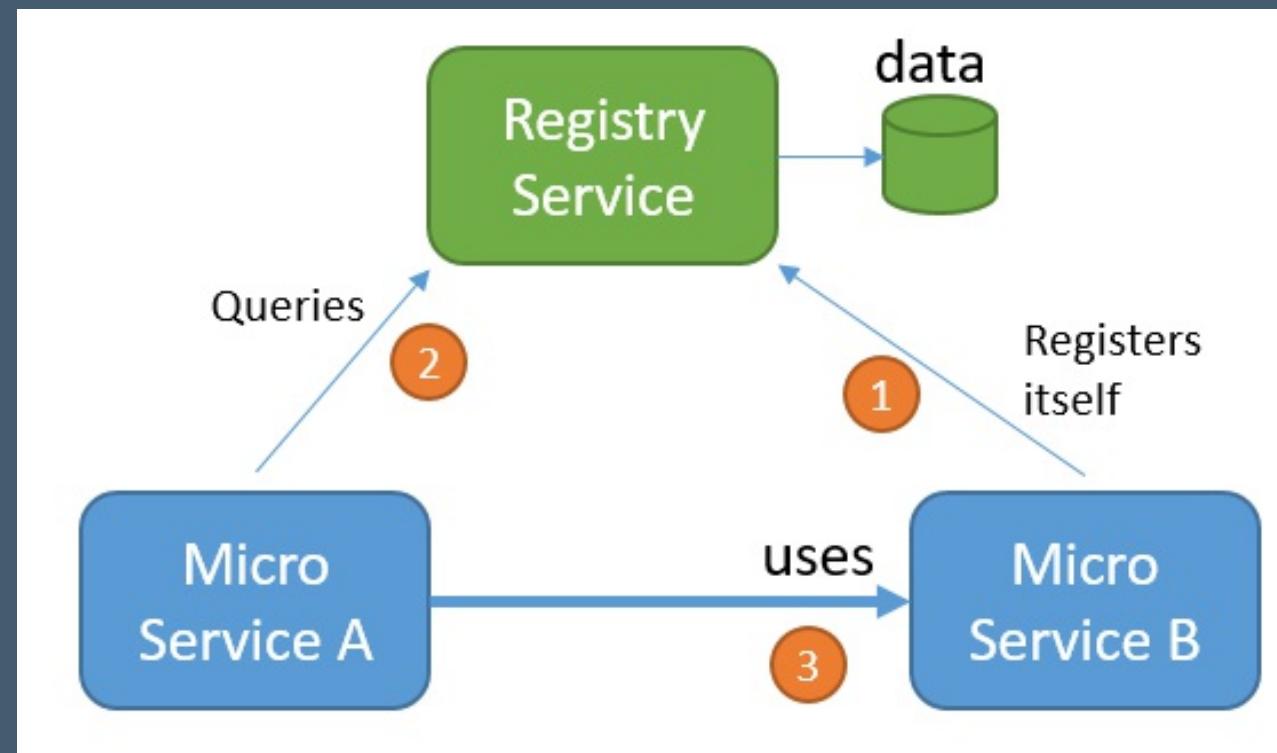
SERVICE DISCOVERY

Bad: Hard Wiring...



SERVICE DISCOVERY

Good: Registry Service

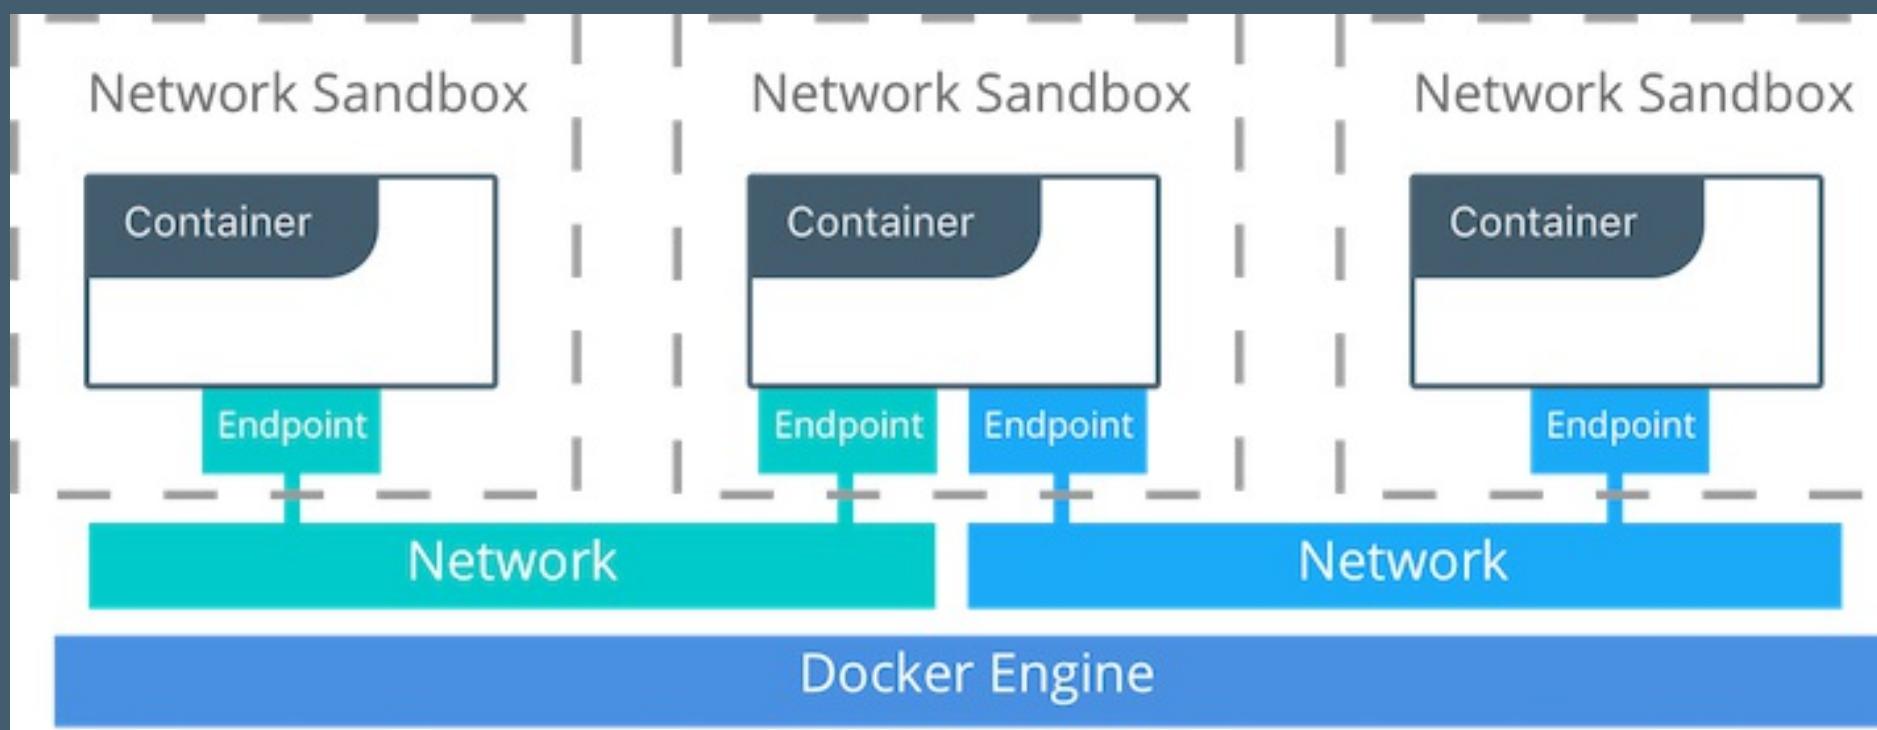


- Docker provides embedded DNS for container and service names
- Address traffic to services by name

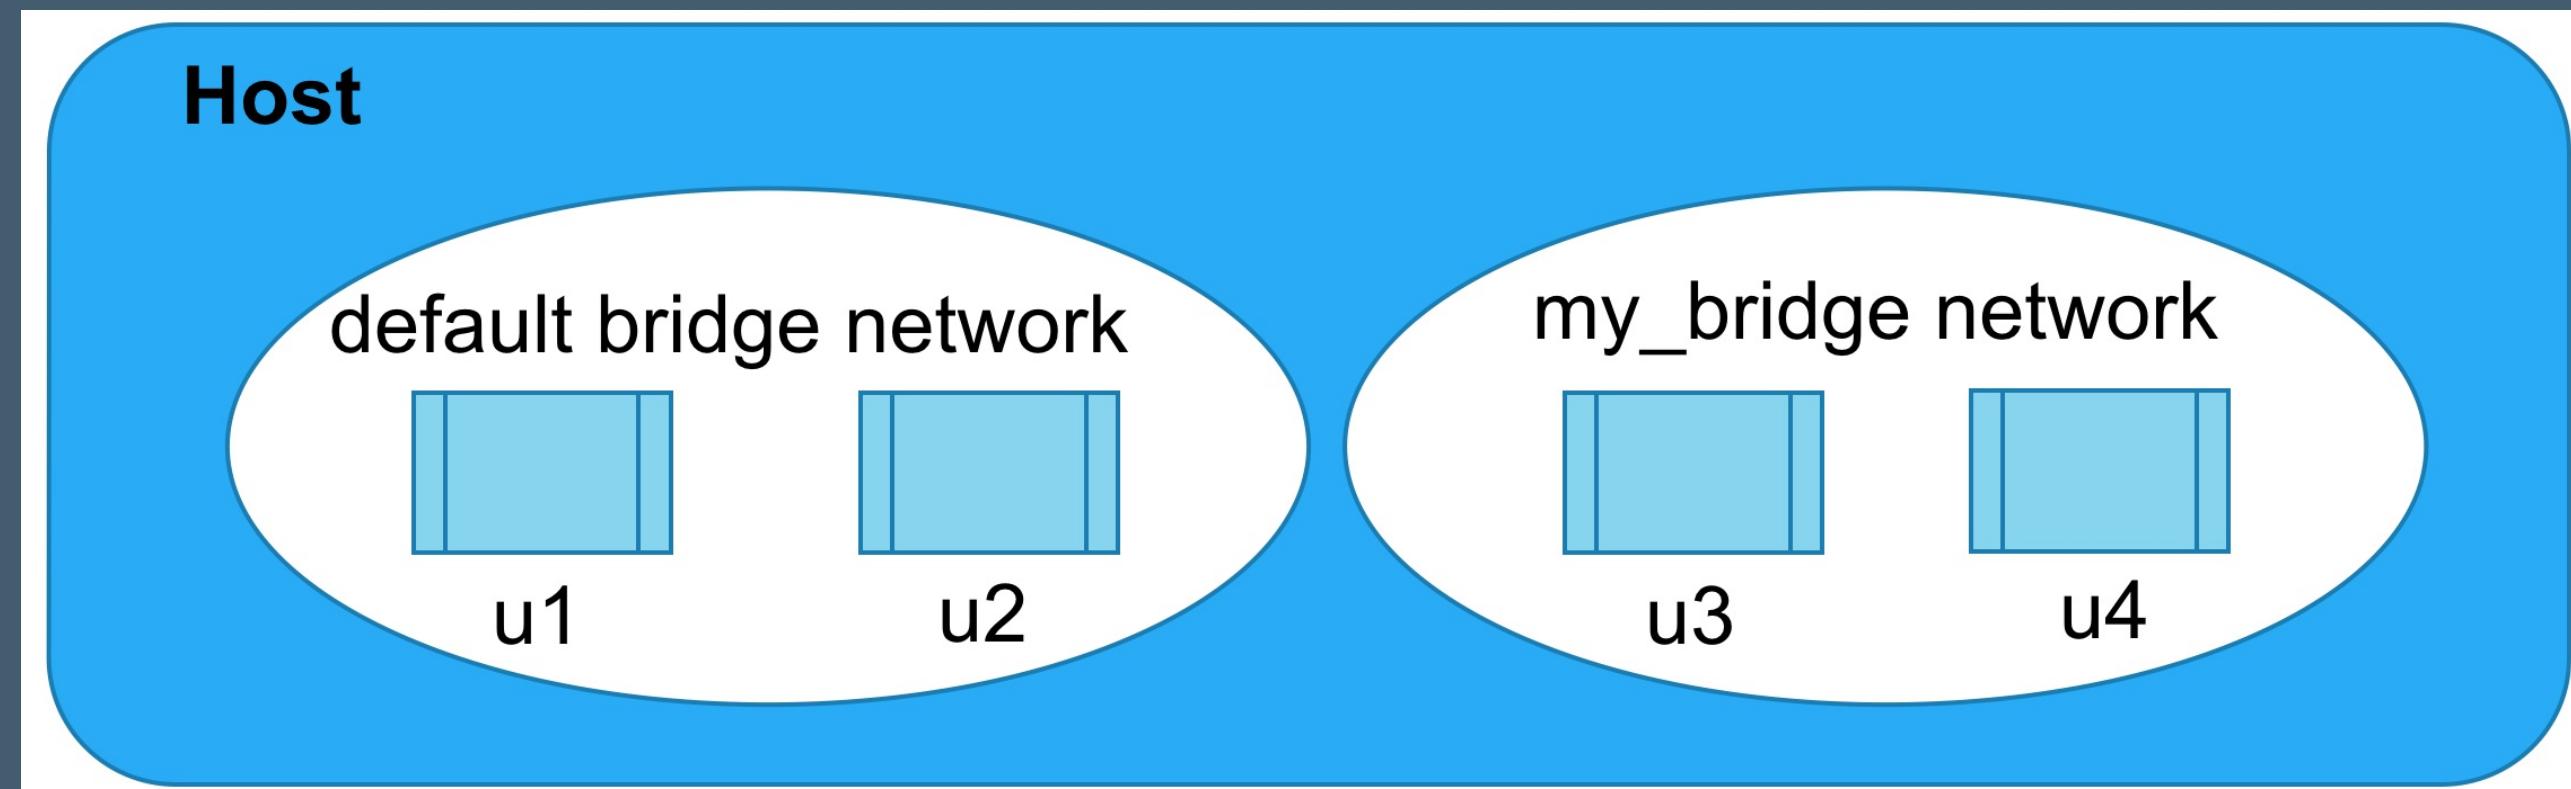


DOCKER SWARM

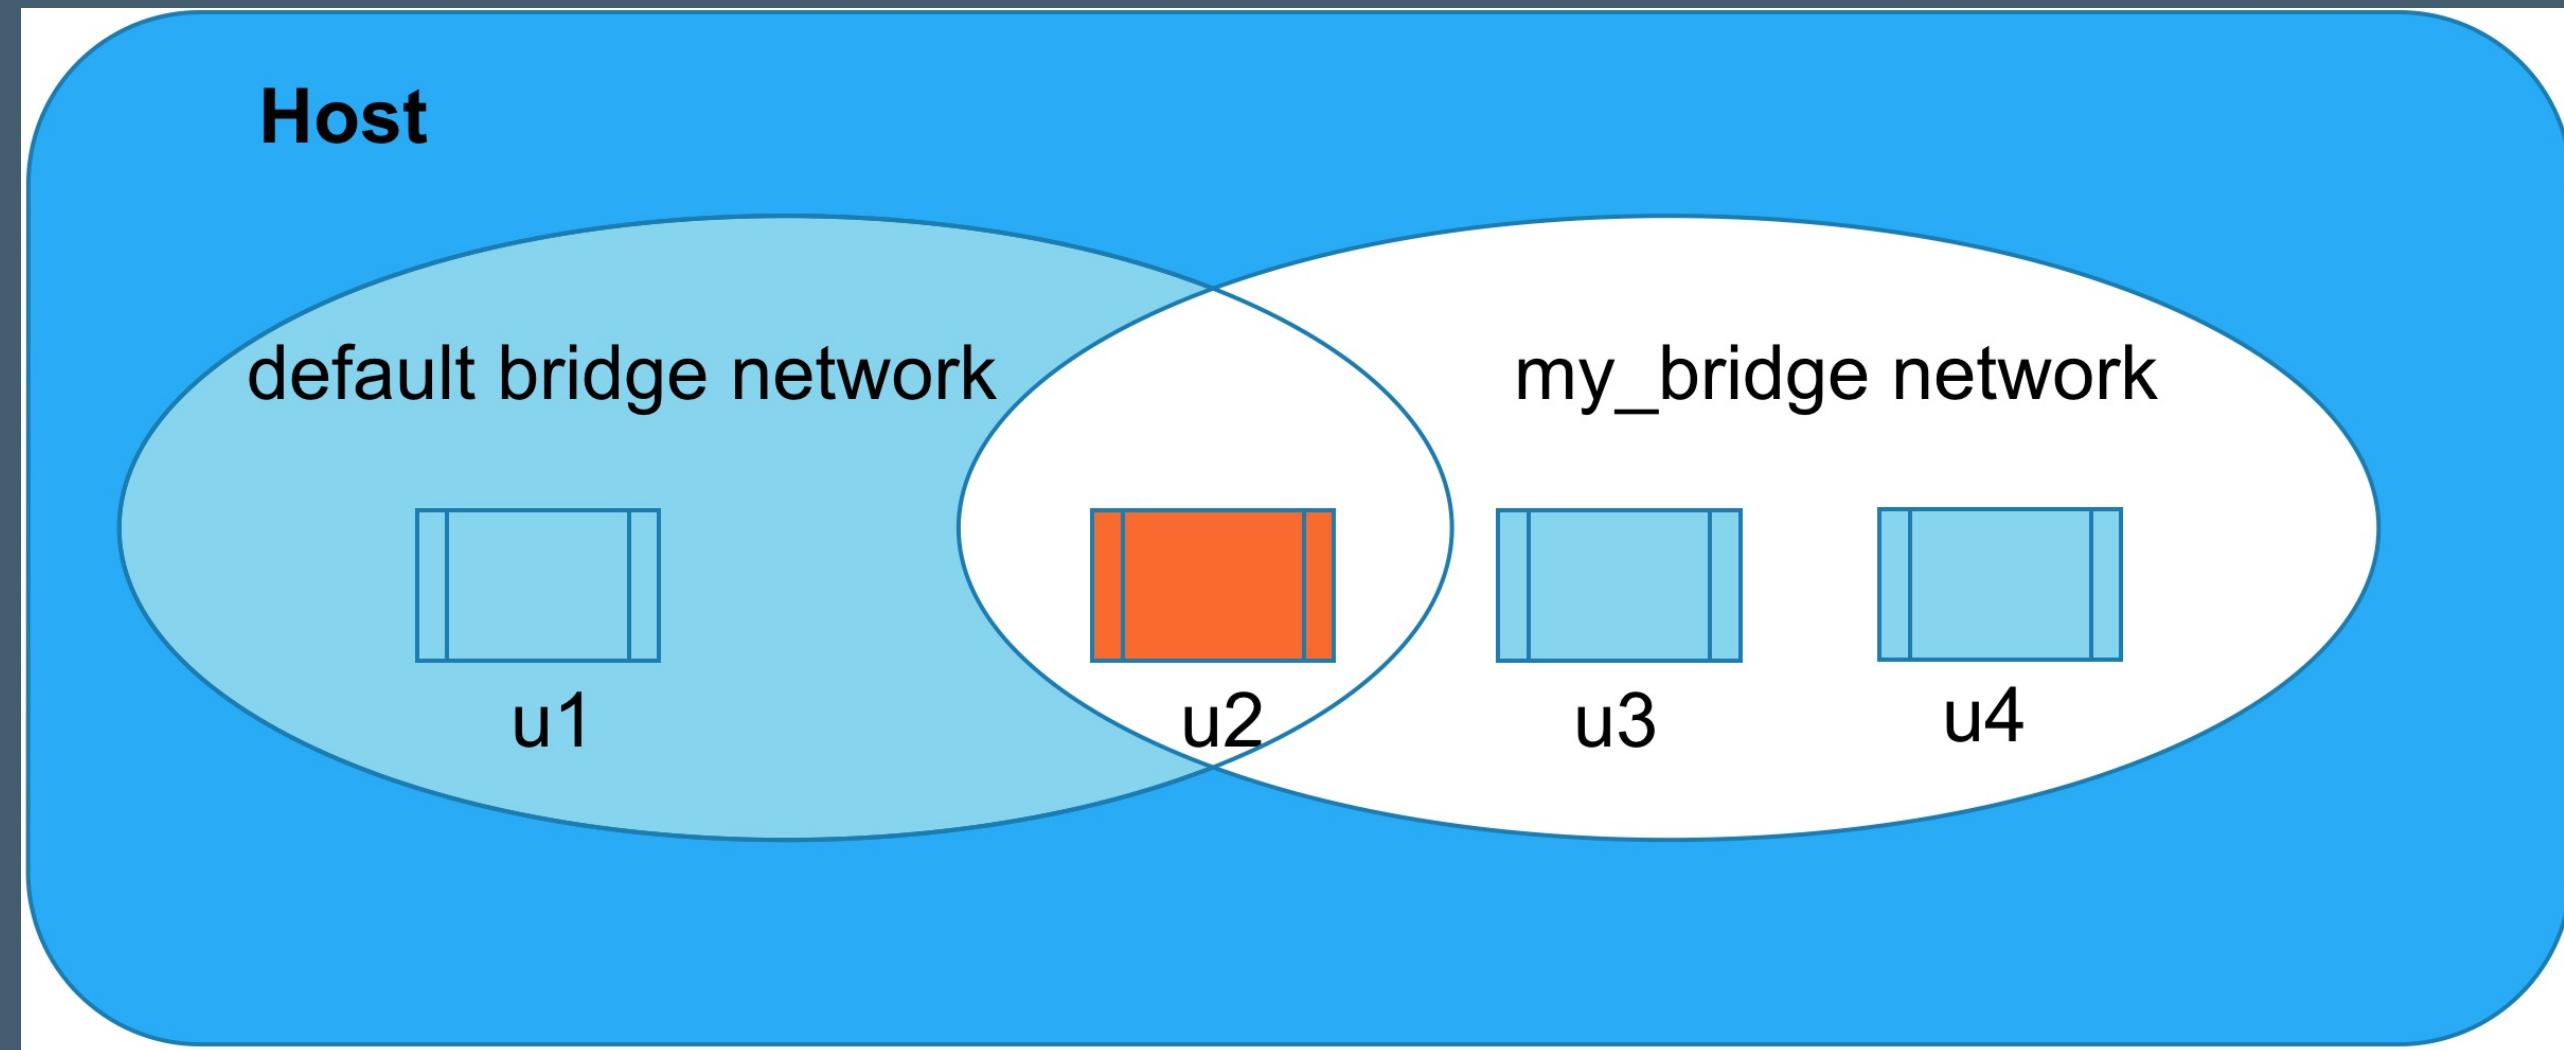
CONTAINER NETWORK MODEL



DOCKER SWARM: NETWORK FIREWALLS



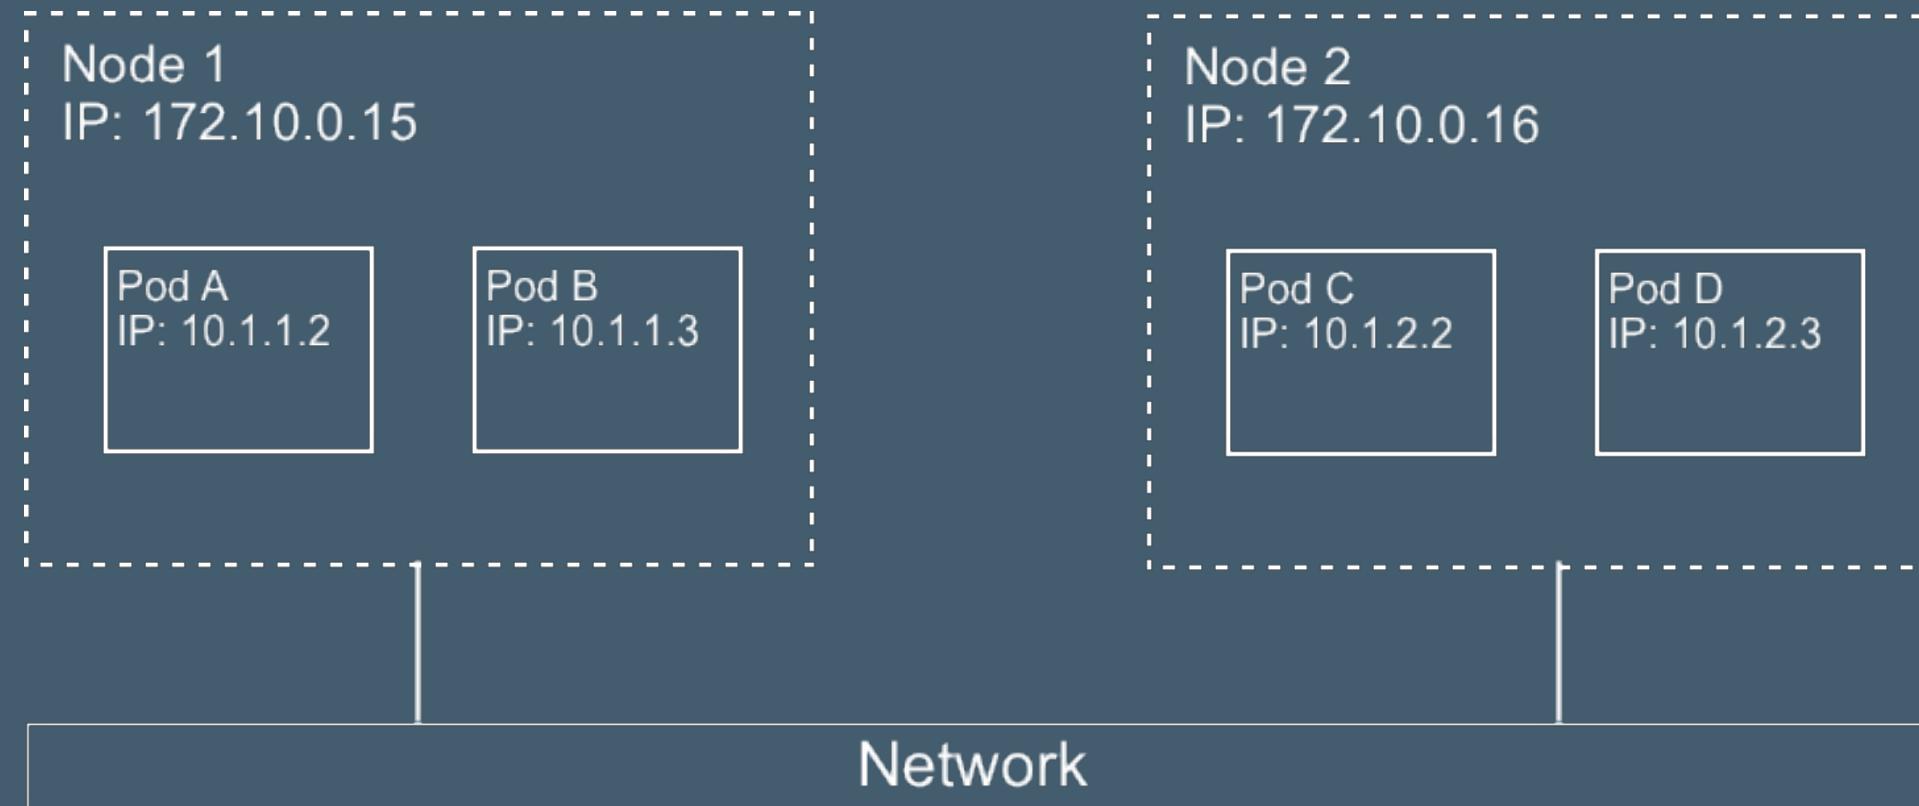
DOCKER SWARM: NETWORK FIREWALLS



docker network connect my_bridge u2



KUBERNETES: NETWORKING

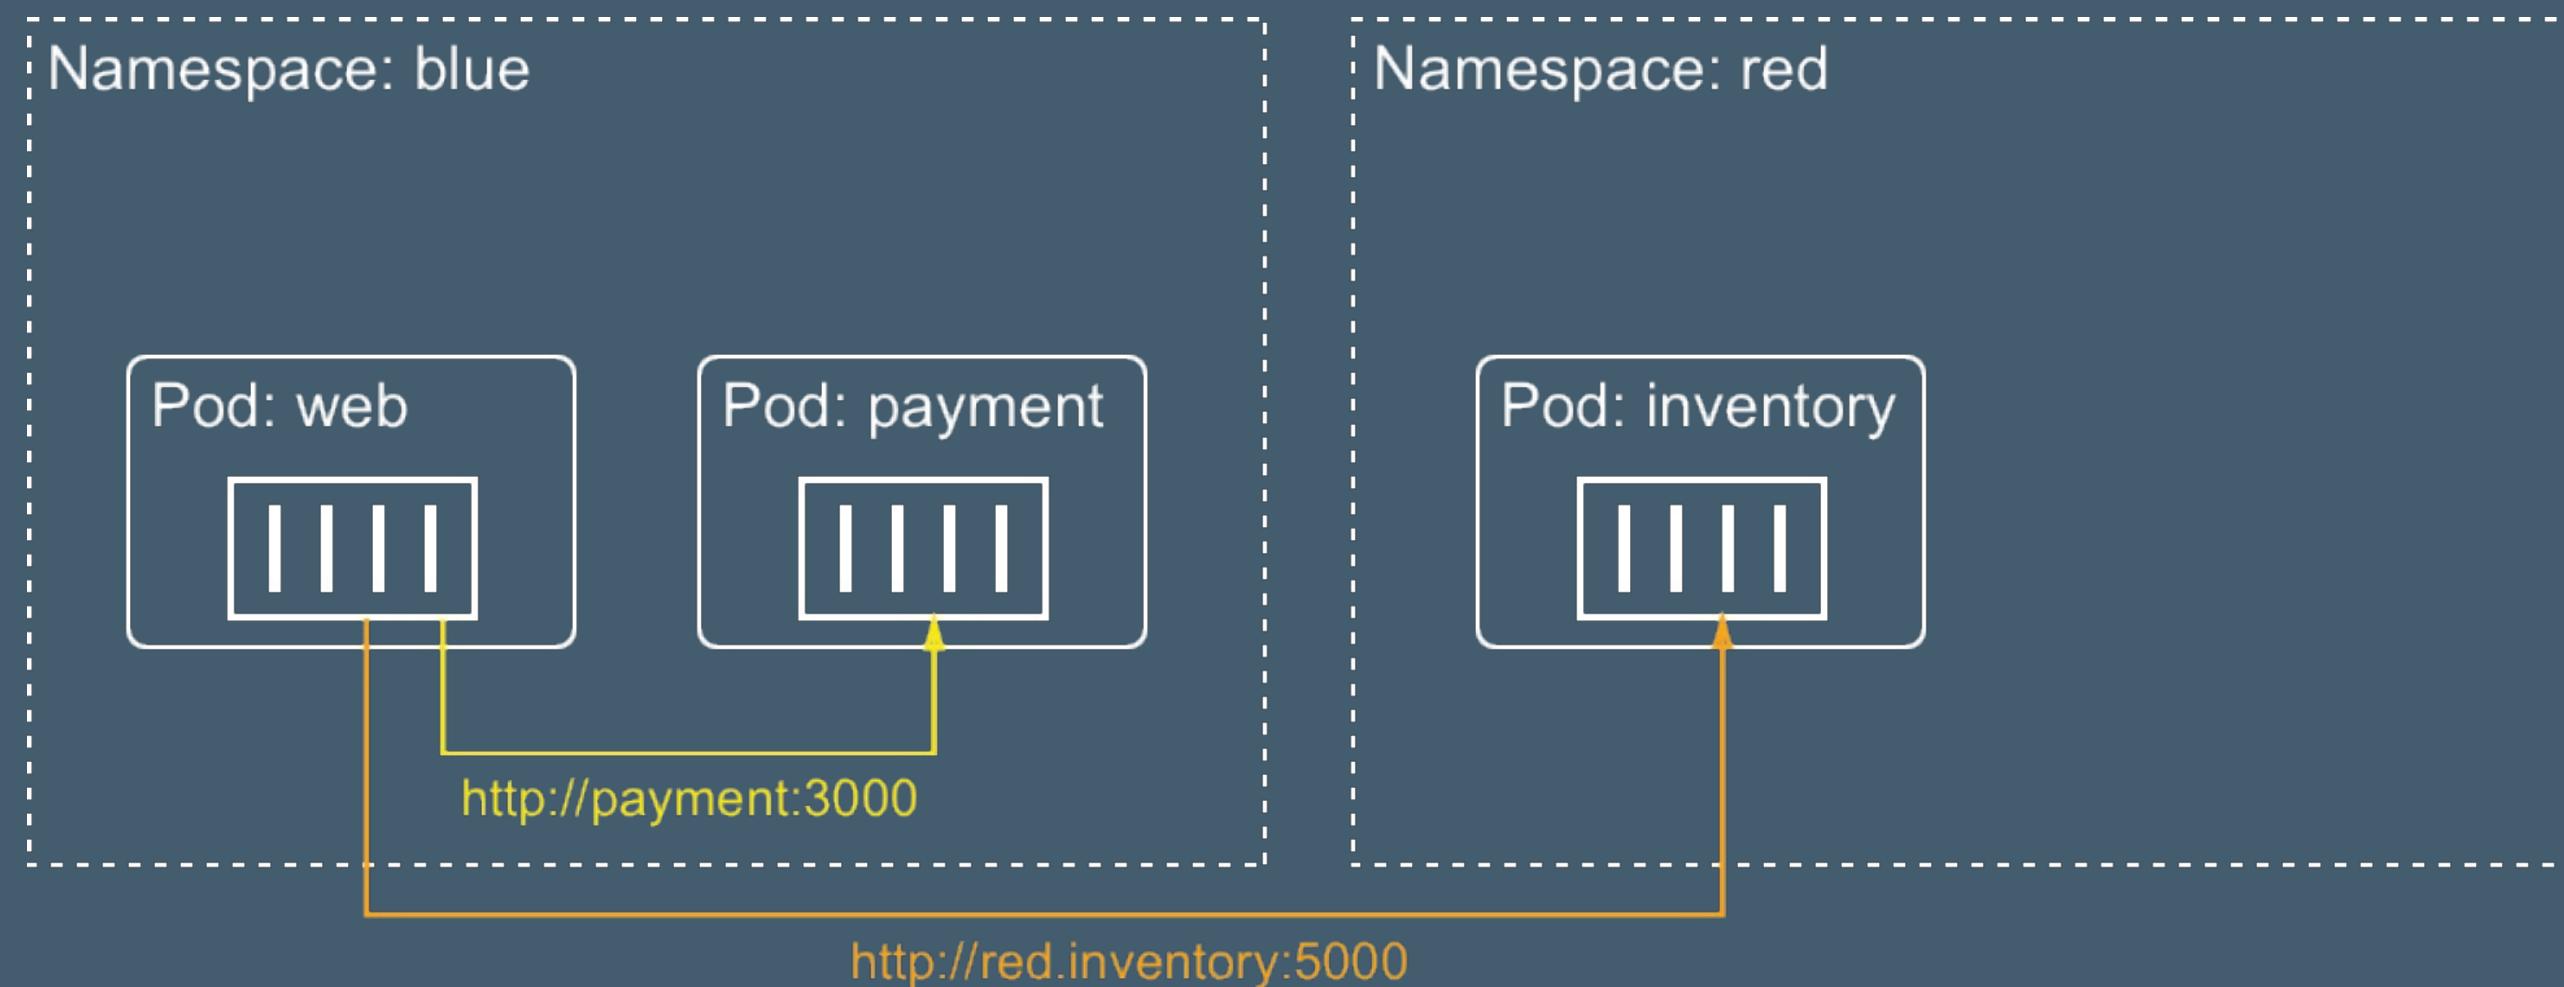


Requirements

- Pod <--> Pod without NAT
- Node <--> Pod without NAT
- Pod's peers find it at the same IP it finds itself



KUBERNETES: NAMESPACES





EXERCISE: SERVICE DISCOVERY

Work through the 'Service Discovery' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What are some pros and cons of orchestrator-based service discovery?
- Questions?



FURTHER READING

- Service discovery and links: <http://dockr.ly/2gQlctq>
- Manage swarm service networks: <http://dockr.ly/2gQjG82>
- Docker Reference Architecture: UCP 2.0 Service Discovery and Load Balancing:
<http://dockr.ly/2rbxDDX>





HEALTH CHECKS



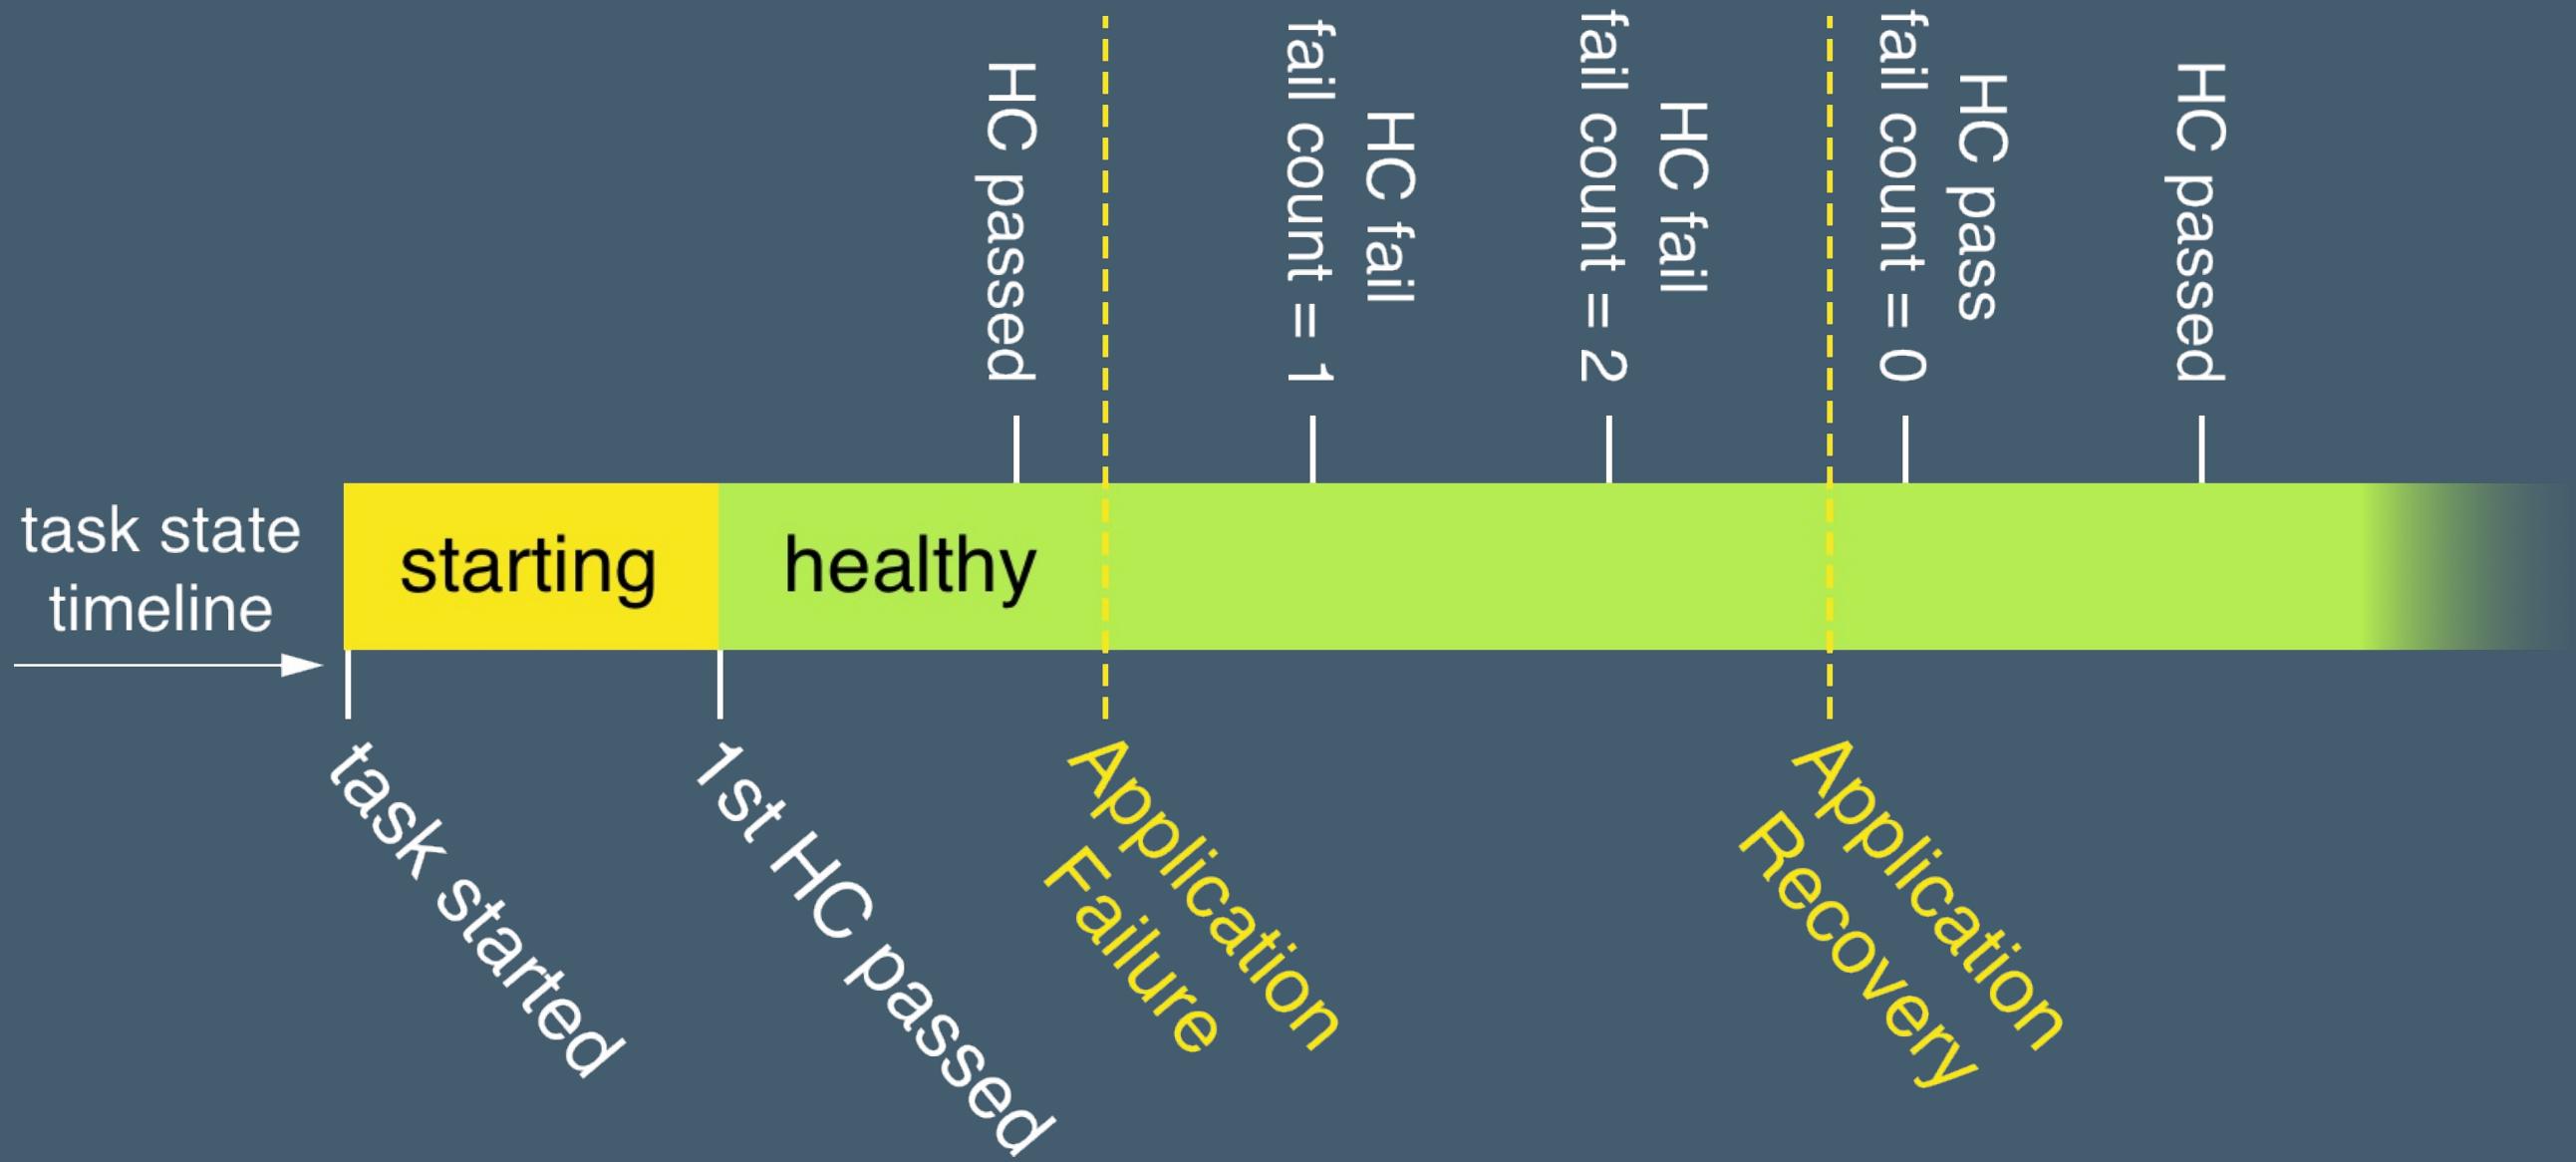
LEARNING OBJECTIVES

By the end of this module, learners will be able to:

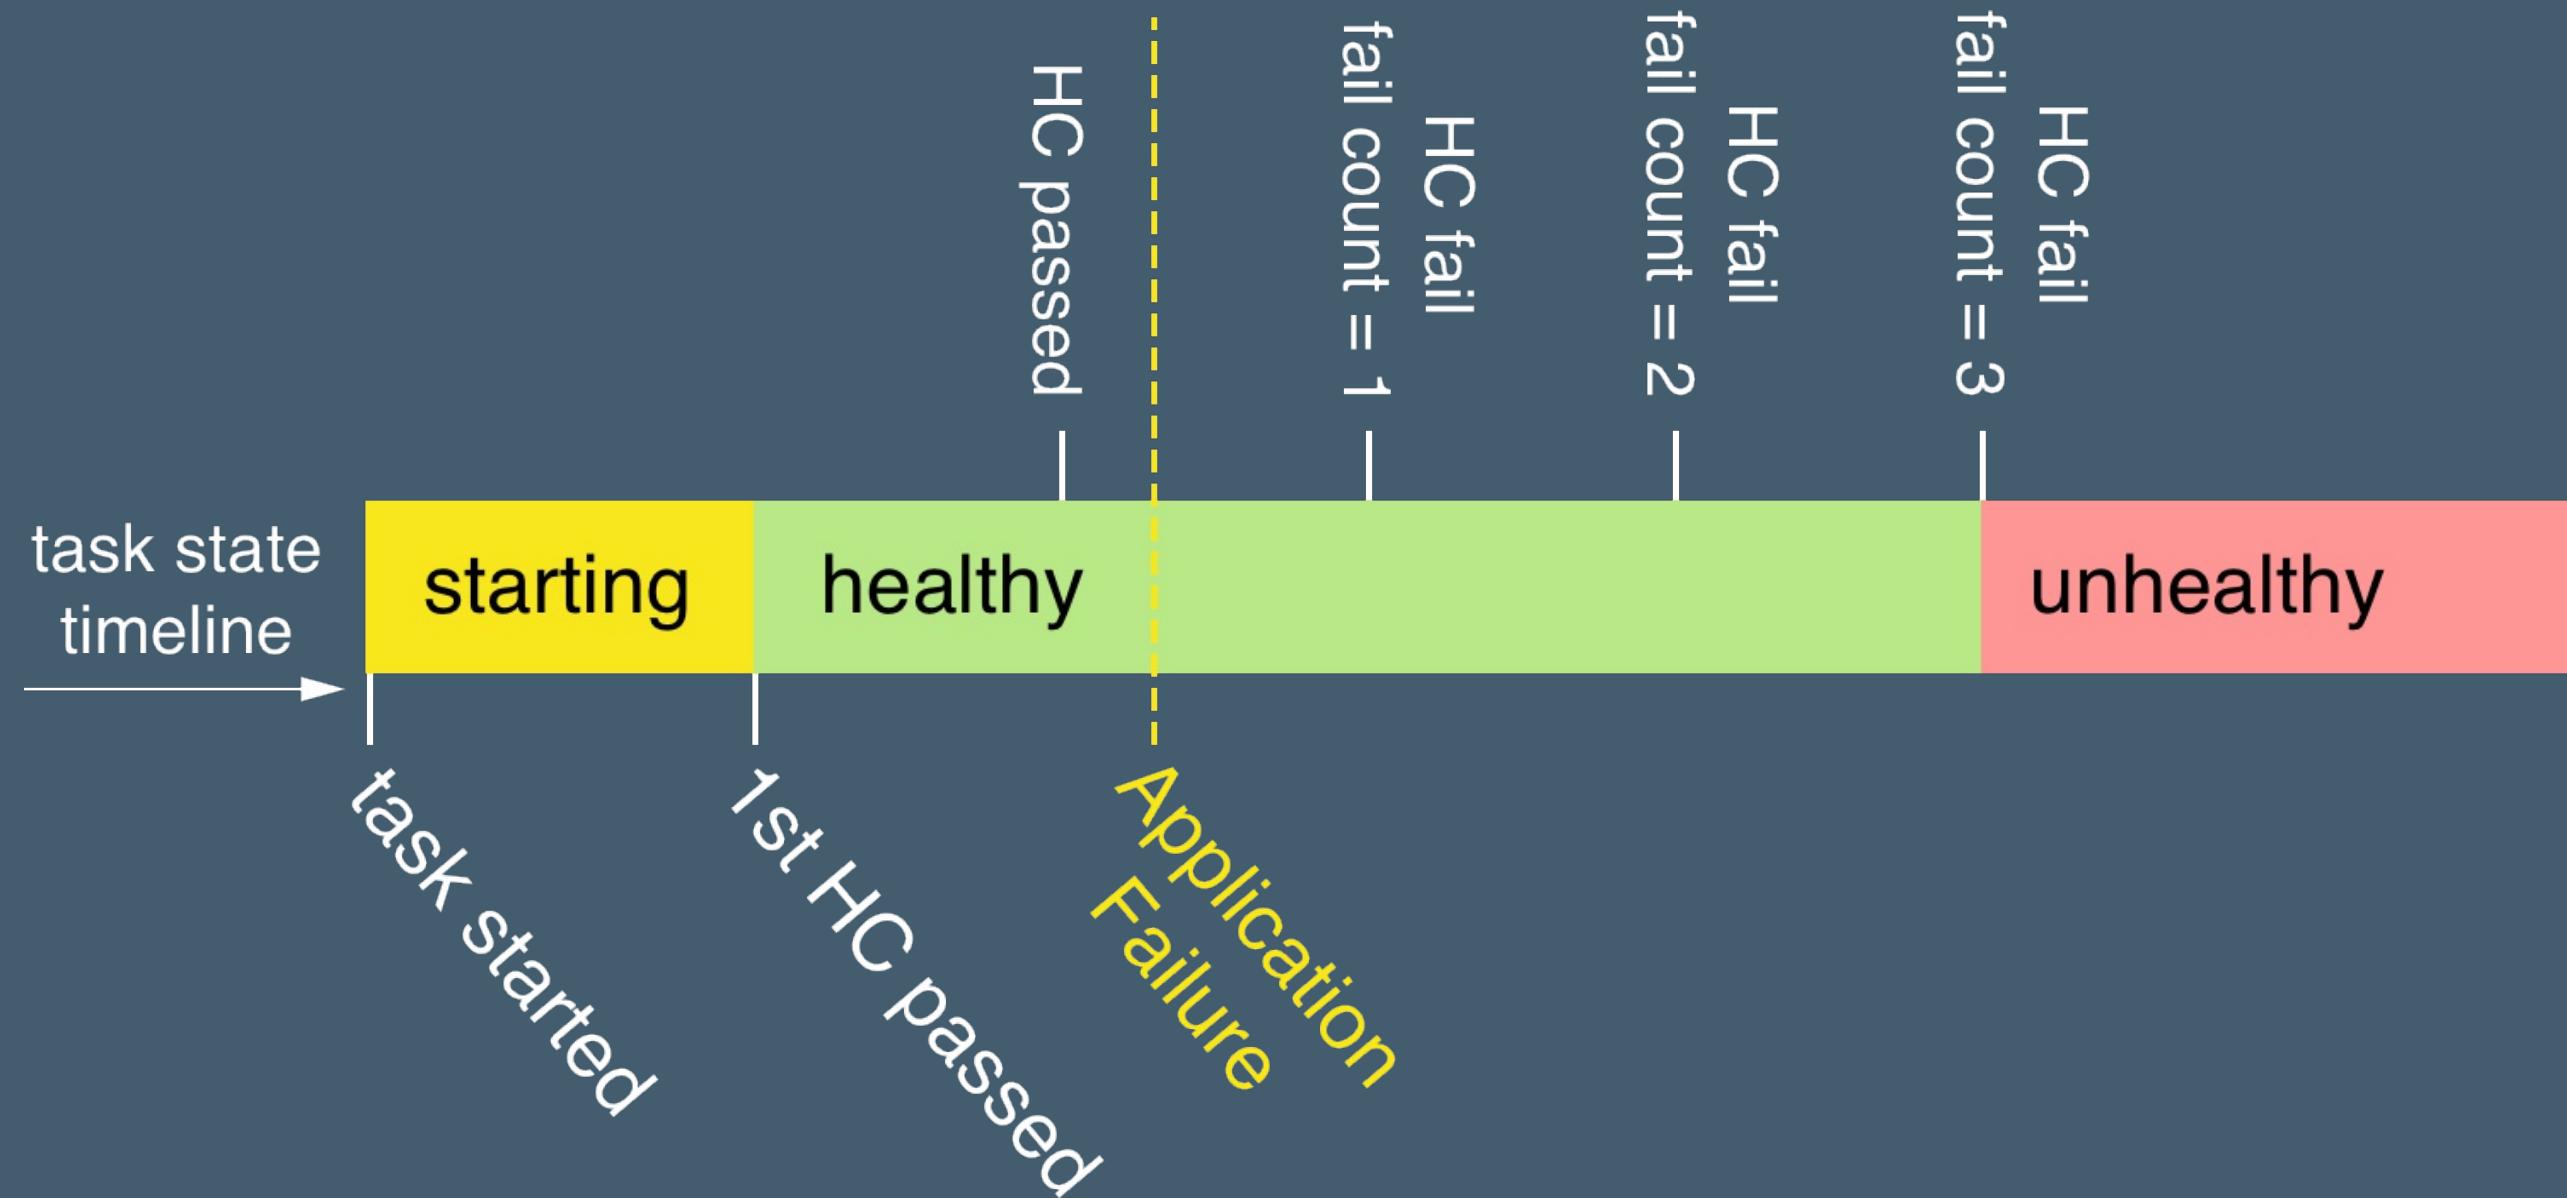
- Implement healthcheck logic in application code for integration into Swarm and Kubernetes' healthcheck protocols
- Configure automated healthchecks in Swarm and Kubernetes



HEALTH CHECK LIFECYCLE - HEALTHY



HEALTH CHECK LIFECYCLE - UNHEALTHY



APPLICATION HEALTH CHECKS

- Orchestration engine periodically runs healthcheck command in service container
- Too many failures -> container/pod killed and rescheduled
- Healthcheck logic supplied by developers
- healthcheck exit 0: healthy
- healthcheck exit 1: unhealthy



IMPLEMENT HEALTH ENDPOINT

/health [GET] endpoint

- Python/Flask

```
@app.route("/health", methods=['GET'])
def get_health():
    if healthy == True:
        return 'OK', 200
    else:
        return 'NOT OK', 500
```

- Node JS/Restify

```
server.get('/health', function create(req, res, next) {
  if(healthy){
    res.send(200, 'OK');
  } else {
    res.send(500, 'NOT OK');
  }
  return next();
});
```

Alternatives: <https://github.com/docker-library/healthcheck>



DOCKER SWARM: HOW TO CHECK HEALTH

- Dockerfile

```
HEALTHCHECK CMD curl --fail http://localhost:5000/health || exit 1
```

Command should return 0 for healthy, 1 for unhealthy

- Docker Compose File

```
healthcheck:  
  interval: 10s  
  timeout: 2s  
  retries: 3
```



KUBERNETES: HOW TO CHECK HEALTH

Liveness Probe:

```
kind: Pod
...
spec:
  containers:
    - name: demo
      image: ...
      livenessProbe:
        periodSeconds: 10
        timeoutSeconds: 2
        failureThreshold: 3
        initialDelaySeconds: 30
        successThreshold: 1
      exec:
        command:
          - cat
          - /tmp/healthy
...
...
```



Readiness Probe:

```
kind: Pod
...
spec:
  containers:
    - name: demo
      image: ...
      readinessProbe:
        periodSeconds: 60
        timeoutSeconds: 2
        ...
        httpGet:
          path: /healthz
          port: 8080
          httpHeaders:
            - name: X-Custom-Header
              value: Awesome
...
...
```

ee2.0-v2.1



EXERCISE: HEALTH CHECKS

Work through the Health Checks exercise in your Docker for Enterprise Developers Exercises book.



DISCUSSION

- Healthcheck endpoints must provide a measure of application health - but what other related logic might you want to implement around your healthcheck?
- Questions?



FURTHER READING

- Dockerfile reference: <http://dockr.ly/1PallBK>





DEFENSIVE PROGRAMMING



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Identify hazards posed by container orchestration that need to be handled effectively
- Write defensive logic that handles microservice outages
- Define and implement circuit breaking patterns in containers



DEFENSIVE PROGRAMMING

- SIGTERM versus SIGKILL
- Availability of external service not guaranteed
- Stale data vs. no data
- Fail fast
- Circuit breakers



SIGTERM VERSUS SIGKILL

- SIGTERM and SIGKILL issued by orchestration engine
- Stateless services are best
- Make failure recoverable:
 - Transactions
 - Compensating actions
 - Idempotency



EXTERNAL SERVICE NOT AVAILABLE

- Distributed services live and die independently
- Docker DNS propagation **eventually consistent**
- Retry
- Increase wait time between calls
- Log warning
- Also see: Circuit Breaker



STALE DATA VERSUS NO DATA

If external service is not available:

- Fail if critical data
- Serve stale data if freshness is less important
- Serve no data if freshness is important



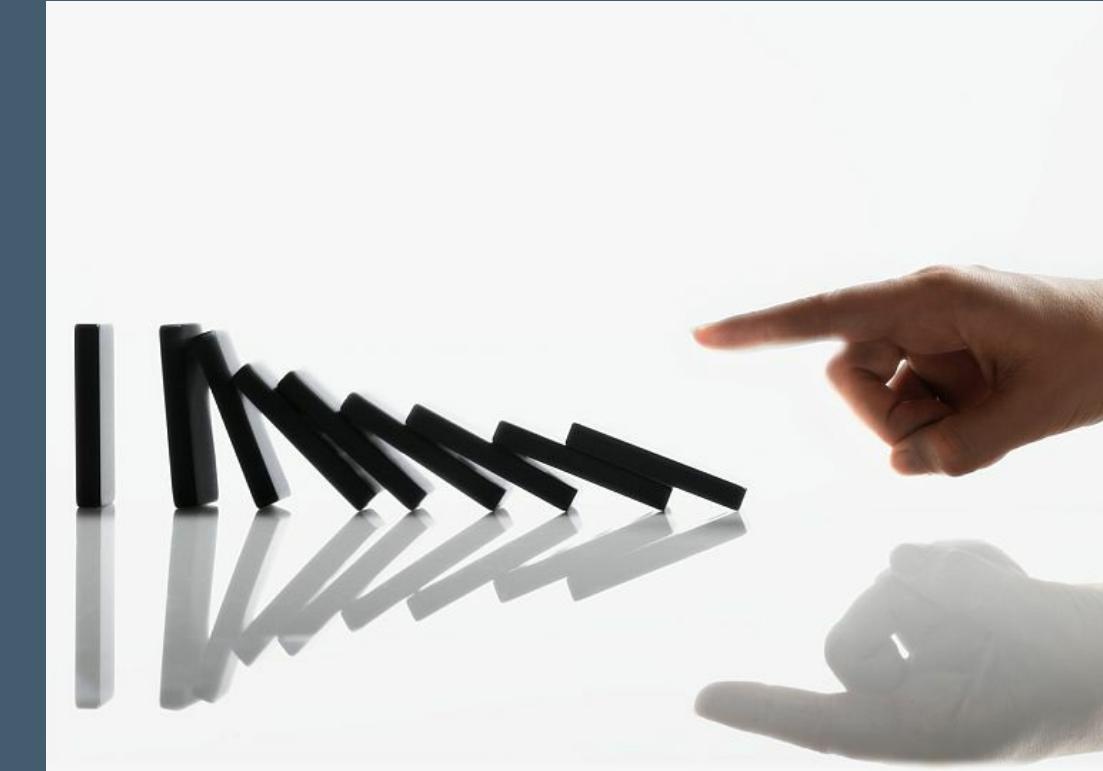
FAIL FAST

- Always check input data first
- Log clear failure reason
- Yet, do not disclose too much
- Take advantage of service rescheduling

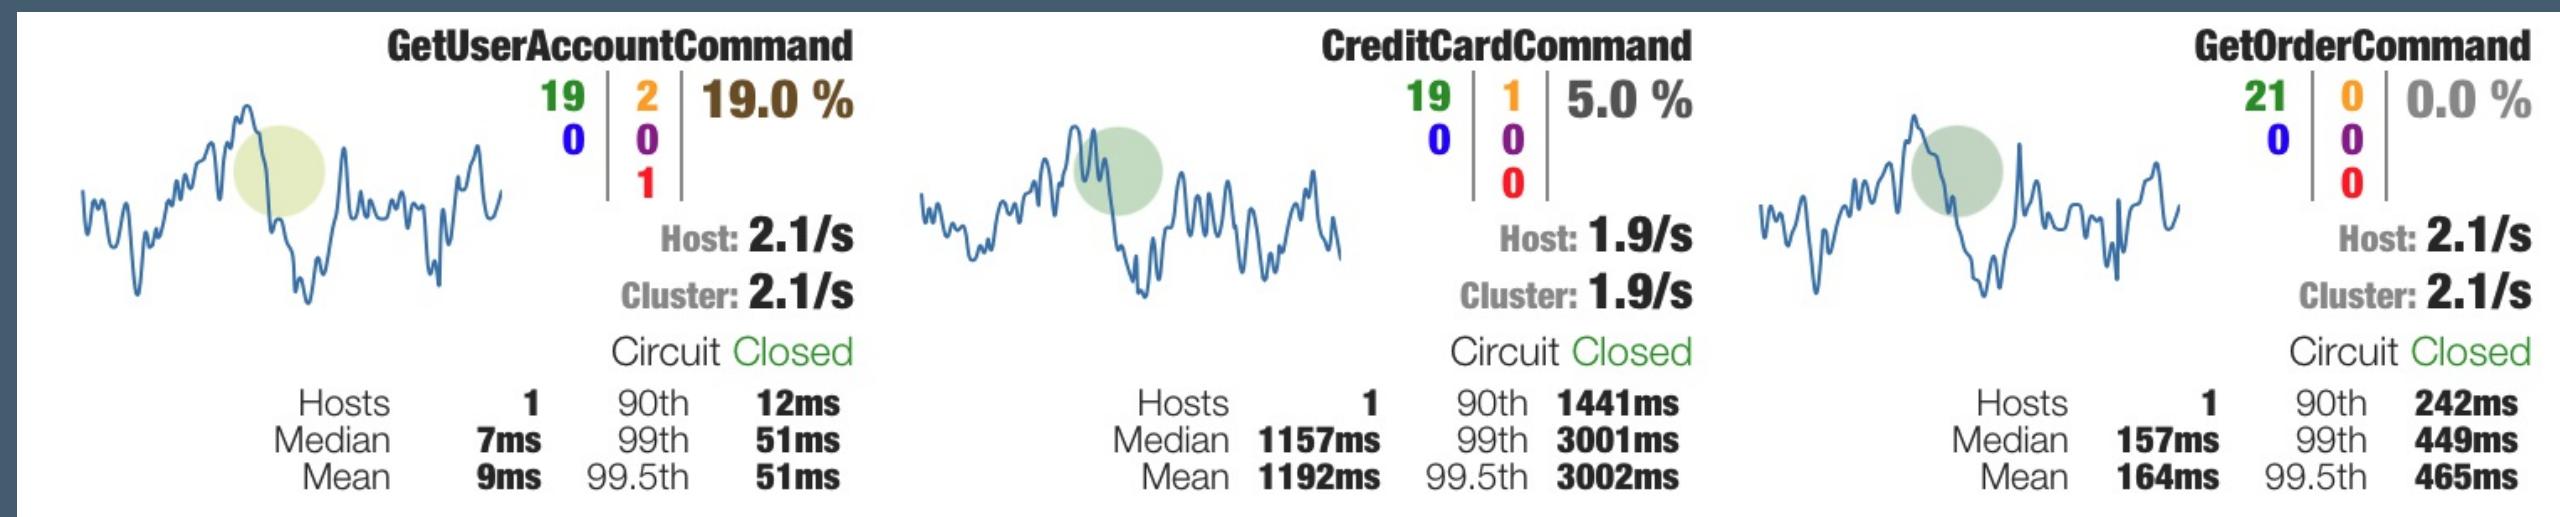


CIRCUIT BREAKERS

- Service separation facilitates damage control
- Avoid Domino Effect
- Give failed service/system time to recover
- Gracefully degrade experience
- Serve stale (or no) data instead



Netflix Hystrix





EXERCISE: DEFENSIVE PROGRAMMING

Work through the 'Defensive Programming' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What are some container-specific situations that should be defended against?
- Questions?



FURTHER READING

- Controlling startup order in Compose: <http://dockr.ly/2vTPjYE>
- Trapping signals in Docker containers: <http://bit.ly/2wZVKZK>
- Netflix Hystrix: <https://github.com/Netflix/Hystrix>
- Managing Microservices with the Istio Service Mesh: <http://bit.ly/2nJjpYH>





LOGGING & ERROR HANDLING



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Instrument a containerized application with logging output
- Configure the Docker daemon's logging driver
- Name principles of error handling fit for orchestrated applications



LOGGING

Why does Logging matter?

- Highly distributed environment
- Many moving parts
- High probability of failure

Why is central Logging crucial?

- Correlate events
- Slice and dice
- No need to access hosts
- No need to backup logs
- Improved security



LOGGING

Where and **What** to log?

- Send logging info to StdOut and StdErr
- Add severity (debug, info, warn, err, fatal)
- Add service relevant unique identifier
- Do **NOT** add sensitive or personal information!



LOGGING

Python

```
from flask import Flask
app = Flask(__name__)
app.logger.setLevel(logging.DEBUG)
...
@app.route('/warning')
def warning():
    app.logger.warning('This is a warning message')
    return "warning"
```

Java

```
public class main{
    static Logger logger = LogManager.getLogger("main");

    public static void main(String[] args) {
        logger.debug("Hello this is a debug message");
        logger.info("Hello this is an info message");
    }
}
```



DOCKER EE LOGGING

- Logs come from three sources:
 - container runtime: system-level
 - kubelet: system-level
 - containers: STDOUT & STDERR
- Globally configured by **/etc/docker/daemon.json**:

```
{  
  "log-driver": "journald",  
  "log-level": "debug",  
  "log-opt": {  
    "tag": "{{.ImageName}}/{{.Name}}/{{.ID}}"  
  }  
}
```

- Override for containers with **--log-opt**



CENTRALIZED LOGGING - ELK STACK



```
version: "3"
services:
  elasticsearch:
    image: elasticsearch:2

  logstash:
    image: logstash
    command: |
      -e '
        input { ... }
        filter { ... }
        output { ... }'
    ports:
      - "12201:12201/udp"

  kibana:
    image: kibana:4
    ports:
      - "5601:5601"
    environment:
      ELASTICSEARCH_URL: http://elasticsearch:9200
```



ERROR HANDLING

- Fail fast!
- Provide **good** quality error messages
 - reason for failure
 - app / service specific unique ID
 - stack trace
- Use defensive coding style
- Sometimes stale results are better than no results
- Implement health endpoint





EXERCISE: LOGGING & ERROR HANDLING

Work through the exercises

- Logging
- Error Handling

in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What additional logging information should your application provide after containerization?
- Questions?



FURTHER READING

- Configure logging drivers: <http://dockr.ly/2gSFMdS>
- Kubernetes Logging Architecture: <http://bit.ly/2veExYy>





BUILDER



ee2.0-v2.1 © 2018 Docker, Inc.

LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- List several Dockerfile best practices concisely
- Author a multi-stage Dockerfile for an application to minimize the final image size



IMAGE SIZE MATTERS

Image	Size (MB)
java:8-jdk	643
java:8-jdk-alpine	145
java:8-jre	311
java:8-jre-alpine	108
node:7.9	665
node:7.9-alpine	59
python:2.7	676
python:2.7-alpine	72
golang:1.7	676
golang:1.7-alpine	241
microsoft/windowsservercore	~9,000
microsoft/nanoserver	~450
busybox	1.1

Smaller images =

- smaller attack surface
- faster download & build times
- often faster container startup times



OPTIMIZE THE DOCKERFILE

Rely on image layer cache

Seldom changing --> beginning of Dockerfile

Frequently changing --> end of Dockerfile

```
FROM node:8-jdk-alpine
RUN yum install <somelibrary>
WORKDIR /app
COPY package.json .
RUN npm install
COPY ..
EXPOSE 3000
CMD npm start
```



MULTI STAGE BUILD

Multiple FROM statements per Dockerfile

Only last FROM results in an image

All layers are cached

```
FROM ubuntu AS build-env
RUN apt-get install make
ADD . /src
RUN cd /src && make

FROM busybox
COPY --from=build-env /src/build/app /usr/local/bin/app
EXPOSE 80
ENTRYPOINT /usr/local/bin/app
```





EXERCISE: BUILDER

Work through the 'Builder' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What are the tradeoffs between more versus fewer image layers?
- Questions?



FURTHER READING

- Best practices for writing Dockerfiles: <http://dockr.ly/22WijO>
- Use multistage builds: <http://dockr.ly/2ewcUY3>
- Containers and layers: <http://dockr.ly/2vOwbap>





DOCKER SWARM & KUBERNETES



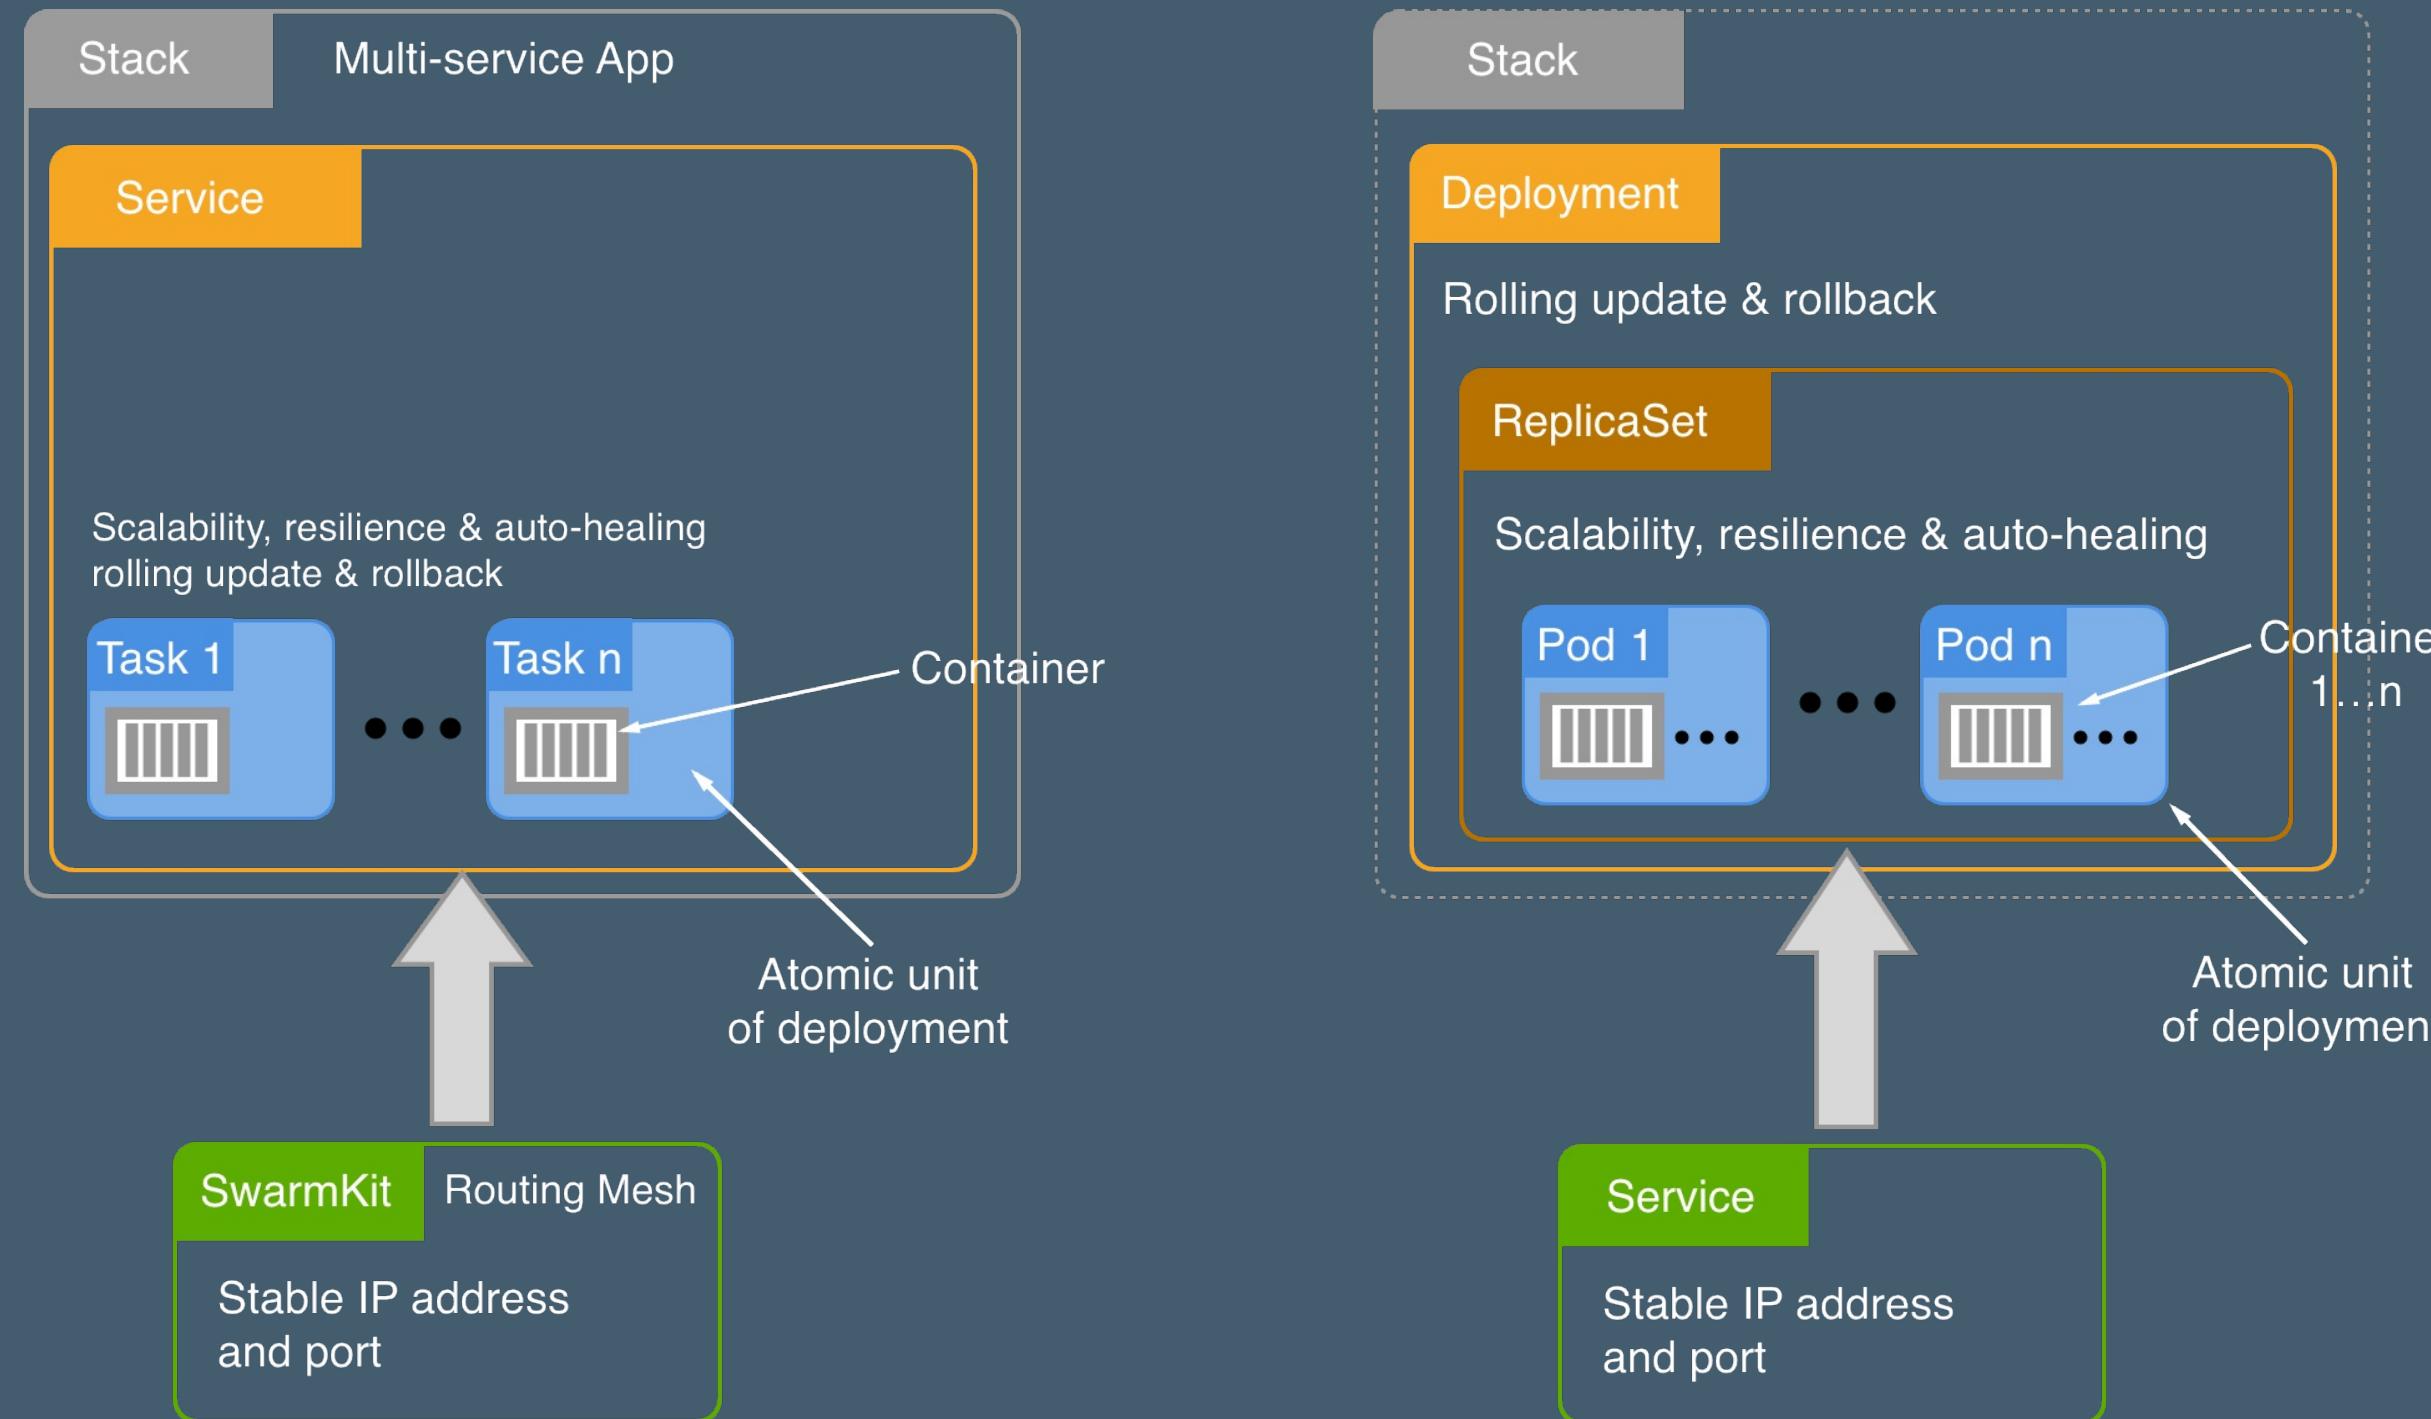
LEARNING OBJECTIVES

By the end of this module, learners will be able to:

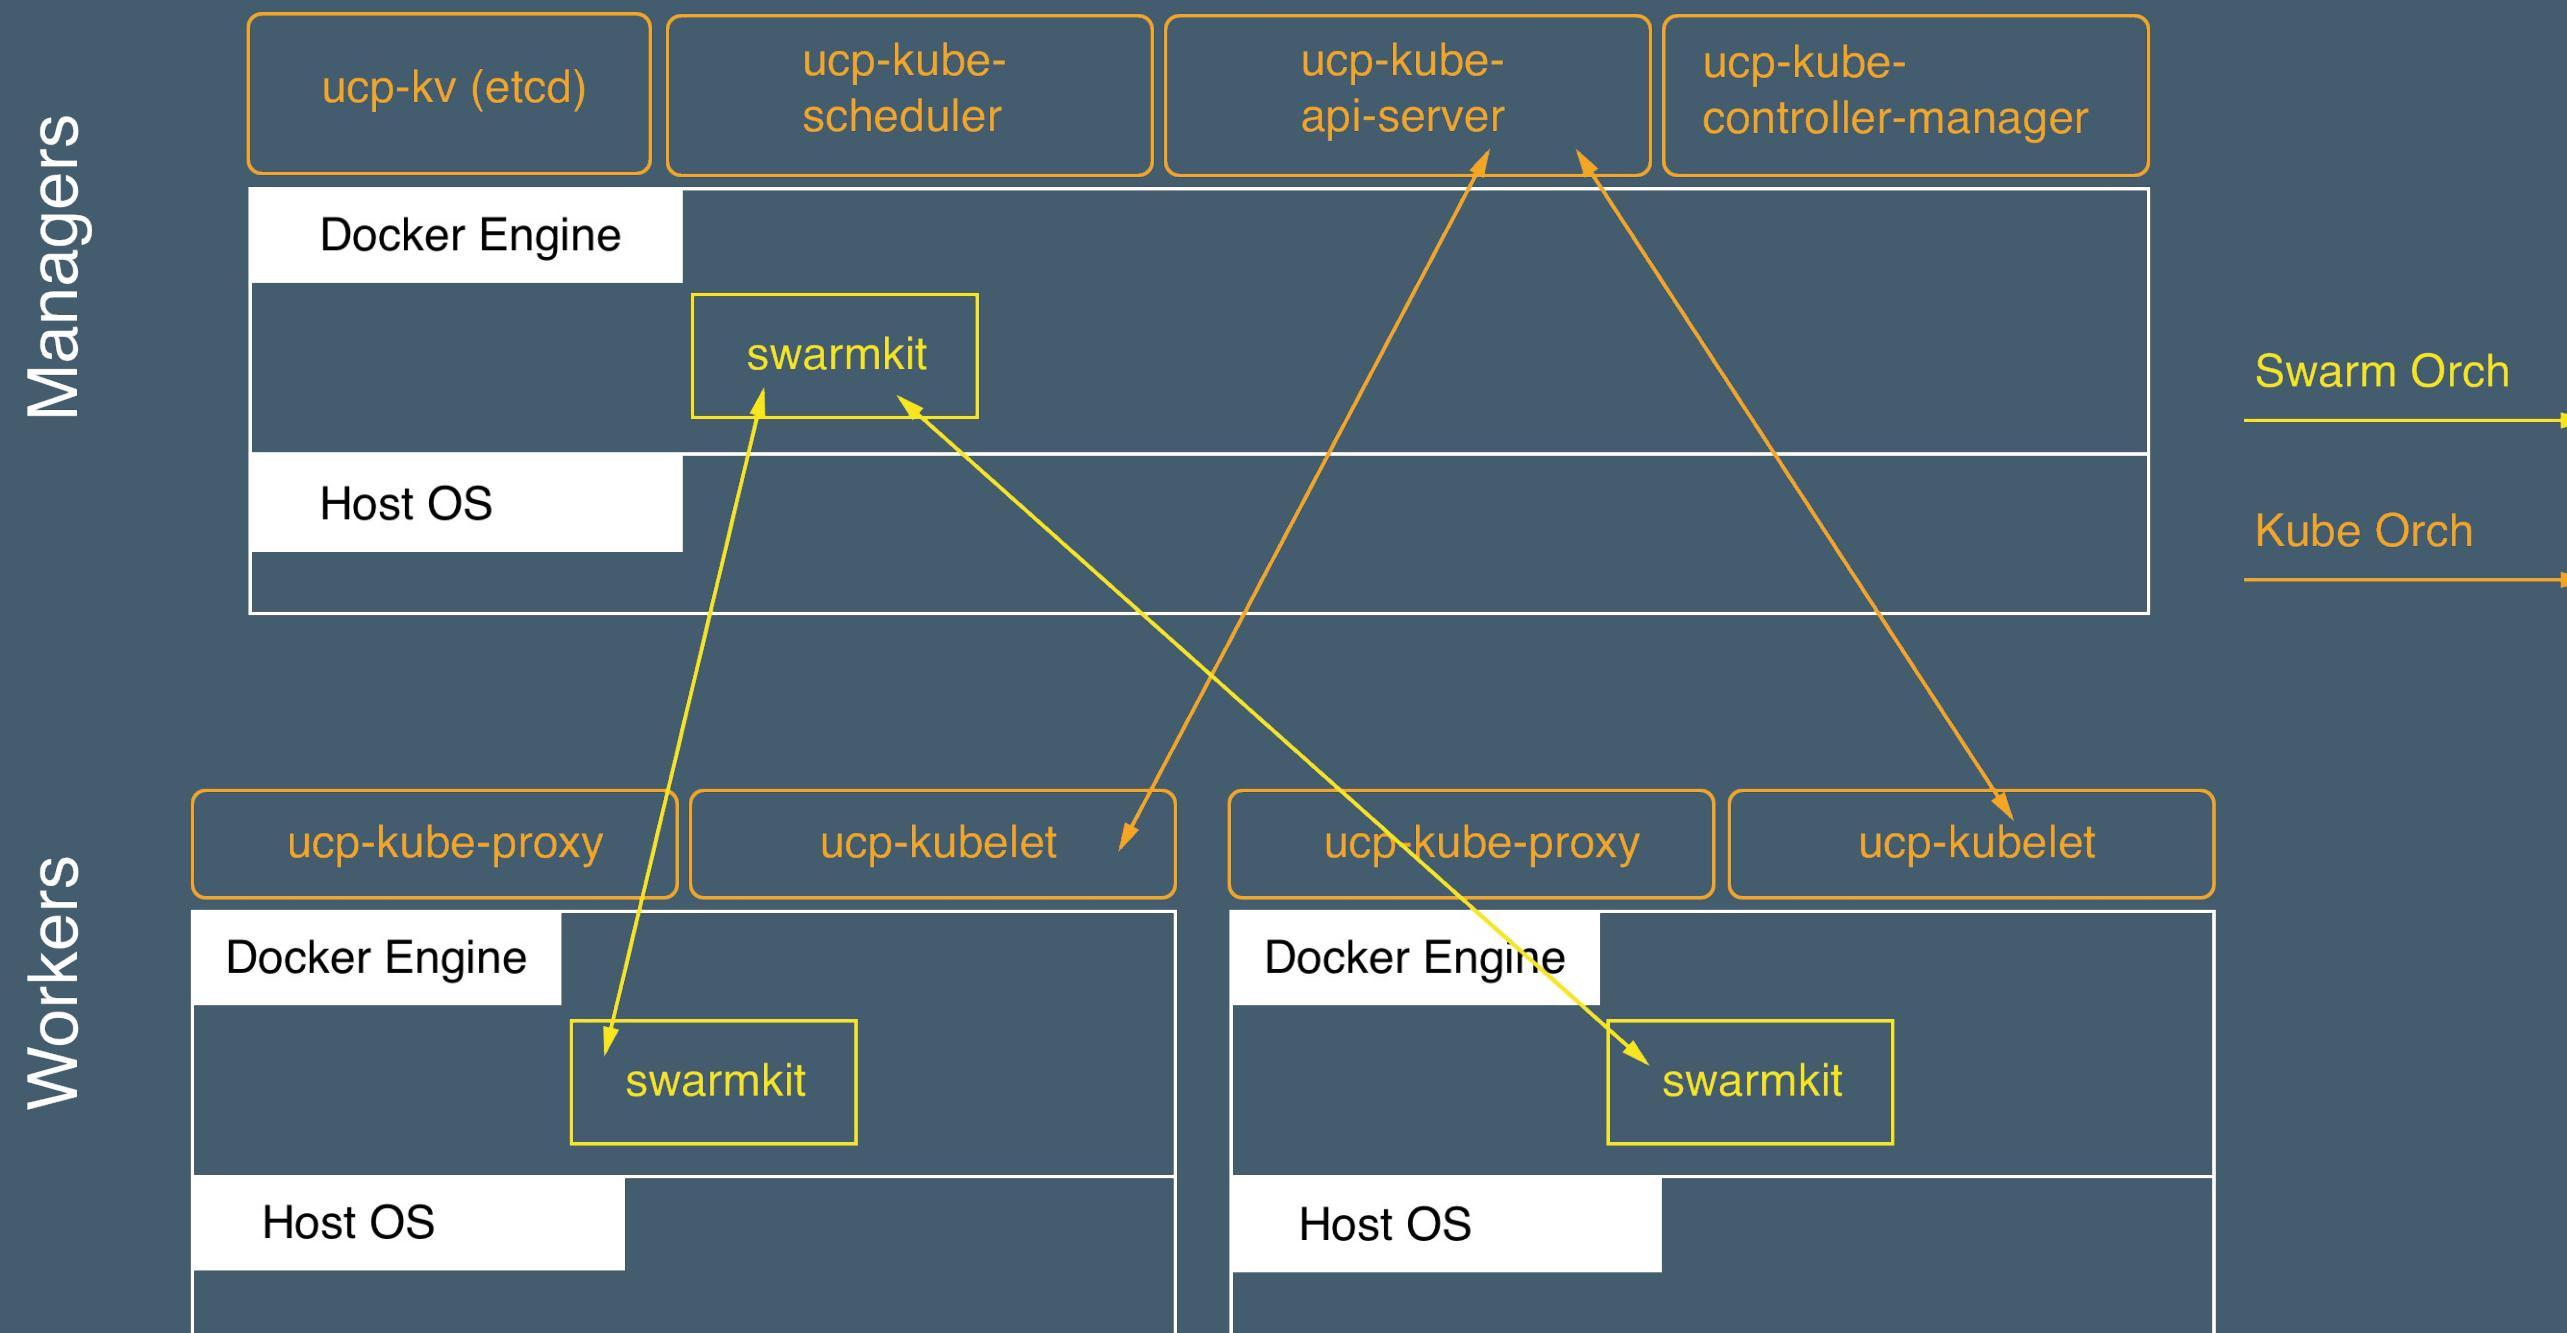
- Draw correct analogies between fundamental Swarm and Kubernetes objects
- Understand how UCP manages both orchestrators
- Implement layer 4 routing to applications in both orchestrators



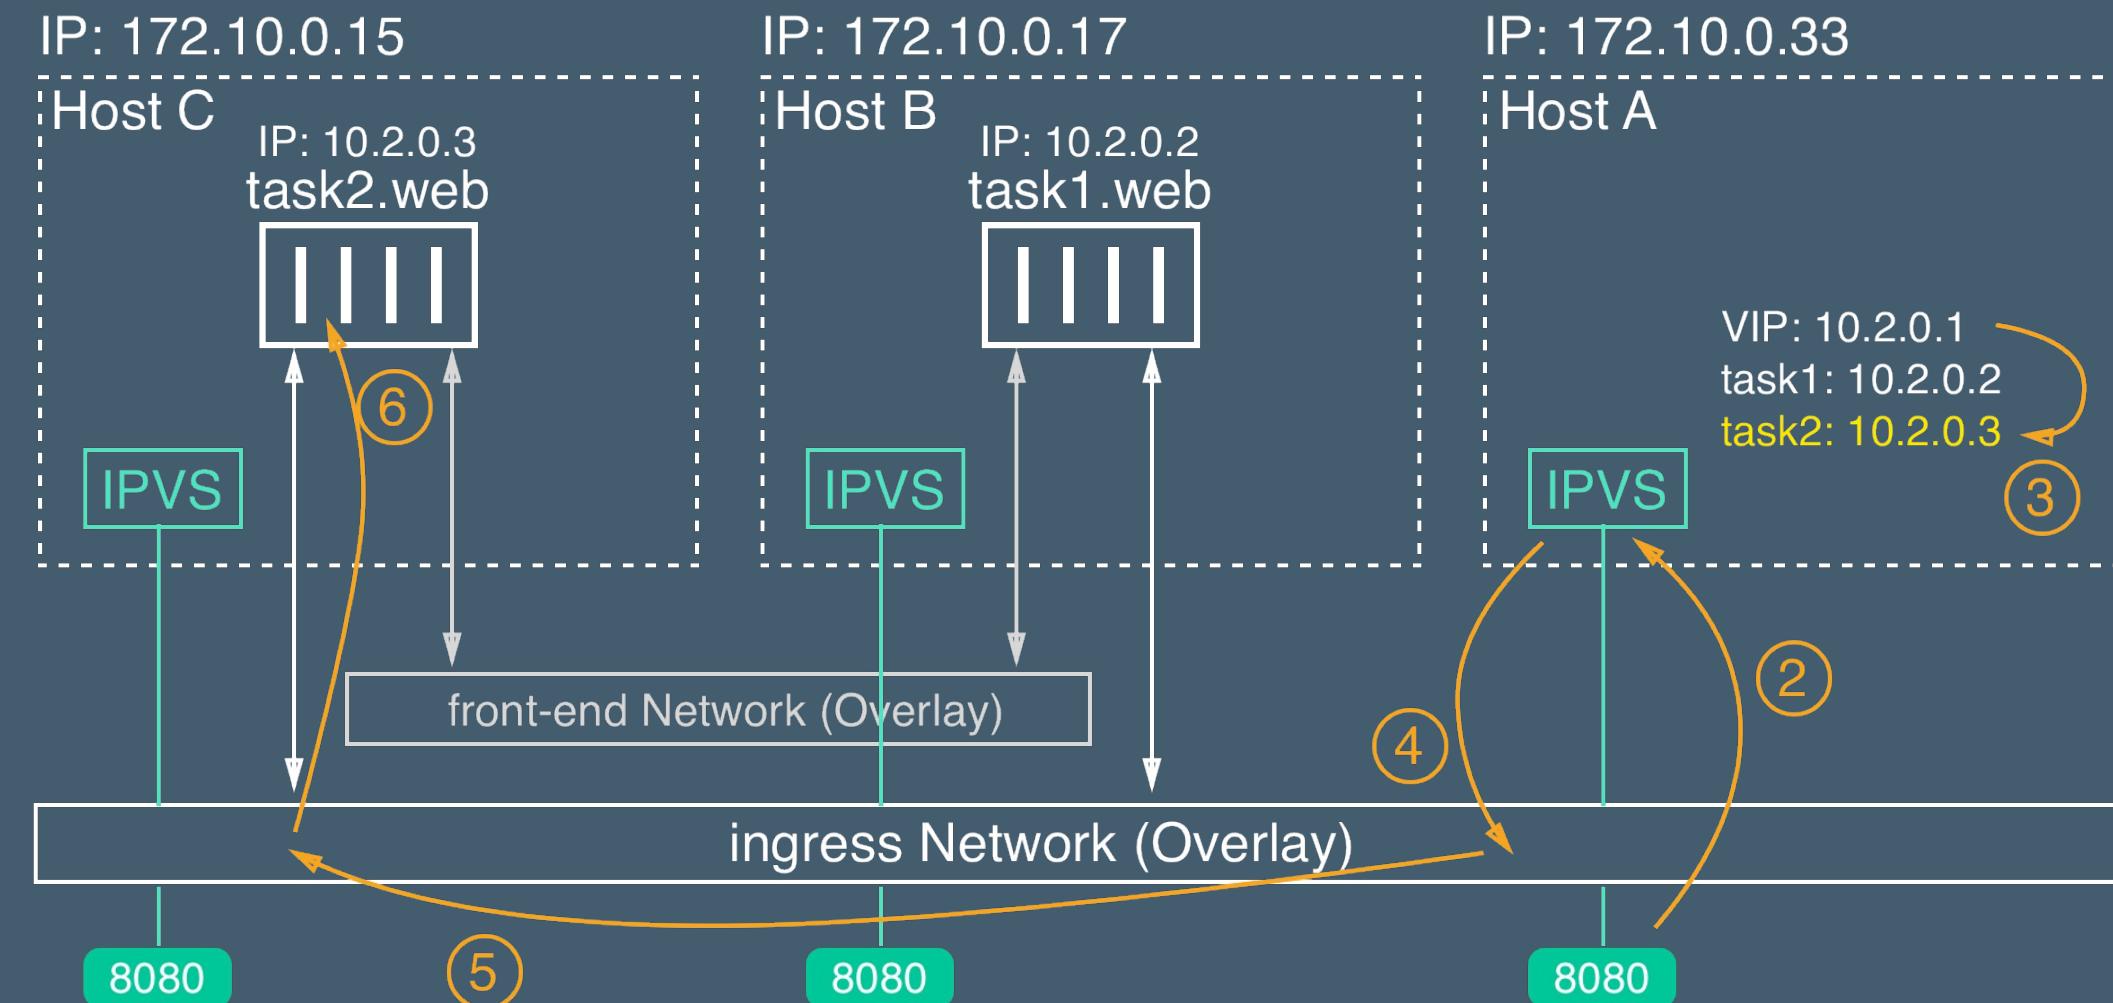
SWARM VERSUS KUBERNETES



ORCHESTRATORS IN UCP



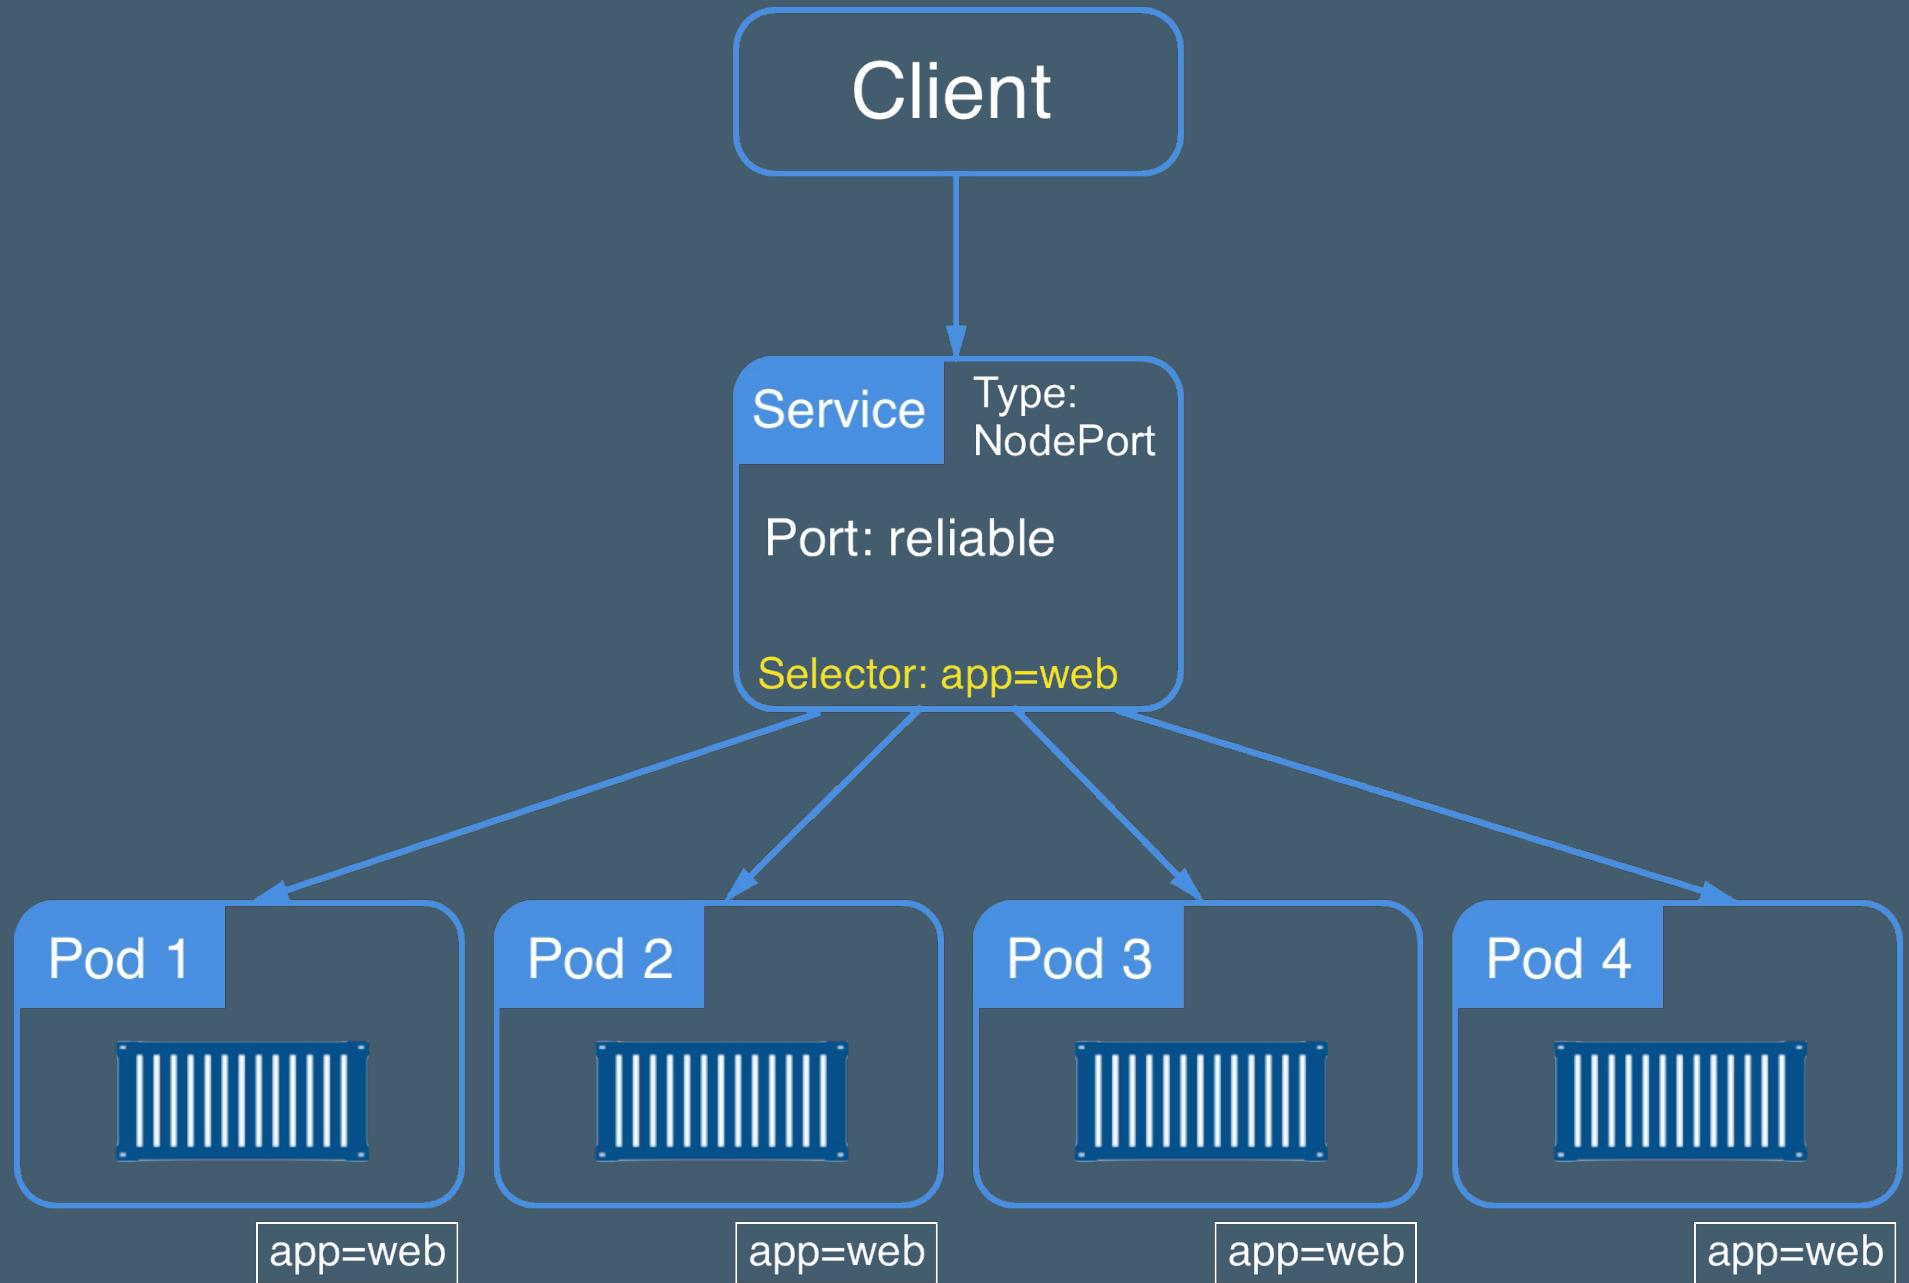
SWARM ROUTING MESH



```
docker service create  
--name web  
--replicas 2 \  
--network front-end \  
-p 8080:80 nginx:alpine
```



L4 ROUTING - KUBERNETES: SERVICE OBJECT





EXERCISE: DOCKER SWARM

Work through the exercises:

- Creating a Docker Swarm
- Routing Mesh

in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- What sort of applications would you prefer to design for Swarm? Kubernetes?
- Questions?



FURTHER READING

- Swarm mode overview: <http://dockr.ly/2jh11Vd>
- Swarm mode key concepts: <http://dockr.ly/2jgMT0f>
- Getting started with swarm mode: <http://dockr.ly/2vYFNU4>
- High availability in Docker swarm: <http://dockr.ly/2gVH6MS>





SECRETS



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Use secrets in Swarm and Kubernetes
- Manage and mock secrets across environments
- Identify the encryption strategies used to protect secrets throughout their lifecycle



DISTRIBUTING SECRETS

- Apps need sensitive data like access tokens, ssh keys etc
- How to securely store and distribute secrets across a cluster?

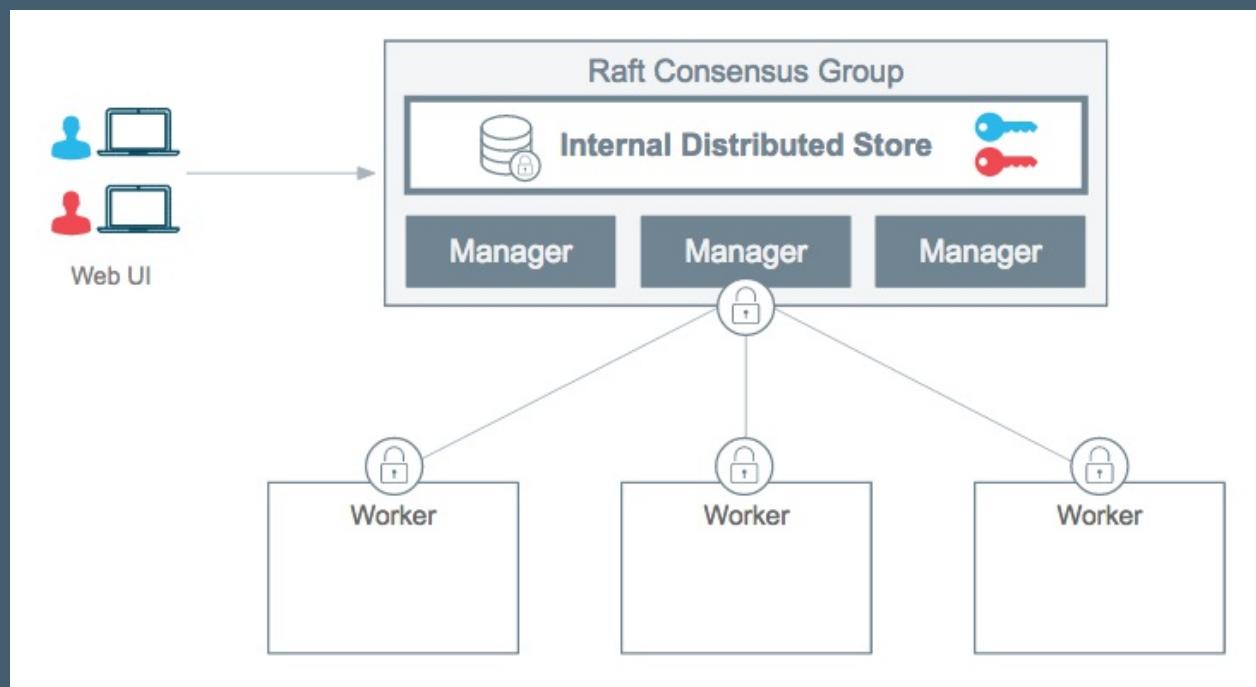


SECRET REQUIREMENTS

- Encrypted at rest
- Encrypted in transit
- Accessible only to intended containers



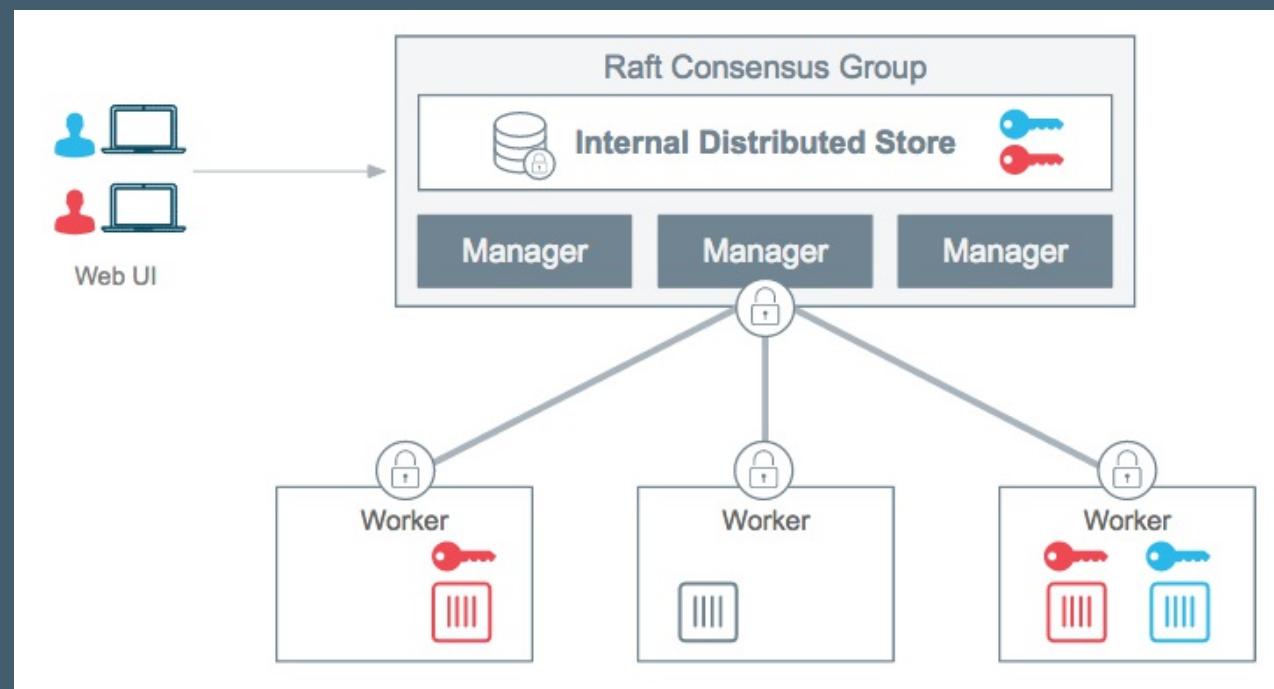
SECRETS WORKFLOW 1: CREATION & STORAGE



- Swarm:
 - Stored in state db
 - Automatically encrypted at rest
- Kubernetes:
 - Stored in etcd
 - Must configure encryption:
<http://bit.ly/2GCoPwE>



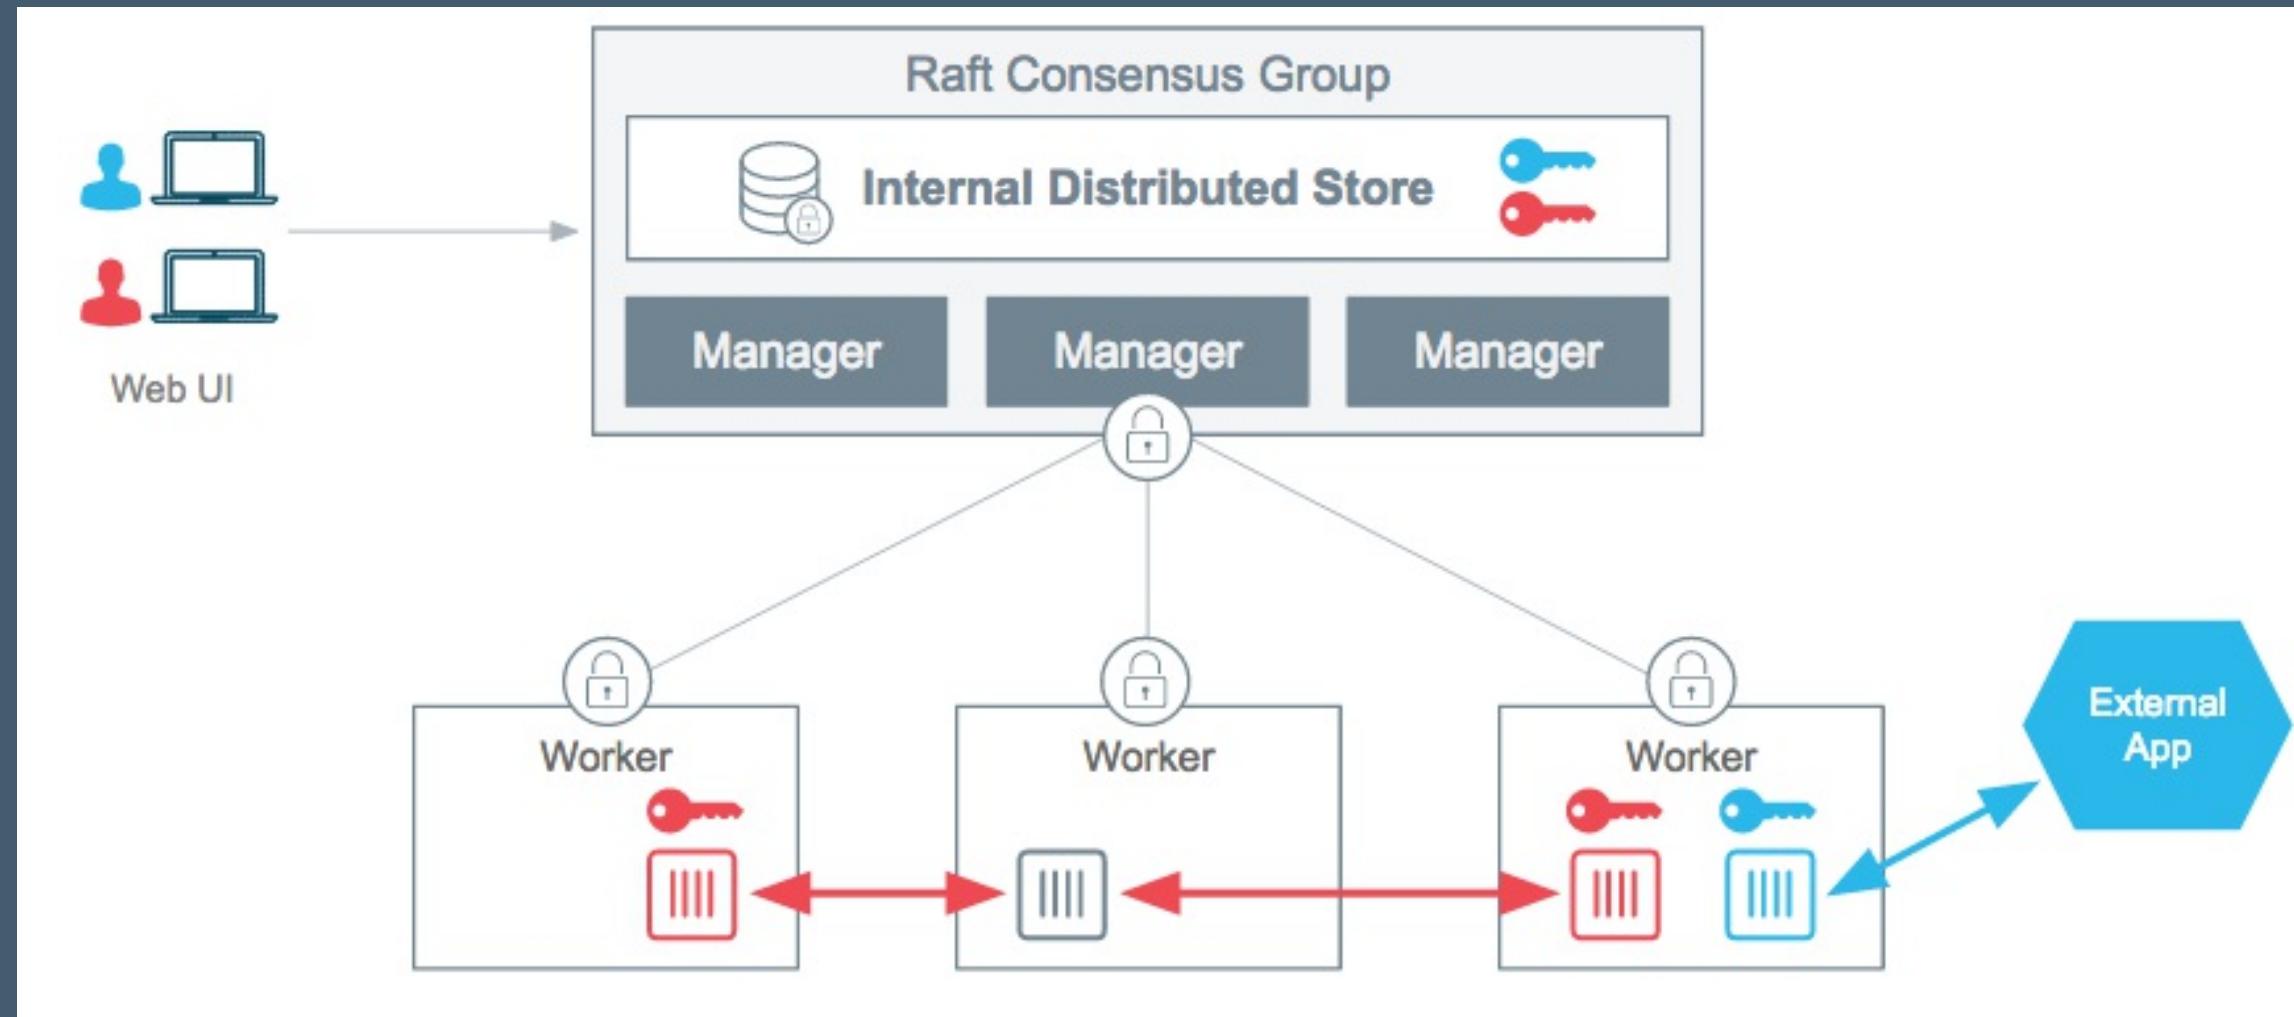
SECRETS WORKFLOW 2: DISTRIBUTION



- Swarm:
 - Default mutual TLS
 - Associated with services
- Kubernetes:
 - TLS optional
 - Associated with pods or deployments
- Both:
 - Least Privilege: secrets available only exactly where needed
 - Deleted along with container



SECRETS WORKFLOW 3: SECRET USAGE



SECRETS IN DEVELOPMENT

Swarm: Service

```
docker service create --name foo \  
  --secret source=DB_PASSWORD_V1,target=DB_PASSWORD \  
  --env DB_PASSWORD_FILE=/run/secrets/DB_PASSWORD \  
  <IMAGE_NAME>
```

Development: Container

```
docker container run --name foo \  
  -v ./dev-secrets/DEV_DB_PASSWORD:/app/my-secrets/DB_PASSWORD \  
  -e DB_PASSWORD_FILE=/app/my-secrets/DB_PASSWORD \  
  ... \  
  <IMAGE_NAME>
```

Code (Node JS)

```
var fs = require('fs');  
var db_password = fs.readFileSync(process.env.DB_PASSWORD_FILE, 'utf8');  
...
```





SECRETS

Work through the 'Secrets' exercise in the Docker for Enterprise Developers Exercises book.



DISCUSSION

- Secrets are one way to inject config into a container. What are some others, and when should each be used?
- Questions?



FURTHER READING

- Manage sensitive data with Docker secrets: <http://dockr.ly/2vUNbuH>
- Docker secret reference: <http://dockr.ly/2iSsNJC>
- Introducing Docker secrets management: <http://dockr.ly/2k7zwzE>
- Securing the AtSea app with Docker secrets: <http://dockr.ly/2wx5MyV>
- Manage secrets: <https://dockr.ly/2HmznE3>
- Kubernetes Secrets: <http://bit.ly/2C6hMZf>





CONFIGURATION MANAGEMENT



LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- List a few places in a containerized enterprise application to put configuration information
- Categorize the configuration data type usually used in enterprise applications
- Explain and justify where to put sensitive configuration data

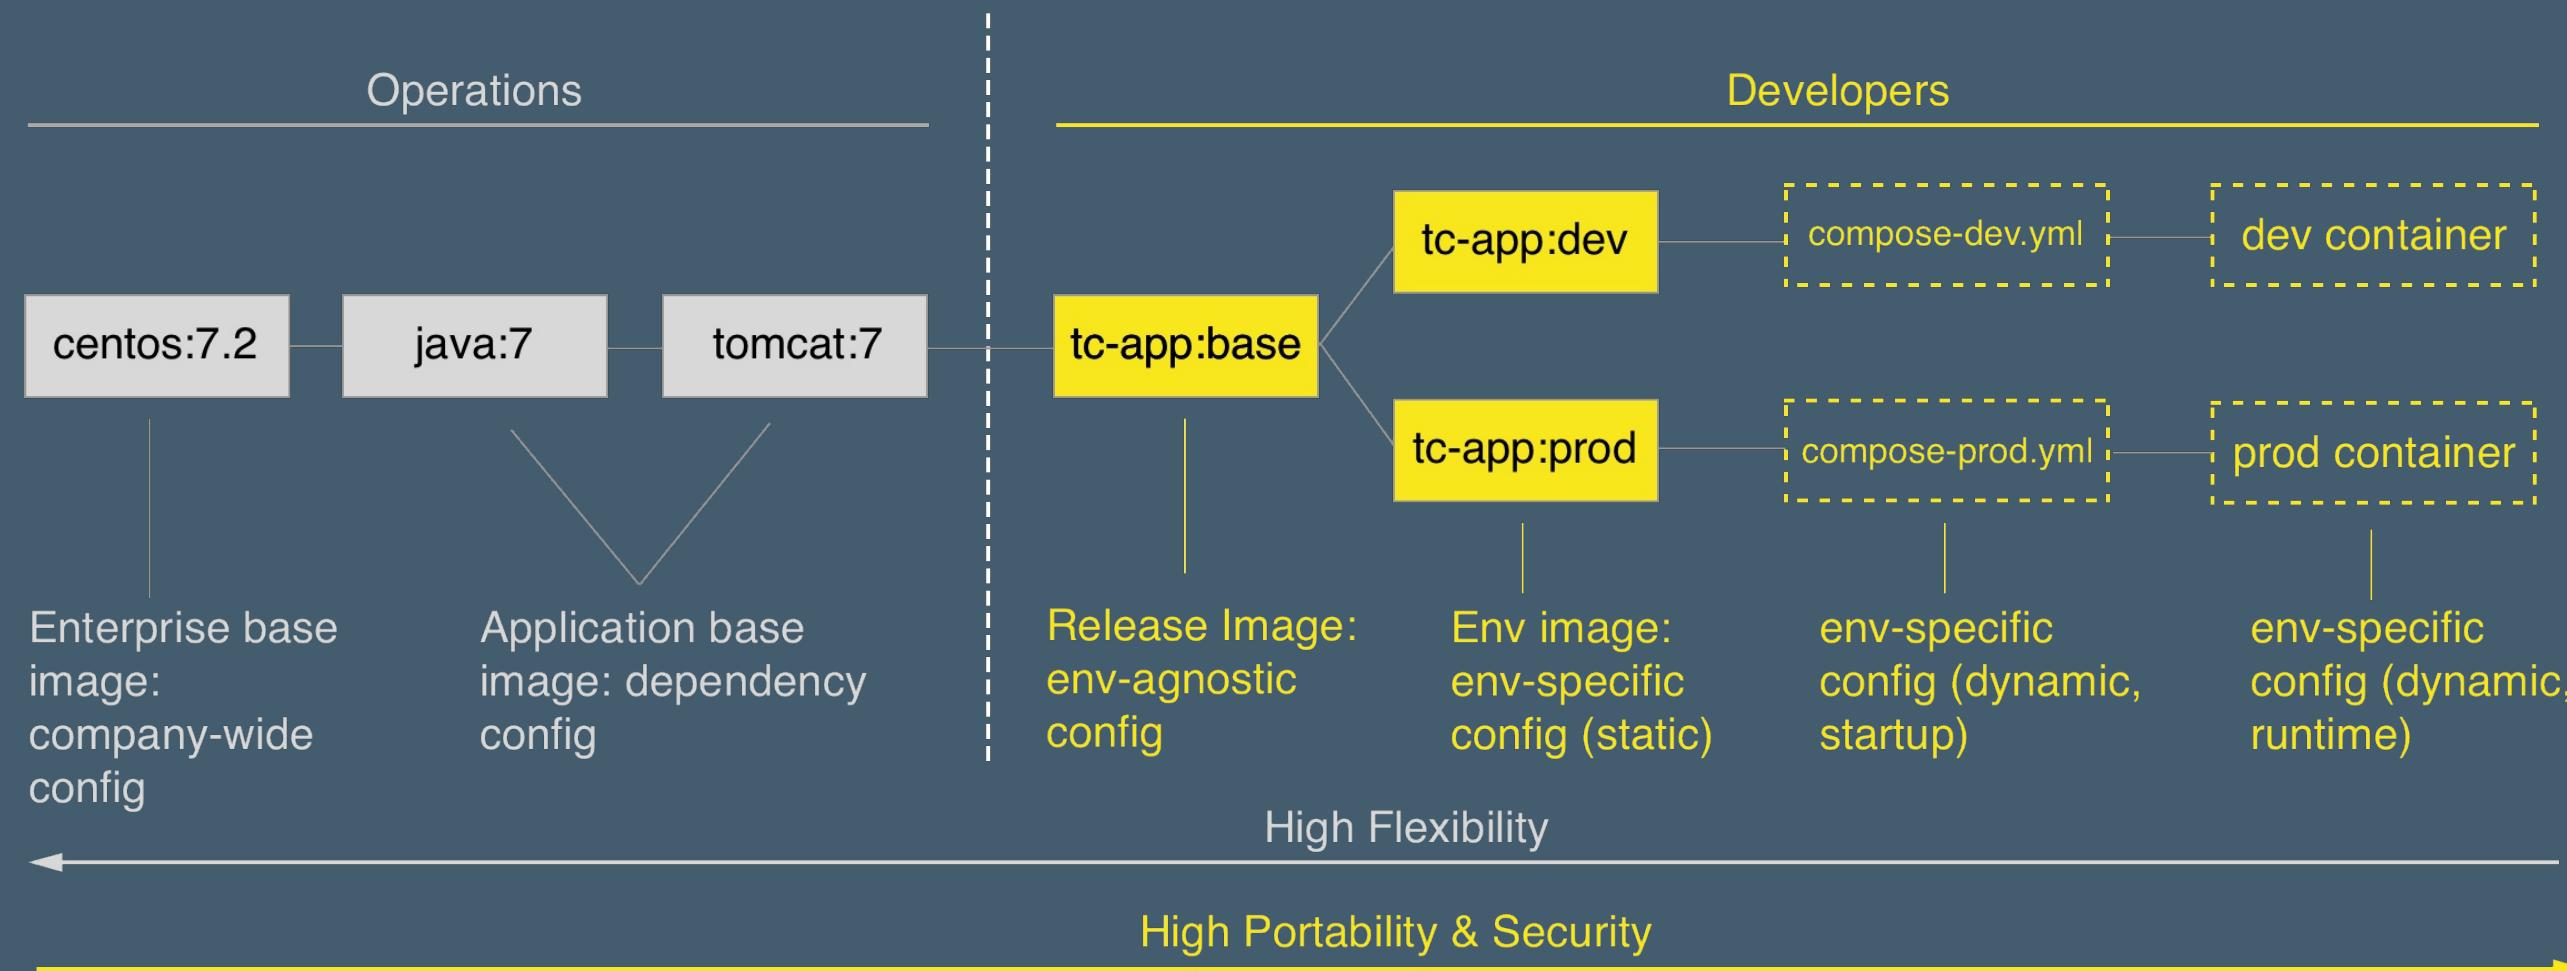


WHERE SHOULD MY CONFIG LIVE?

- Base image
- Derived image
- Stack or Object YAML files
- parameter file in image
- parameter file mounted externally
- entrypoint script
- external service



CONFIGURATION BUCKETS



ENTERPRISE BASE IMAGE

- Provided by Ops (CaaS)
- Start with Official Image
- Add company wide Tools & Utilities
- Push to DTR
- Scan image --> zero vulnerabilities
- Mind Image Size!

```
FROM centos:7
RUN yum -y --noplugins install bzip2 tar sudo curl net-tools
```



APPLICATION BASE IMAGE

- Define as **Dockerfile**
- Start with Enterprise Base Image
- Add App Type specific Tools & Utilities
- Mind Image Size!

```
FROM <DTR_FQDN>/<ORG_Name>/enterprise-base:1.5
COPY files/dynatrace-agent-6.1.0.7880-unix.jar /opt/dynatrace/
```

...



RELEASE IMAGE

- Define as **Dockerfile**
- Start with Application Base Image
- Add Application Artifacts
- Add environment-agnostic config
- Mind Image Size!

```
FROM <DTR_FQDN>/<ORG_Name>/app-base:2.2
COPY files/MY_APP_1.3.1-M24_1.war /opt/jboss/standalone/deployments/
...

```



ENVIRONMENT IMAGE

- Define as **Dockerfile**
- Start with Release Image
- Add environment-specific config
- Mind Image Size!
- Could alternatively mount config at runtime



DYNAMIC CONFIG: STARTUP

- Config defined at startup
- provided in stack or object YAML files
 - environment variables
 - Docker EE-managed secrets
 - Volume mounts
- or defined in **entrypoint.sh**
 - Hashicorp Vault-managed secret
 - ```
$ curl -H "X-Vault-Token: f3b09679-3001-009d-2b80-9c306ab81aa6" -X GET \
https://vlt.example.com:8200/v1/secret/db
```



# DYNAMIC CONFIG: RUNTIME

- Config to be adjusted in-flight
- example: Consul KV service discovery

```
$ consul-template -consul consul.example.com:6124 \
-template "/tmp/nginx.ctmpl:/var/nginx/nginx.conf:service nginx restart"
```





## EXERCISE: CONFIGURATION MANAGEMENT

Work through the 'Configuration Management' exercise in the Docker for Enterprise Developers Exercises book.



# DISCUSSION

- What are some pitfalls of configuration mismanagement?
- Questions?



## FURTHER READING

- Store configuration data using Docker configs: <http://dockr.ly/2wTjFrv>
- Kubernetes Configuration Management with Containers: <http://bit.ly/2E8Atl9>





# DEVELOPMENT PIPELINE OVERVIEW



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Explain what container-as-a-service means in the context of a typical enterprise using DDC
- Sketch a diagram containing all essential elements of a Docker development pipeline
- List a handful of important differences between a traditional and a containerized CI/CD pipeline



# TRADITIONAL CI/CD

- Push code to a repo
- CI/CD manager pulls code through tests in multiple environments
- Code (automatically?) deployed to production when passing

But what does this look like with containers and images?

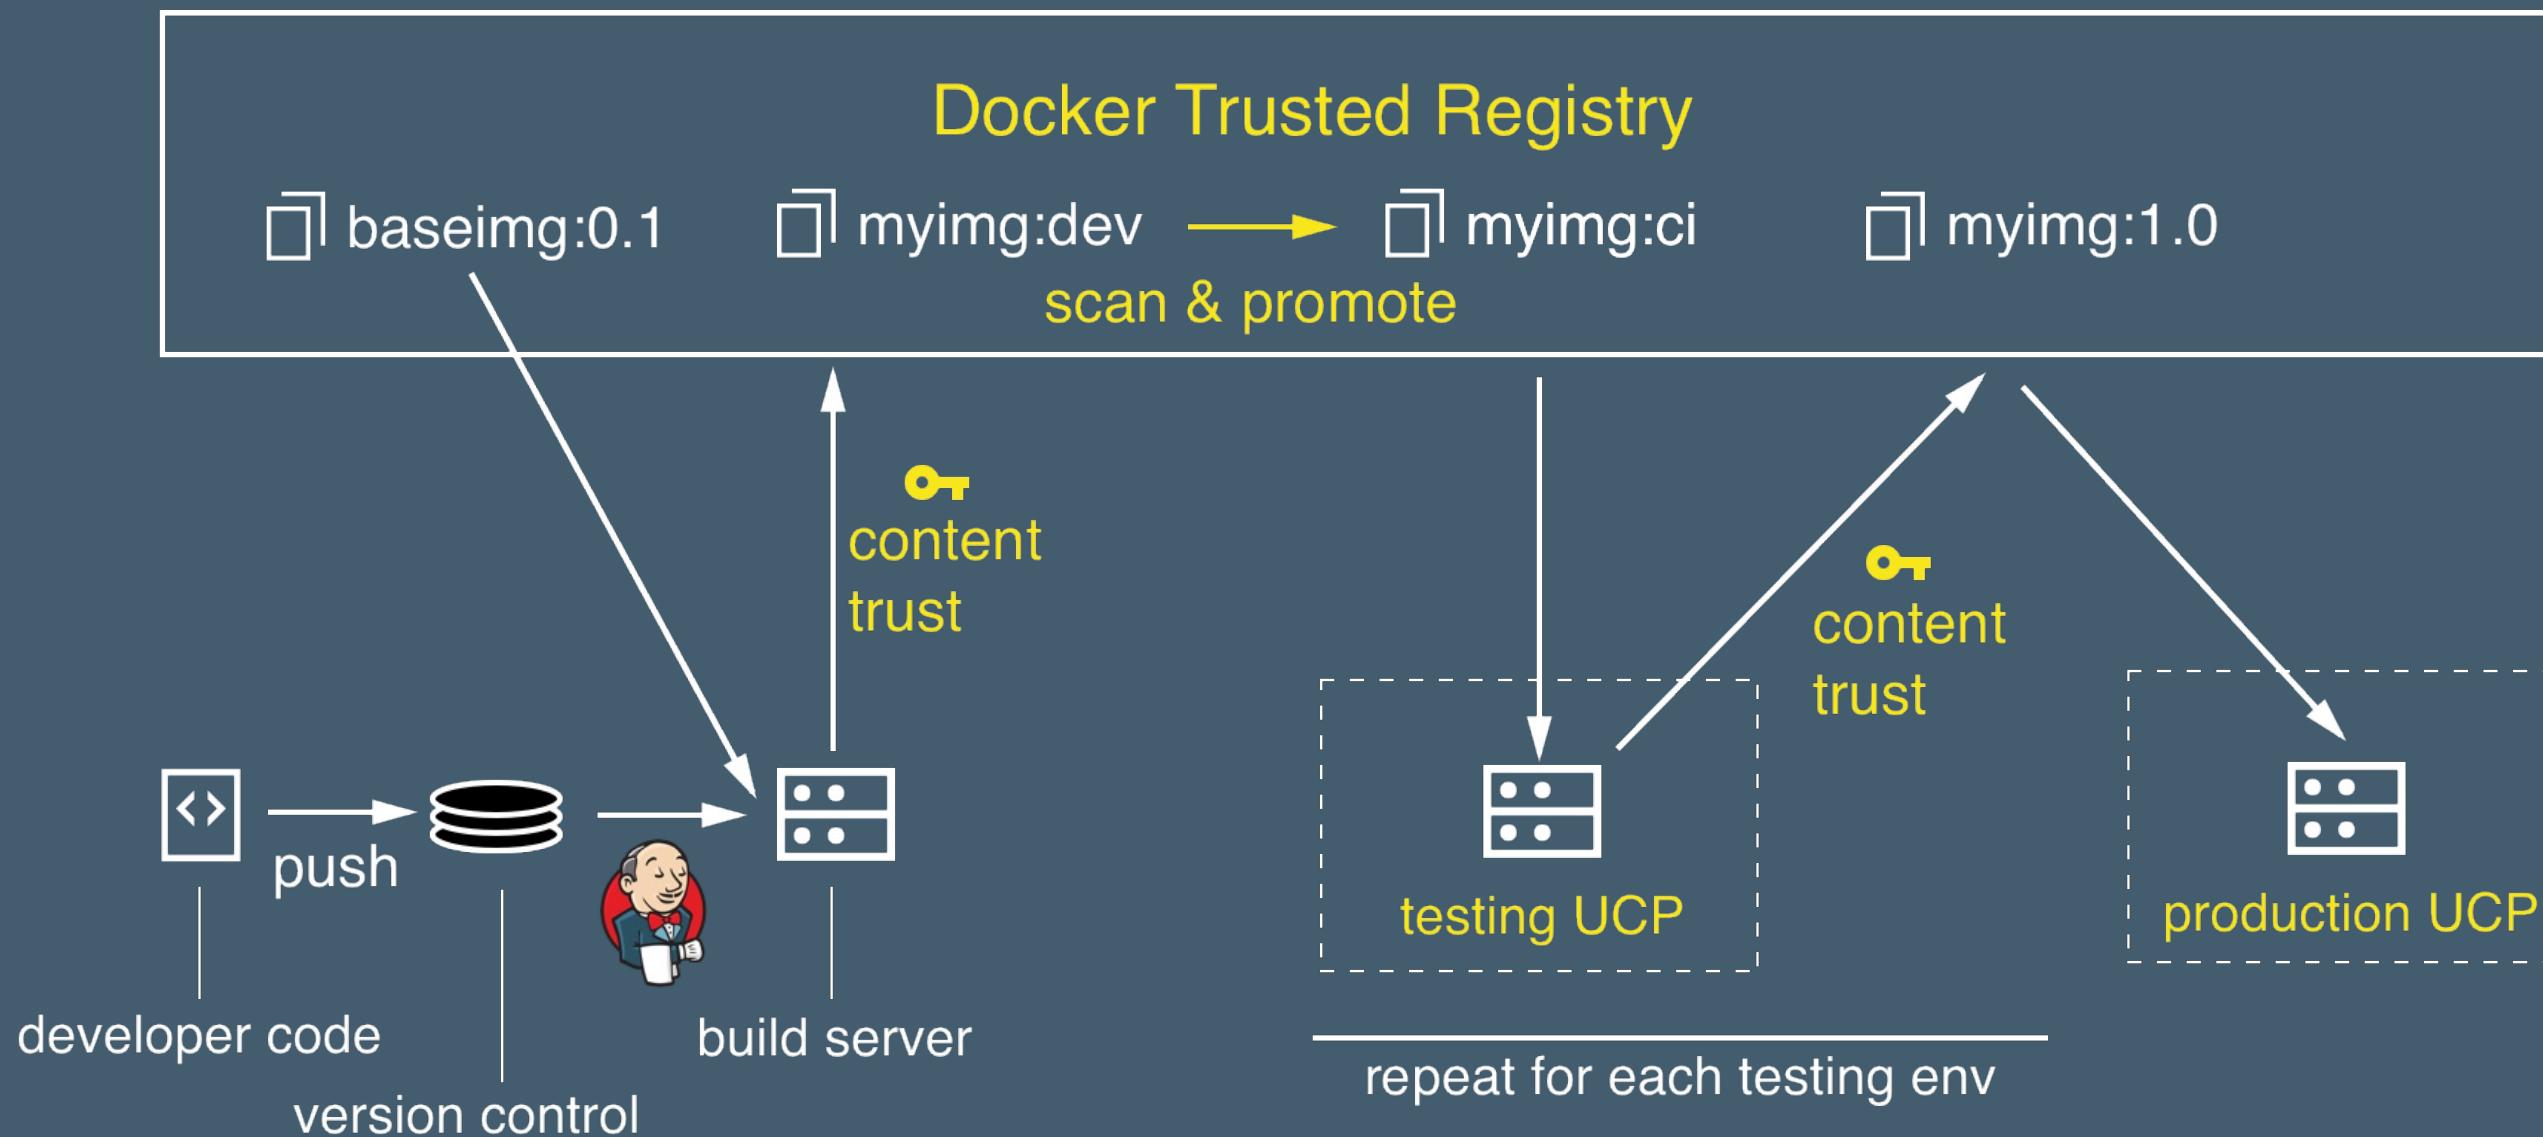


# CONTAINERS IN CI/CD

- Images are full stack, not code alone
- Must check for dependency vulnerabilities
- Must ensure provenance
- UCP Clusters must reflect pipeline environments



# A DOCKER PIPELINE



# PIPELINE ALTERNATIVES

- Dev UCP
- Node RBAC instead of separate UCPs
- Build-serverless pipeline
- DTR Webhooks



# DEPLOYMENT PHASE

- All familiar deployment strategies possible
- Need to think about how load balancing works on a swarm (external and internal)
- Consider routing meshes (L4 vs L7)



# CONTAINERS AS A SERVICE

- (Dev)ops teams can provide base images
- Typically contain OS, middleware, enterprise credentials
- Developers use these as starting point



# FURTHER READING

- Dockerized CI/CD Reference Architecture: <http://dockr.ly/2tETN0V>
- Docker EE Best Practices Reference Architecture: <http://dockr.ly/2ohLad6>
- Docker Reference Architecture: Development Pipeline Best Practices Using Docker EE <http://dockr.ly/2tETN0V>





# UNIVERSAL CONTROL PLANE



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Install and uninstall a highly available UCP cluster

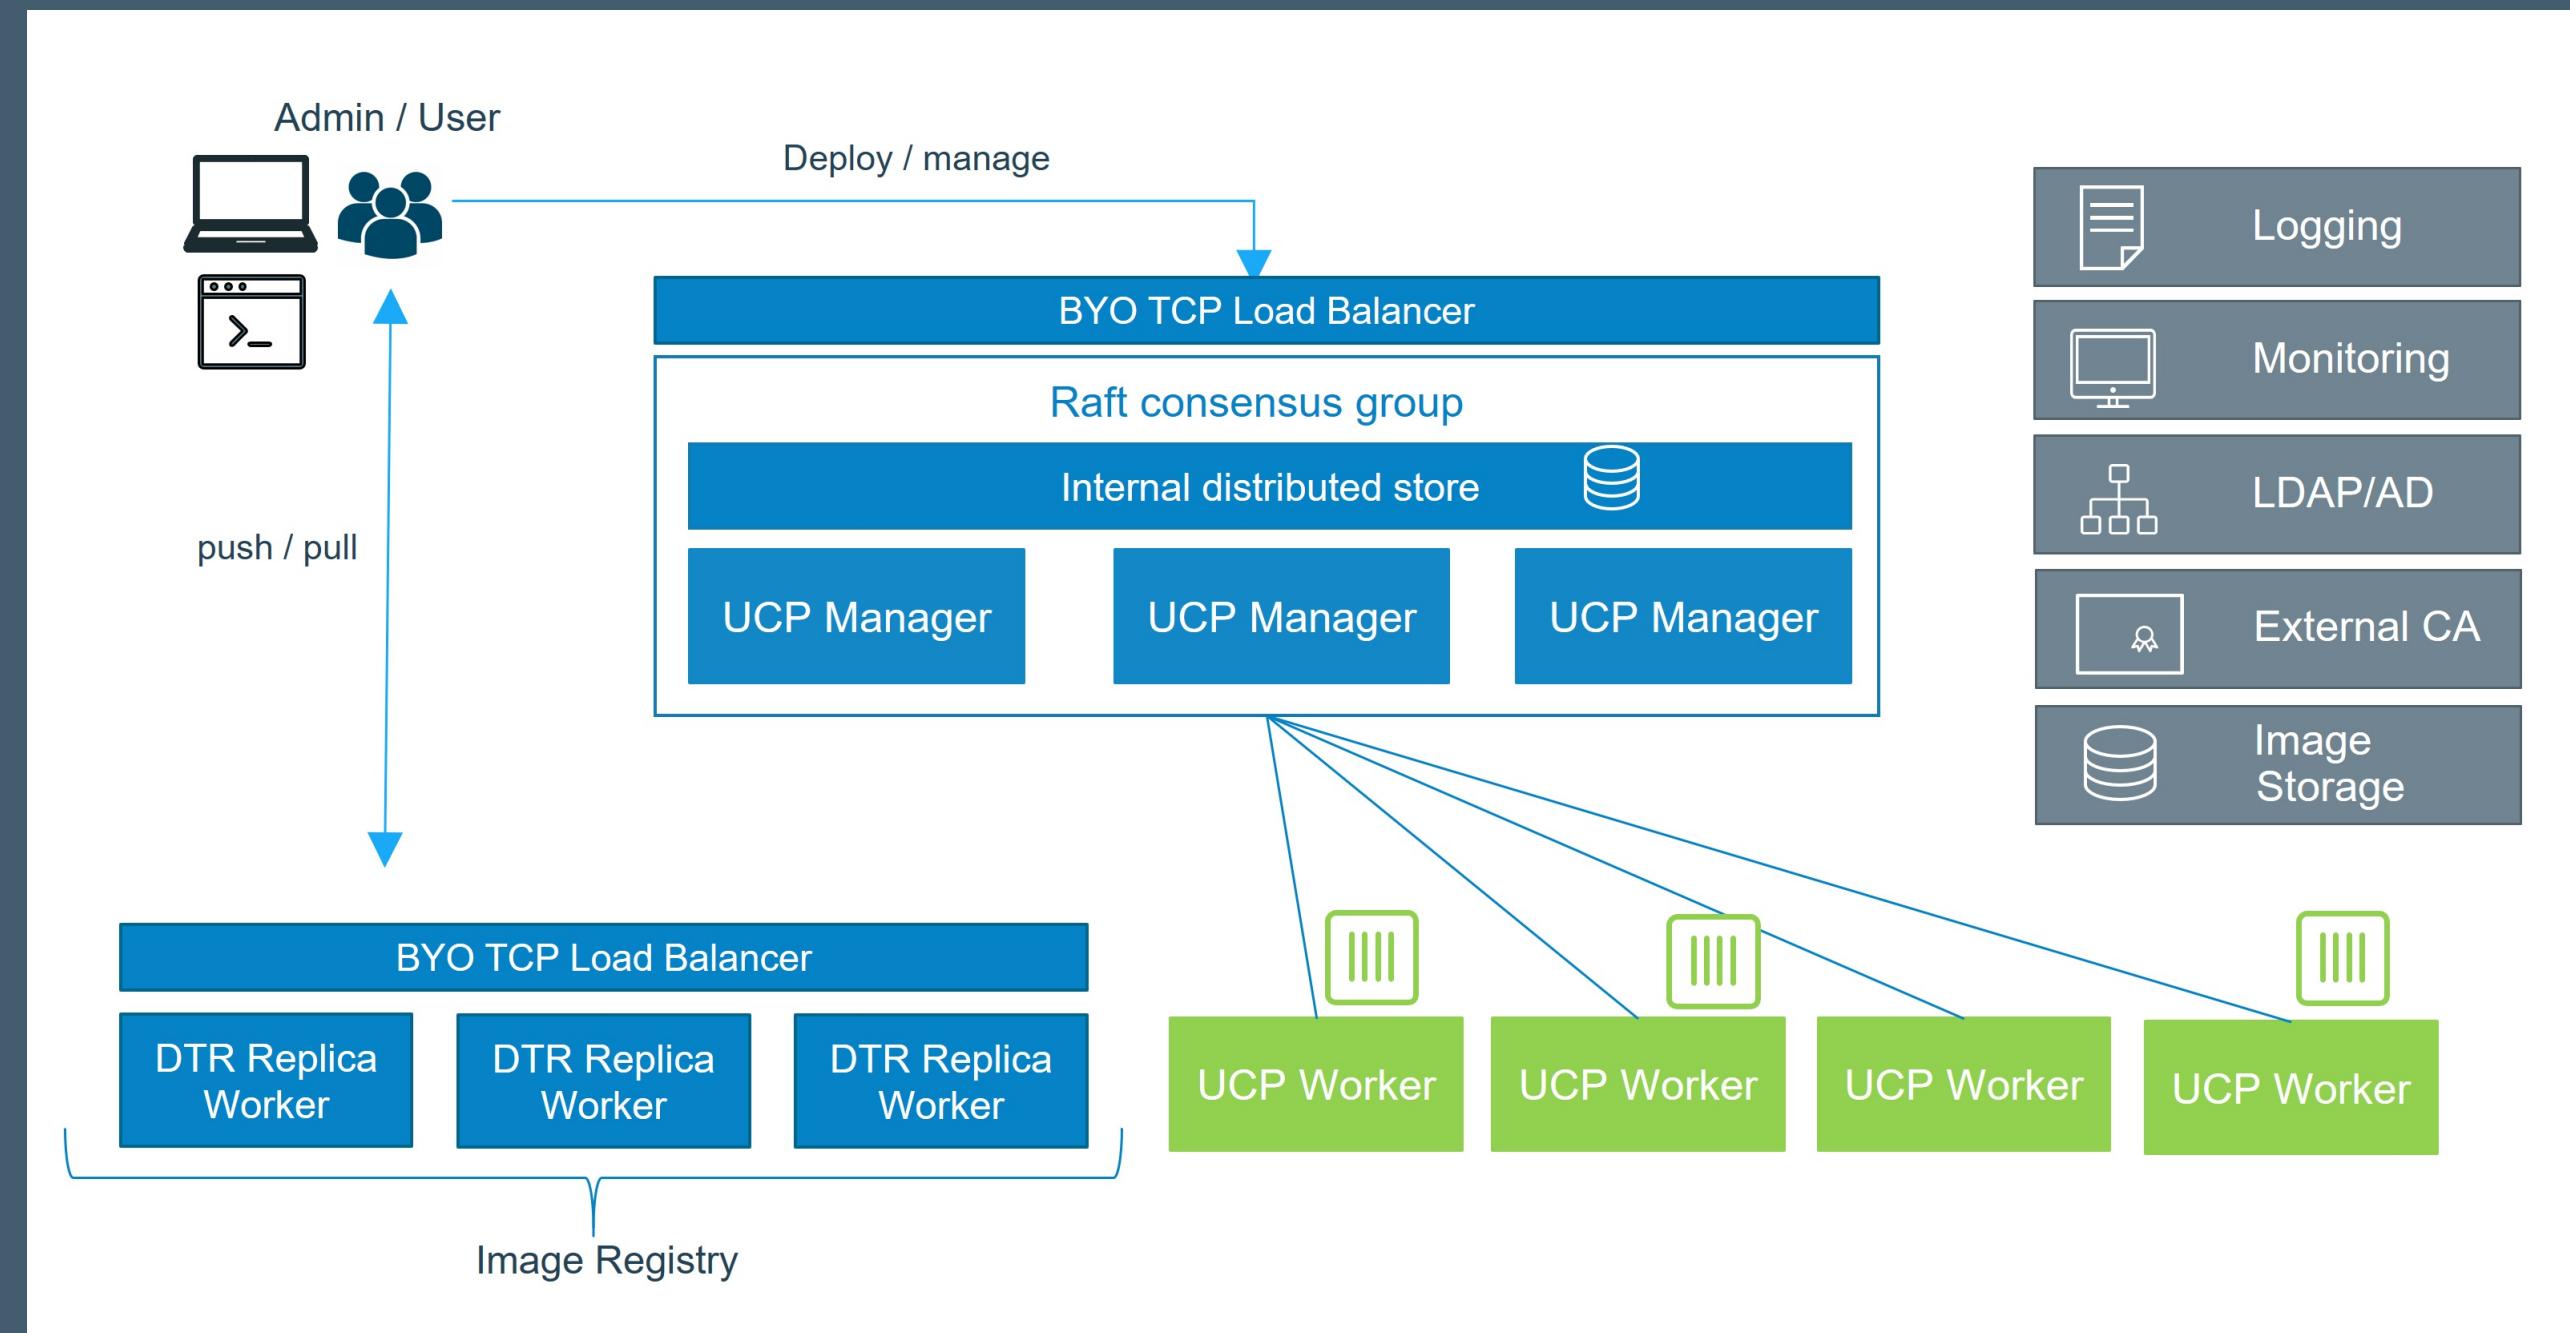


# DOCKER ENTERPRISE EDITION

- Security compliance tooling
- Containerized CI/CD
- Secure image distribution
- Role based access control



# DOCKER EE ARCHITECTURE



# UCP IS A SWARM

- UCP sits on top of a Swarm
- Automatically create Kubernetes Cluster
- Adds UI, RBAC, registry integration...
- Start with one UCP manager == Swarm manager leader
- New UCP nodes joined exactly as new Swarm nodes joined

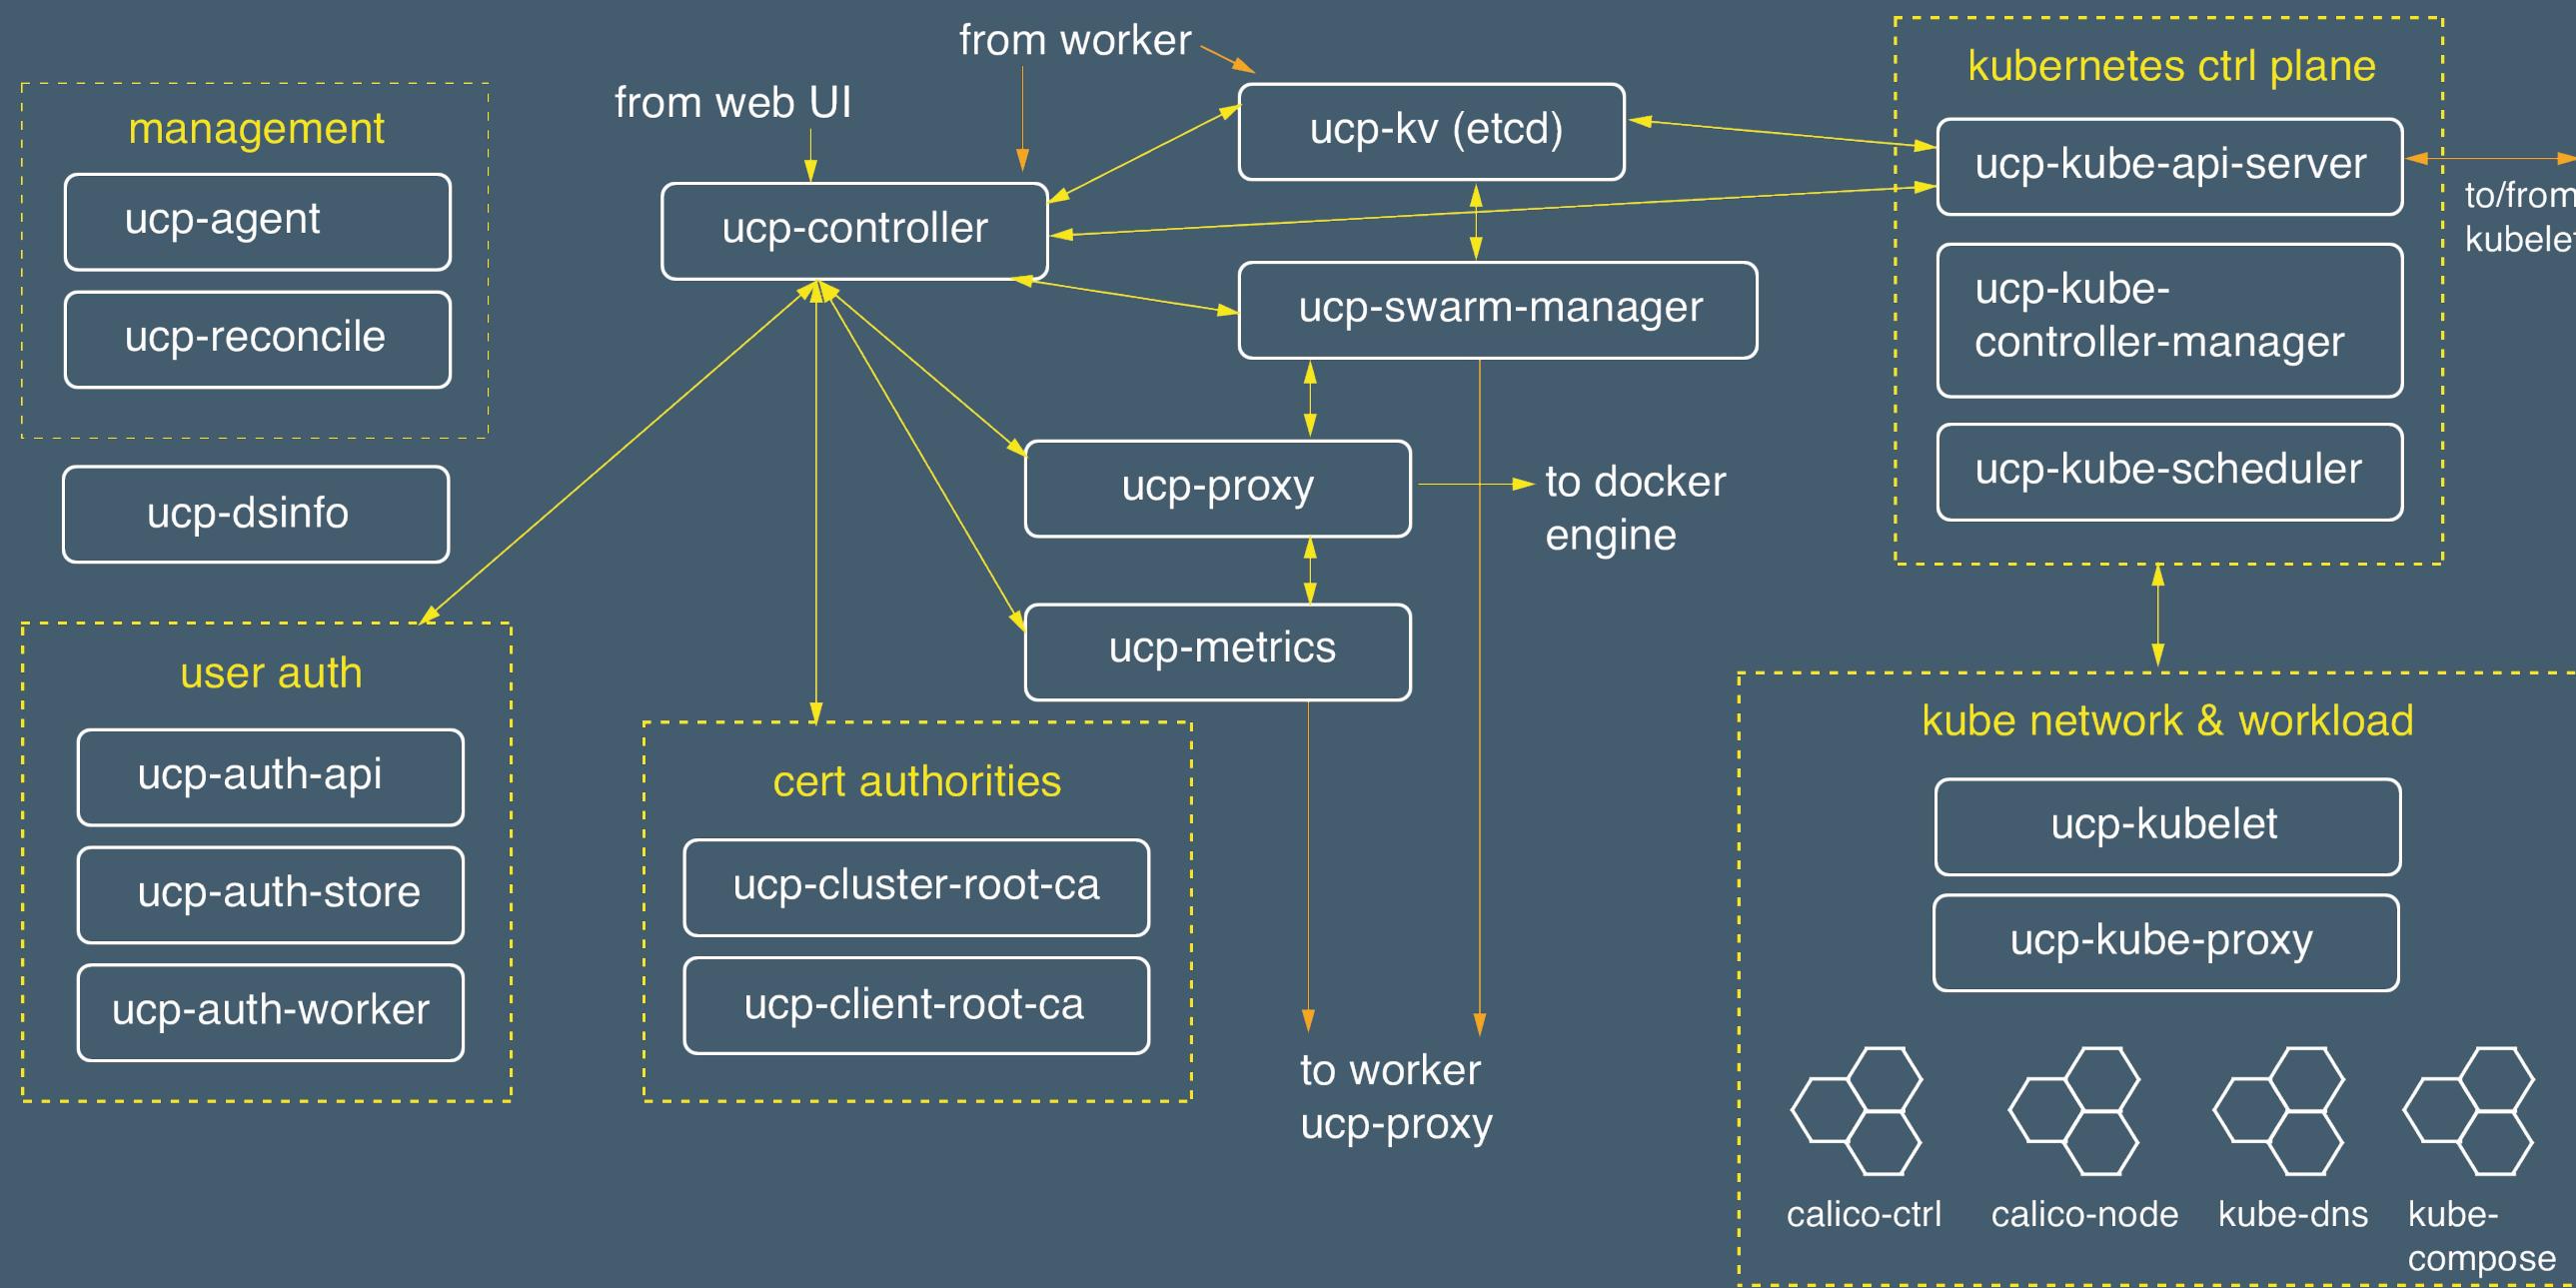


# CHOOSE YOUR ORCHESTRATION ENGINE

- Swarm
- Kubernetes (containerized and running on top of the Swarm)
- Classic Swarm (legacy only)



# UCP MANAGER ARCHITECTURE OVERVIEW





## EXERCISE: INSTALLING UCP

Work through the 'Installing UCP' exercise in the Docker for Enterprise Developers Exercises book.



## FURTHER READING

- High availability architecture and apps with DDC: <http://dockr.ly/1sqPrIH>
- UCP architecture: <https://dockr.ly/2qTOgTb>





# CONTEXT BASED ROUTING



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Diagram the components that provide layer 7 routing in UCP and Kubernetes
- Configure Swarm services and Kubernetes deployments to use layer 7 routing



# HOW IT WORKS

We don't like:

```
https://example.com:8080
https://example.com:8081
...
...
```

Much better:

```
https://foo.example.com
https://bar.example.com
or
https://example.com/cats
https://example.com/dogs
...
...
```



# LAYER 7 ROUTING

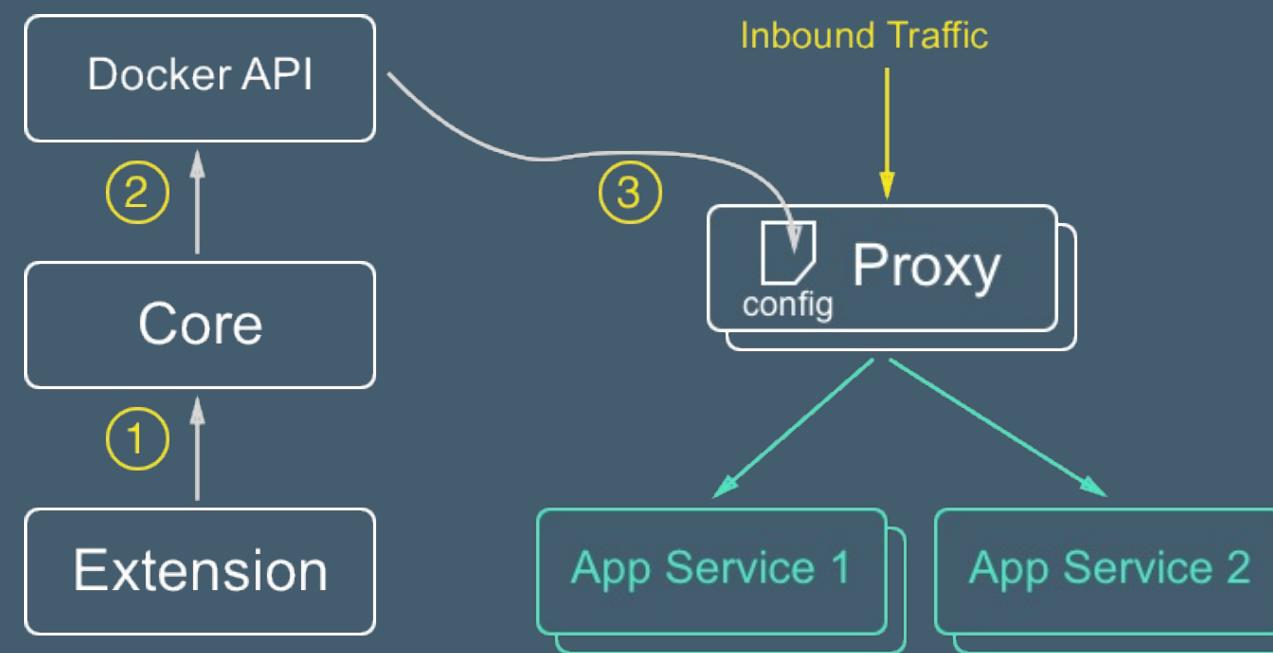
- L4 routing requires LB be aware of each service
- Application layer (L7) load balancing avoids this
- Route traffic to services based on
  - header information
  - path

```
Accept:/*
Accept-Encoding:gzip, deflate, sdch
Accept-Language:en-US,en;q=0.8,de;q=0.6,nb;q=0.4
Cache-Control:no-cache
Connection:keep-alive
Cookie:...
DNT:1
Host:pets.my-company.com
Pragma:no-cache
Referer:...
```

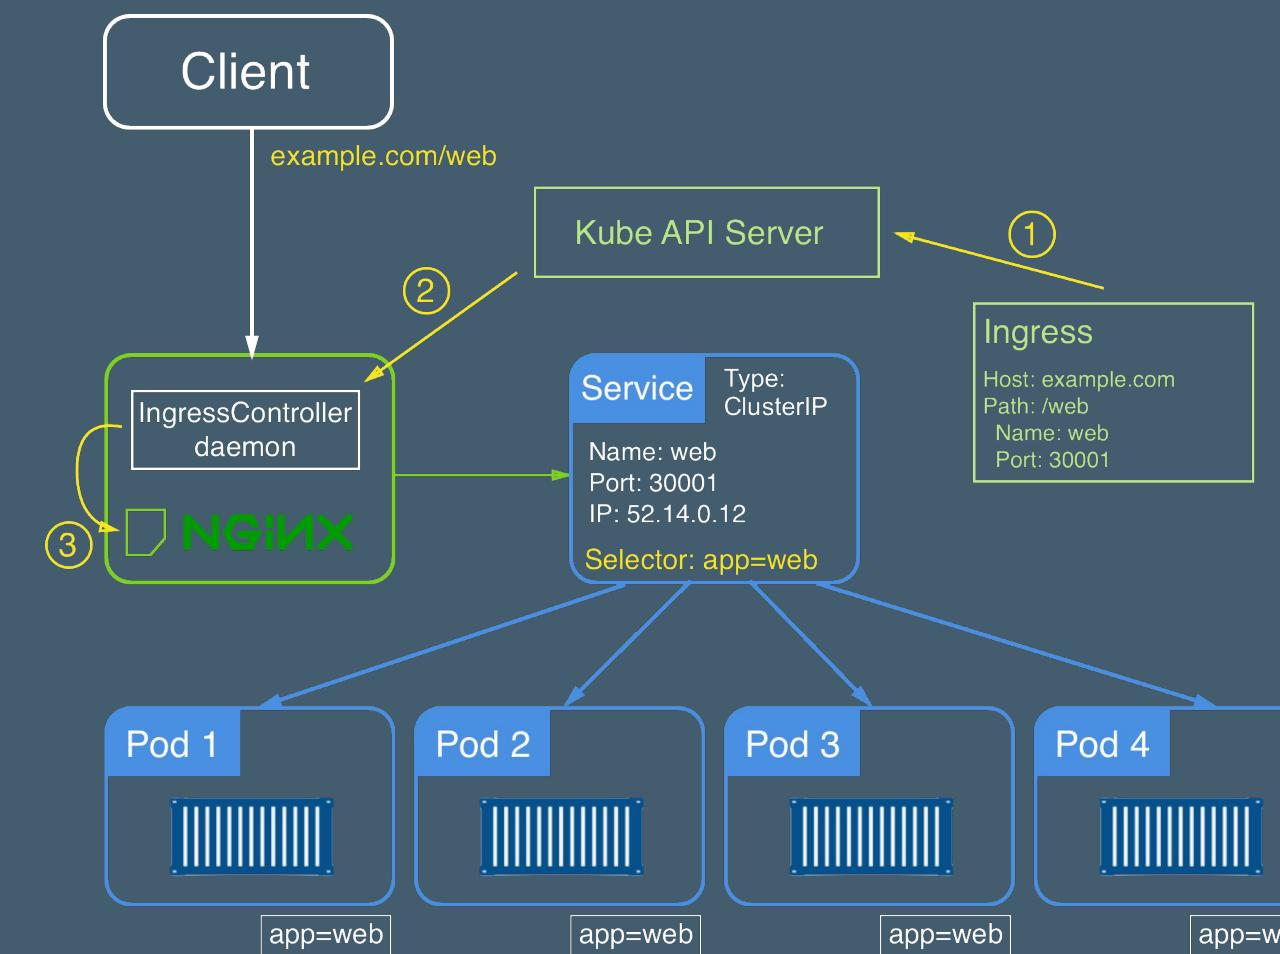


# INTERLOCK 2 & INGRESSCONTROLLER

Docker Swarm



Kubernetes





## EXERCISE: LAYER 7 LOAD BALANCING

Work through the 'Layer 7 Load Balancing' exercise in the Docker for Enterprise Developers Exercises book.



# DISCUSSION

- In what situations would L7 routing be preferable to L4 routing? Vice-versa?
- Questions?



## FURTHER READING

- Interlock 2: <http://bit.ly/2ETb4cx>
- Kubernetes IngressController: <http://bit.ly/2nMFaqL>
- Nginx IngressController: <http://bit.ly/2skeOko>
- Use swarm mode routing mesh: <http://dockr.ly/2ePQeOM>
- Docker Reference Architecture: UCP 2.0 Service Discovery and Load Balancing:  
<http://dockr.ly/2rbxDDX>
- Get to know the DDC networking updates: <http://dockr.ly/2fbZ61F>





# DOCKER TRUSTED REGISTRY



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Install a highly available DTR
- Configure a UCP node so it trusts DTR
- Push an arbitrary Docker image to a DTR repository
- Uninstall DTR



# SECURE SOFTWARE SUPPLY CHAIN

- Image Creation
- Image Distribution
- Container Execution



# DTR KEY FEATURES

- Image Creation:
  - Image Security Scanning
  - Repository Automation
  - Image Promotion
- Image Distribution:
  - Content Trust / Notary
  - Content Cache
- Image Storage:
  - Pluggable Storage Drivers



# PUSHING IMAGES TO DTR

- Configure your Docker engine to trust DTR

```
$ sudo curl -k https://<DTR_FQDN>/ca \
-o /etc/pki/ca-trust/source/anchors/<DTR_FQDN>.crt
$ sudo update-ca-trust
$ sudo /bin/systemctl restart docker.service
```

- Login to DTR:

```
$ docker login <DTR_FQDN>
```

- Tag your images with:

```
<DTR_FQDN>/<Org Name>/<Image Name>:<Tag>
```

- Push image

Note: **<Image Name>** and **<Repository Name>** are equivalent.





## EXERCISE: INSTALLING DTR

Work through the 'Installing Docker Trusted Registry' exercise in the Docker for Enterprise Developers Exercises book.



## FURTHER READING

- DTR overview: <https://dockr.ly/2F84rkN>
- Troubleshoot DTR: <https://dockr.ly/2JhT78m>
- Integrate with multiple registries: <https://dockr.ly/2HqBzX7>





# CONTENT TRUST



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Identify the kinds of attacks and countermeasures considered by The Update Framework
- Describe the differences in Docker's behavior with content trust enabled versus disabled
- Establish content trust for a repository in DTR
- Demand UCP only runs images signed by a set of teams

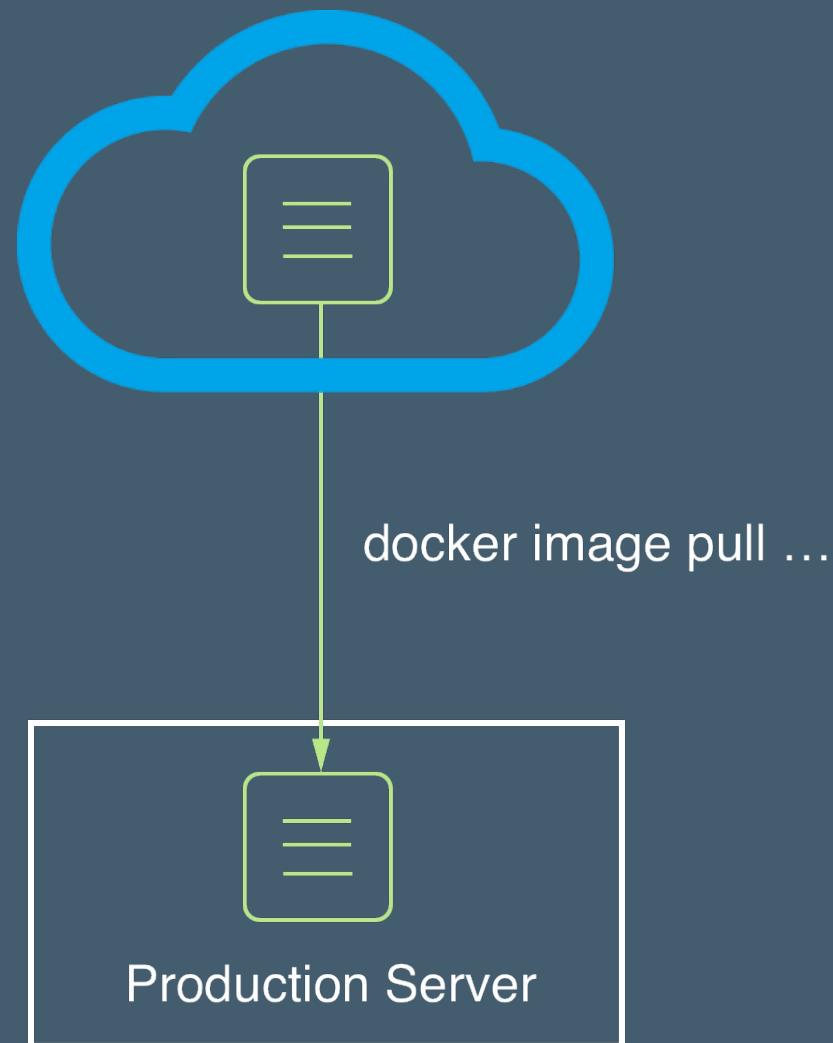


# WHAT IMAGE DO YOU MEAN?

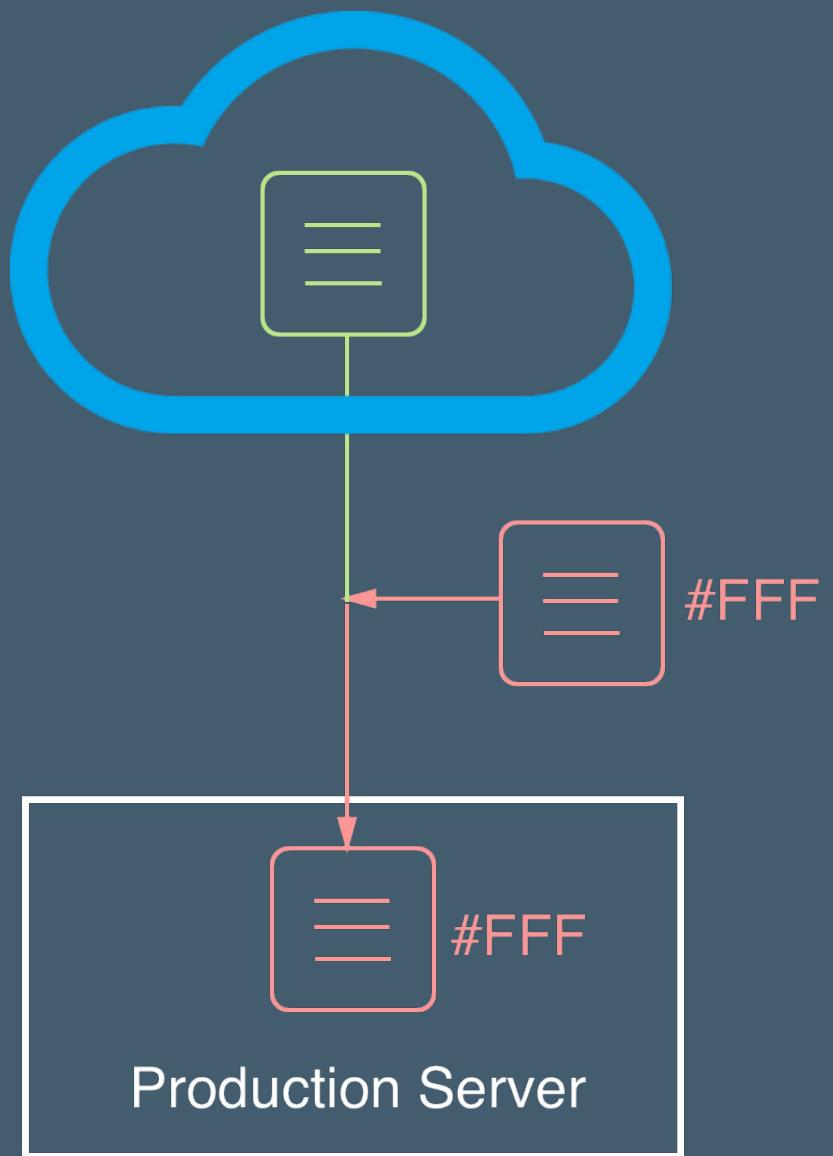
- Mutable: ref by tag: **alpine:latest**
- Immutable: ref by sha: **alpine@sha256:FFFF....**



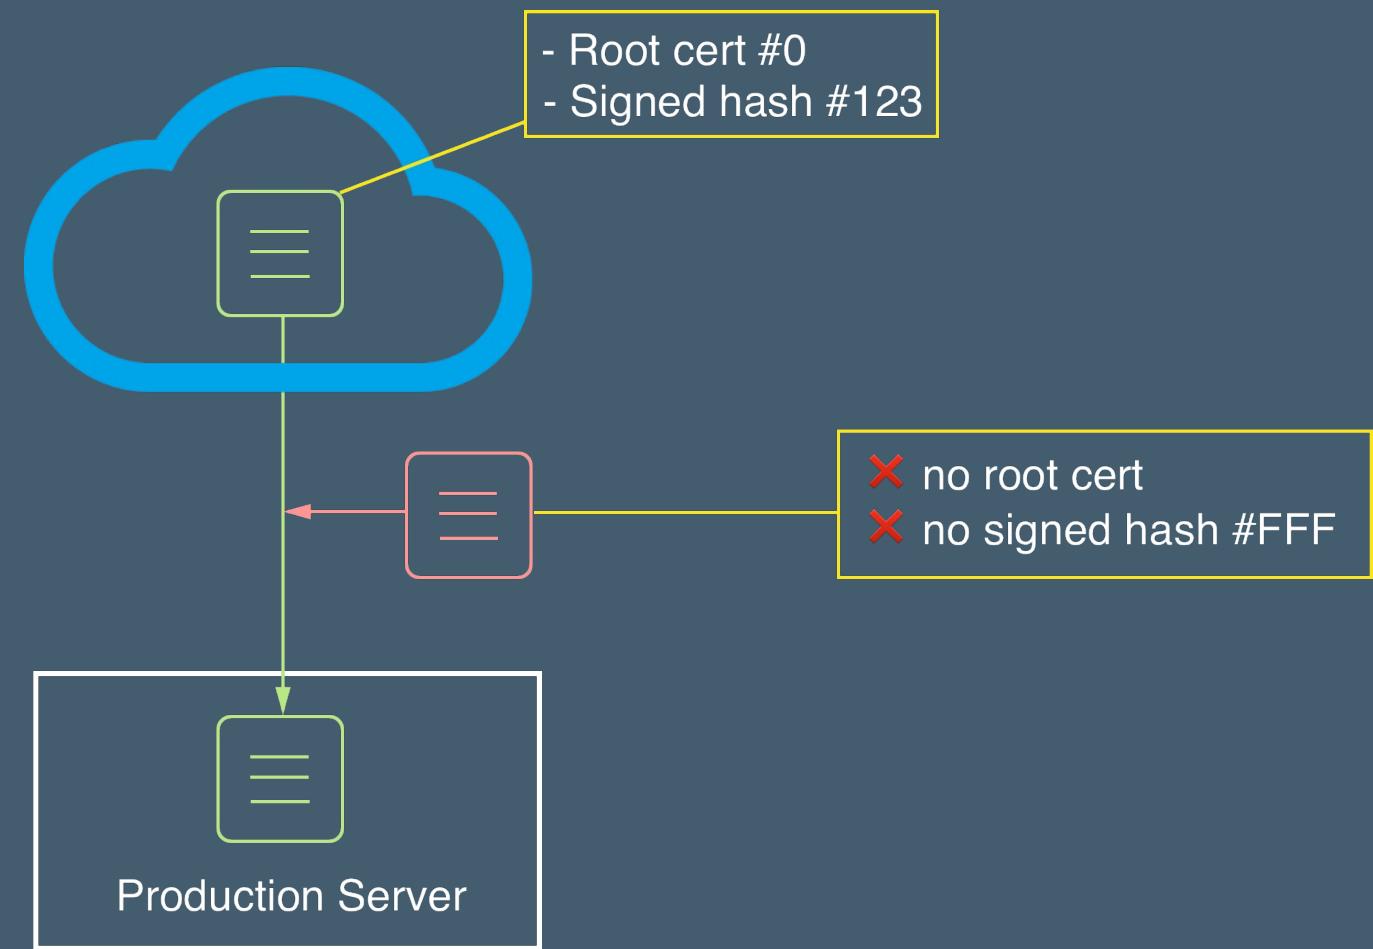
# DOWNLOADING SOFTWARE



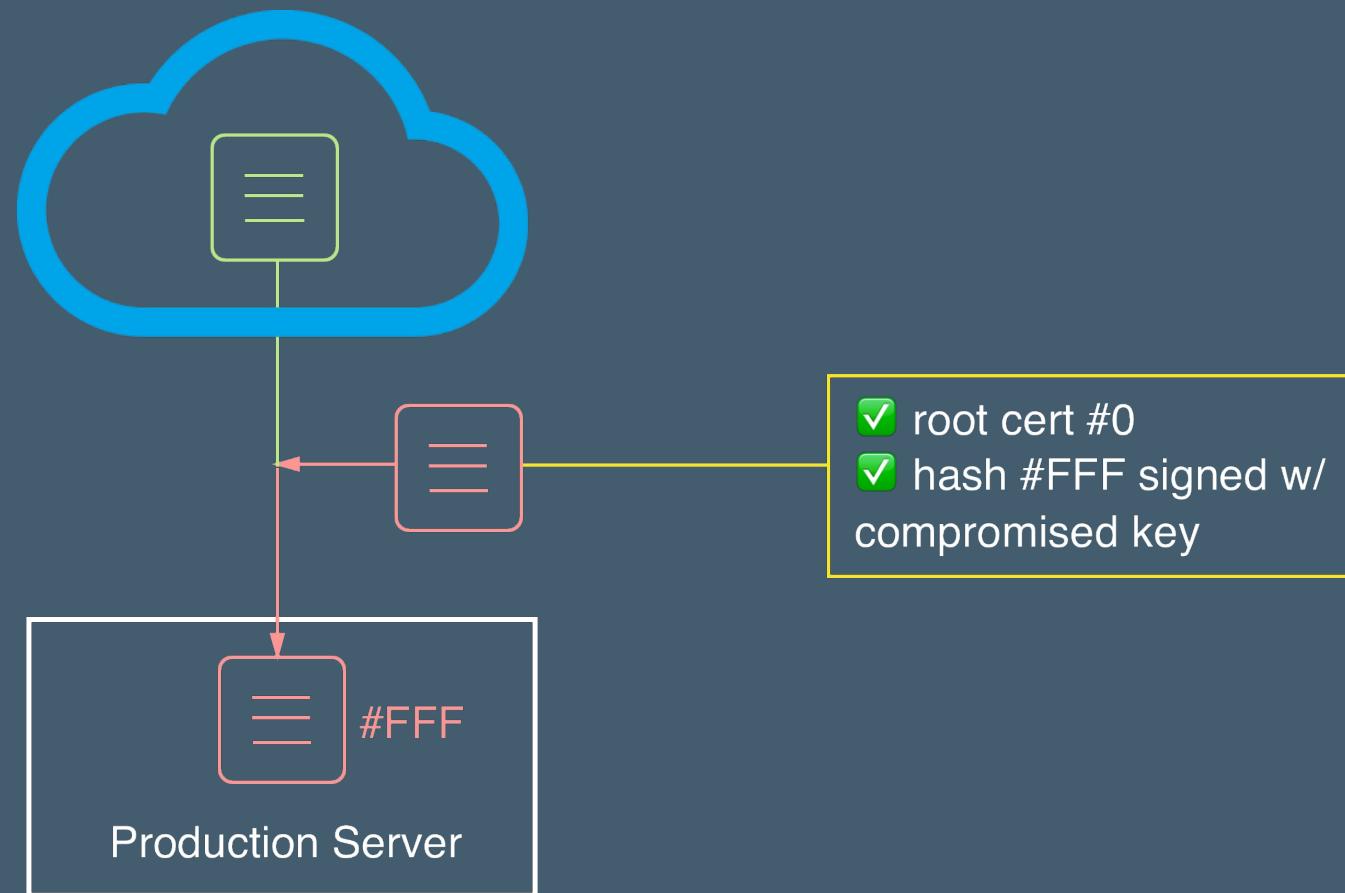
# SIMPLE INJECTION ATTACK



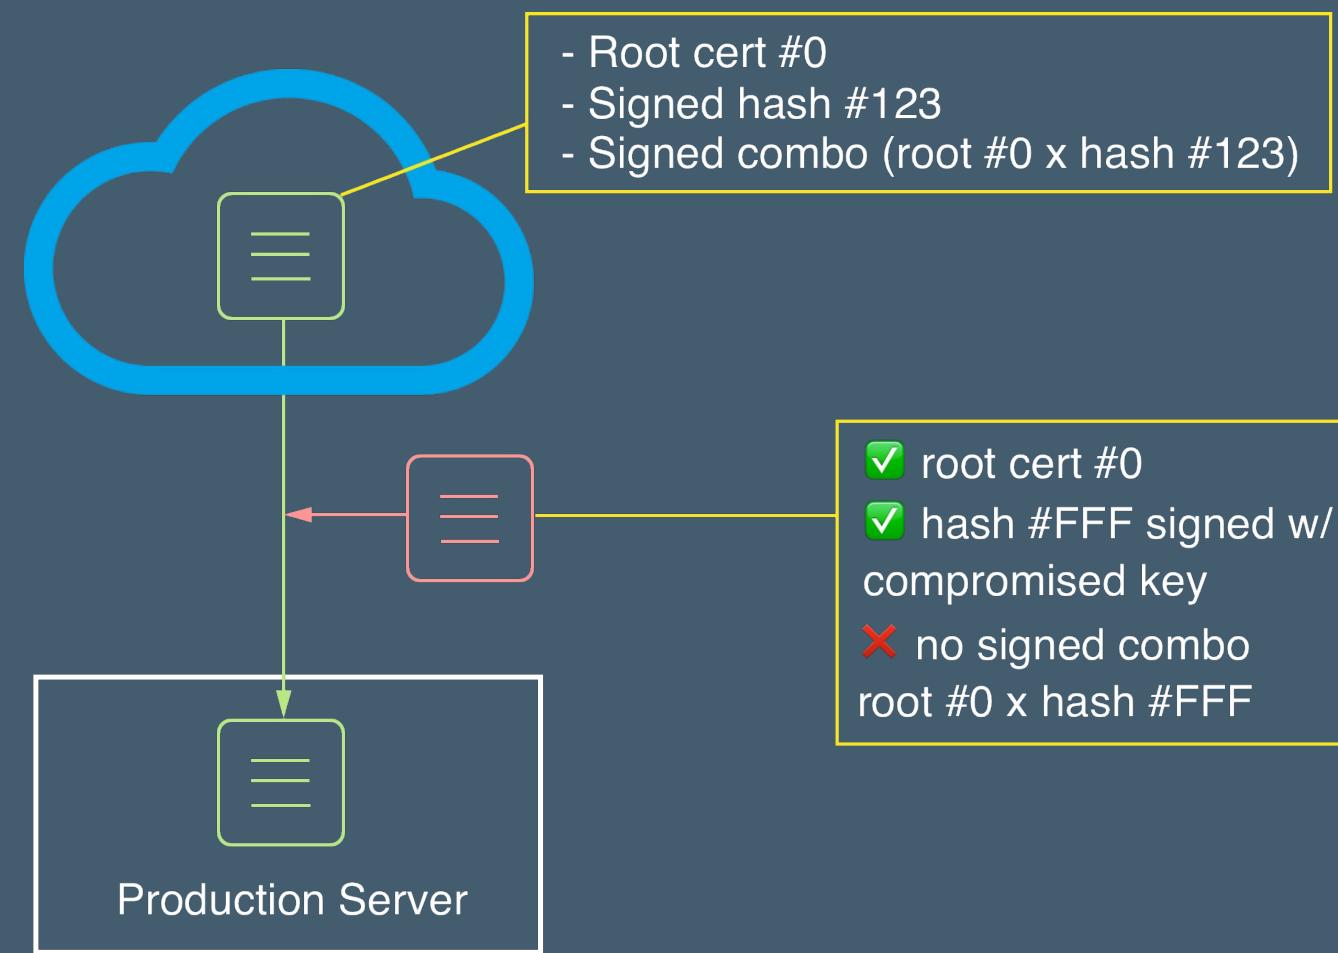
# SIGNED CONTENT



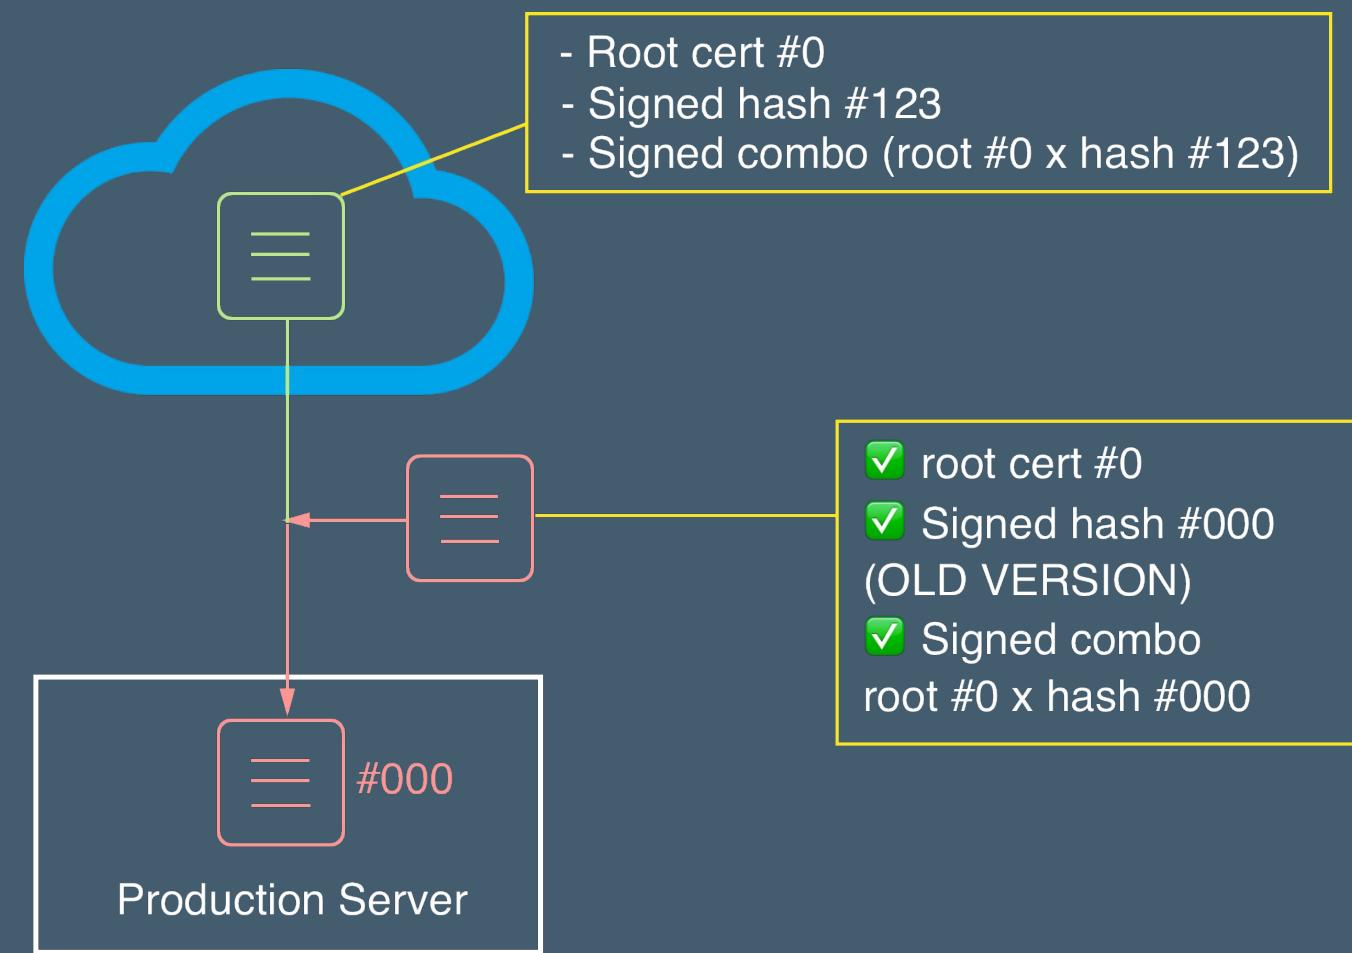
# MIX AND MATCH ATTACK



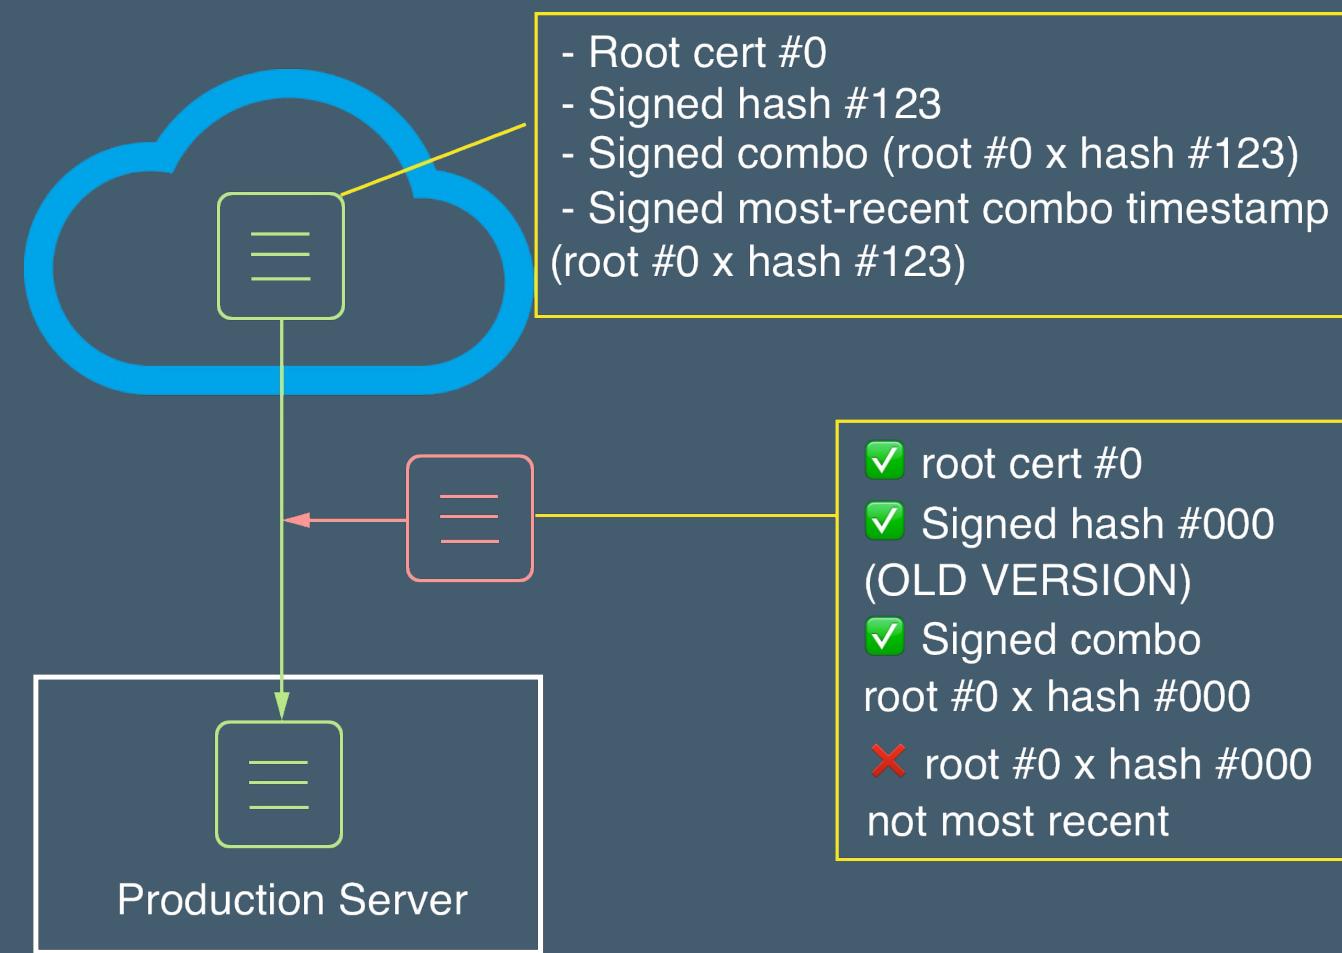
# METADATA VALIDATION



# DOWNGRADE ATTACK



# FRESHNESS GUARANTEED



# DOCKER CONTENT TRUST

- Image publishers sign their images
- Image consumers can ensure their images are signed
- Implements The Update Framework (TUF) into Docker as Notary
  - <http://theupdateframework.com/>
  - <https://github.com/docker/notary>



# CONTENT TRUST FOR IMAGE PUBLISHERS

- Metadata generated and signed on image push certifying authenticity
- Four (+1) different keys are used:
  - Root key: fundamental signing authority
  - Target key: repository content integrity
  - Snapshot key: metadata integrity
  - Timestamp key: repository freshness
  - Delegation key: multiple signing entities (optional)

(Jargon aside: the target key and snapshot key are sometimes collectively called the repository key, and the root key is sometimes called the offline key).



# CONTENT TRUST FOR IMAGE CONSUMERS

If content trust is enabled, only signed images are available for use with:

- docker image push
- docker image pull
- docker image build
- docker container create
- docker container run
- docker service create



# CONTENT TRUST IN DTR

- DTR comes with a built in Notary server
- Can enable trust directly from Docker Engine
- Optional: install Notary client side for finer-grained control



# INTEGRATION WITH UCP

- UCP can be configured to only run signed images
- Checked during application deployment
- Any unsigned image will be rejected
- Multiple signatures possible

**Admin Settings**

The screenshot shows the 'Admin Settings' interface with a sidebar on the left containing links like Swarm, Certificates, Routing Mesh, Cluster Configuration, Authentication & Authorization, Logs, License, Docker Trusted Registry, and Docker Content Trust (which is selected). The main content area is titled 'Content Trust Settings' and contains a checked checkbox for 'Run Only Signed Images'. Below it is a note: 'Select an Org and then Team in that Org. Signatures from all Teams listed required.' There is a 'Add Team +' button and two dropdown boxes: one containing 'engineering' and another containing 'qa'. A trash icon is next to the 'qa' box.

**Content Trust Settings**

Run Only Signed Images ?

Select an Org and then Team in that Org. Signatures from all Teams listed required.

Add Team +

engineering qa





## EXERCISE: CONTENT TRUST

Work through the 'Content Trust' exercise in your exercise book.



# DISCUSSION

- What is one simple alternative to content trust?
- Questions?



# FURTHER READING

- Intro to content trust: <http://dockr.ly/2gAglo3>
- Features of Docker content trust: <http://dockr.ly/1EyCxrR>
- Automation with content trust: <http://dockr.ly/2x5lLaC>





# IMAGE SECURITY SCANNING



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Describe the process DTR takes to scan images for vulnerabilities
- Set up security scanning on DTR, and configure it for an individual repository
- Use DTR's UI to find out an image's vulnerabilities, and more information about that vulnerability



# WHAT'S IN YOUR IMAGES?

- Components installed in base layer
- Components installed in custom layers
- Known vulnerabilities?
- (Currently) unknown vulnerabilities?



# DTR SECURITY SCANNING FLOW

1. New layer pushed to DTR
2. Scan triggered (auto or manual)
3. Bill-of-materials generated and saved (**slow, resource intensive**)
4. BoM re-checked against CVE database every time the db is updated (**fast**)
5. Regular database updates from <https://dss-cve-updates.docker.com/>



# SCAN RESULTS PER LAYER

engineering/engineering-app private : latest

linux/amd64 1eb35305b7 47.15 MB Pushed 10 minutes ago by pixel 6 critical 10 major 14 minor All layers already scanned

[Layers](#) [Components](#) [Delete](#) [Promote](#) [Scan](#)

| Layer                                                                                                                                                                                                                  | Content                                                                    | Actions              |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------|----------------------|
| 1 ADD                                                                                                                                                                                                                  | file:96db69a1ba6c80f604d07b14bcbf84445624ad3eb5b0471eddabf09cb7925366 in / | <a href="#">show</a> |
| 2 set -xe && echo '#!/bin/sh' > /usr/sbin/policy-rc.d && echo 'exit 101' >> /usr/sbin/policy-rc.d && chmod +x /usr/sbin/policy-rc.d && dpkg-divert --local --rename --add /sbin/initctl && cp -a /usr/sbin/policy-rc.d | 47.15 MB                                                                   | <a href="#">show</a> |
| 3 rm -rf /var/lib/apt/lists/*                                                                                                                                                                                          |                                                                            |                      |

**ADD**  
file:96db69a1ba6c80f604d07b14bcbf84445624ad3eb5b0471eddabf09cb7925366 in /

**47.15 MB** [show](#)

**COMPONENTS (46)** **VULNERABILITIES (30) ▾**

**zlib 1.2.8.dfsg-2ubuntu4.1** **2 critical 2 major**

**ncurses 6.0+20160213-1ubuntu1** **2 critical 2 major**



# SCAN RESULTS PER COMPONENT

 engineering/engineer/app private : latest

 linux/amd64 1eb35305b7 47.15 MB Pushed 11 minutes ago by pixel ! 6 critical 10 major 14 minor All layers already scanned

[Layers](#) [Components](#) [Delete](#) [Promote](#) [Scan](#)

| Component                                                            | Version                              | License           | Vulnerabilities                                                                                                                                                                     |
|----------------------------------------------------------------------|--------------------------------------|-------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>zlib</b><br>1.2.8.dfsg-2ubuntu4.1<br><b>2 critical 2 major</b>    | <b>zlib</b><br>1.2.8.dfsg-2ubuntu4.1 | <b>PERMISSIVE</b> | <a href="#">CVE-2016-9841</a> <b>critical</b> inffast.c in zlib 1.2.8 might allow context-dependent attackers to have unspecified impact by leveraging improper pointer arithmetic. |
| <b>ncurses</b><br>6.0+20160213-1ubuntu1<br><b>2 critical 2 major</b> |                                      |                   | <a href="#">CVE-2016-9843</a> <b>critical</b> The crc32_big function in crc32.c in zlib 1.2.8 might allow context-                                                                  |
| <b>systemd</b><br>229-4ubuntu17<br><b>1 critical 2 major</b>         |                                      |                   |                                                                                                                                                                                     |



# VULNERABILITY OVERRIDES

pdevine/ubuntu: latest private

linux / amd64 56e96244fd 47.61 MB Pushed 20 hours ago by admin 12 critical 37 major 4 minor 1 hidden All layers already scanned

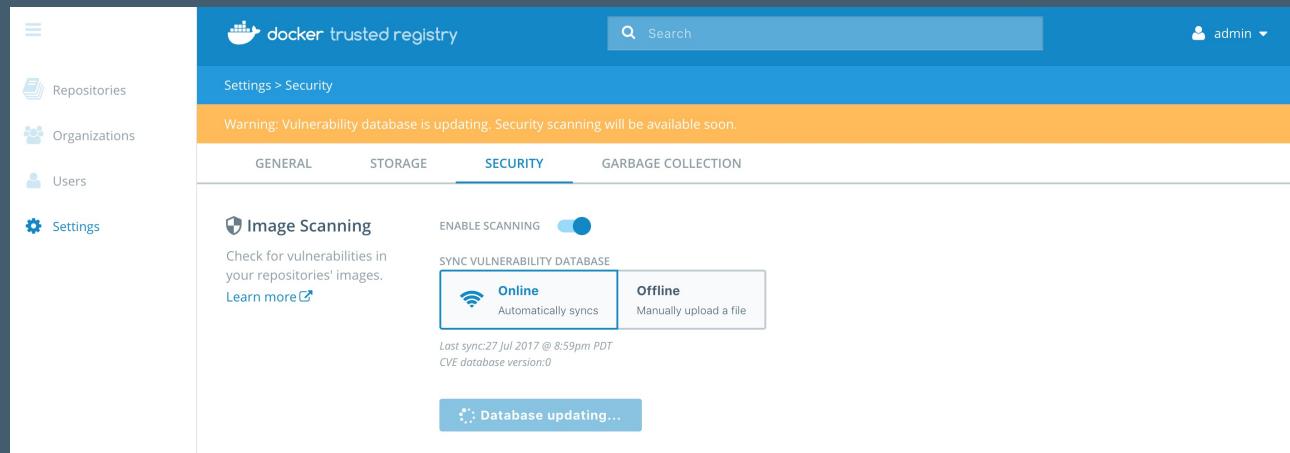
Layers Components Delete Promote Scan

| Component                        | Vulnerability ID | Severity | Description                                                                                                                                                                    | Action |
|----------------------------------|------------------|----------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|
| glibc<br>2.23-0ubuntu3           | CVE-2017-0325    | major    | The <code>makecontext</code> function in the GNU C Library (aka glibc or libc6) before 2.25 creates execution contexts incompatible with the...                                | hide   |
| systemd<br>229-4ubuntu6          | CVE-2015-5180    | major    | <code>res_query</code> in <code>libresolv</code> in glibc before 2.25 allows remote attackers to cause a denial of service (NULL pointer dereference and proce...              | hide   |
| ncurses<br>6.0+20160213-1ubuntu1 | CVE-2017-15671   | major    | The <code>glob</code> function in <code>glob.c</code> in the GNU C Library (aka glibc or libc6) before 2.27, when invoked with <code>GLOB_TILDE</code> , could skip freeing... | hide   |
| zlib<br>1.2.8.dfsg-2ubuntu4      | CVE-2017-12133   | major    | The DNS stub resolver in the GNU C Library (glibc) before version 2.26, when EDNS support is enabled, will solicit large UDP...                                                | hide   |
|                                  | CVE-2016-10228   | major    | The <code>iconv</code> program in the GNU C Library (aka glibc or libc6) 2.25 and earlier, when invoked with the <code>-c</code> option, enters an infinite lo...              | show   |
|                                  | CVE-2017-12132   | major    | The DNS stub resolver in the GNU C Library (aka glibc or libc6) before version 2.26, when EDNS support is enabled, will solicit...                                             | hide   |

- Allows vulnerabilities to be hidden
- Matches based on different components



# CVE DATABASE UPDATES



- DTR -> System -> Security
- Automatic updates: daily 3 AM UTC; must be able to reach <https://dss-cve-updates.docker.com/> on port 443.
- Manual updates: uploaded through DTR, downloaded through store.docker.com.



# SCANNING AUTOMATION

The screenshot shows the Docker Trusted Registry interface for repository settings. The left sidebar includes links for Repositories, Organizations, Users, and Settings. The main header indicates the path: Repositories > engineering/enterprise-app > Settings. The General tab is selected. Under Visibility, 'Private' is chosen (highlighted with a blue border). Under Immutability, 'Off' is selected (highlighted with a blue border). A description field is present but empty. In the Image scanning section, the 'Scan on push & Scan manually' option is selected (highlighted with a blue border), and a note states: 'Every image gets automatically scanned on push.' The footer displays the Docker Trusted Registry version 2.3.0-beta1 and the UCP logo.

Repositories > engineering/enterprise-app > Settings

**General**

VISIBILITY

**Public**  
Visible to everyone

**Private**  
Hide this repository

IMMUTABILITY

**On**  
Tags are immutable

**Off**  
Tags can be overwritten

DESCRIPTION

**Image scanning**

Check for vulnerabilities in your images.

[Learn more](#)

Scan on push &  
Scan manually

Scan manually  
(only)

Every image gets automatically scanned on push.

Docker Trusted Registry  
2.3.0-beta1  
UCP



# COMMON SCANNING MISTAKES

- Make sure initial CVE database is downloaded before starting scans
- Make sure all DTR replicas can reach storage backend
- Consider manual-only scans as part of pipeline





## EXERCISE: IMAGE SCANNING

Work through the 'Image Scanning in DTR' exercise in your exercise book.



# DISCUSSION

- Scanning identifies a vulnerability. What are some generic steps to mitigate?
- Questions?



## FURTHER READING

- Set up security scanning in DTR: <https://dockr.ly/2HMgp9d>
- Scan images for vulnerabilities: <https://dockr.ly/2HNOdCK>





# REPOSITORY AUTOMATION



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Automatically retag an image from one DTR repo to another
- Define webhooks triggered by DTR events
- Integrate DTR into a CI/CD chain using the above

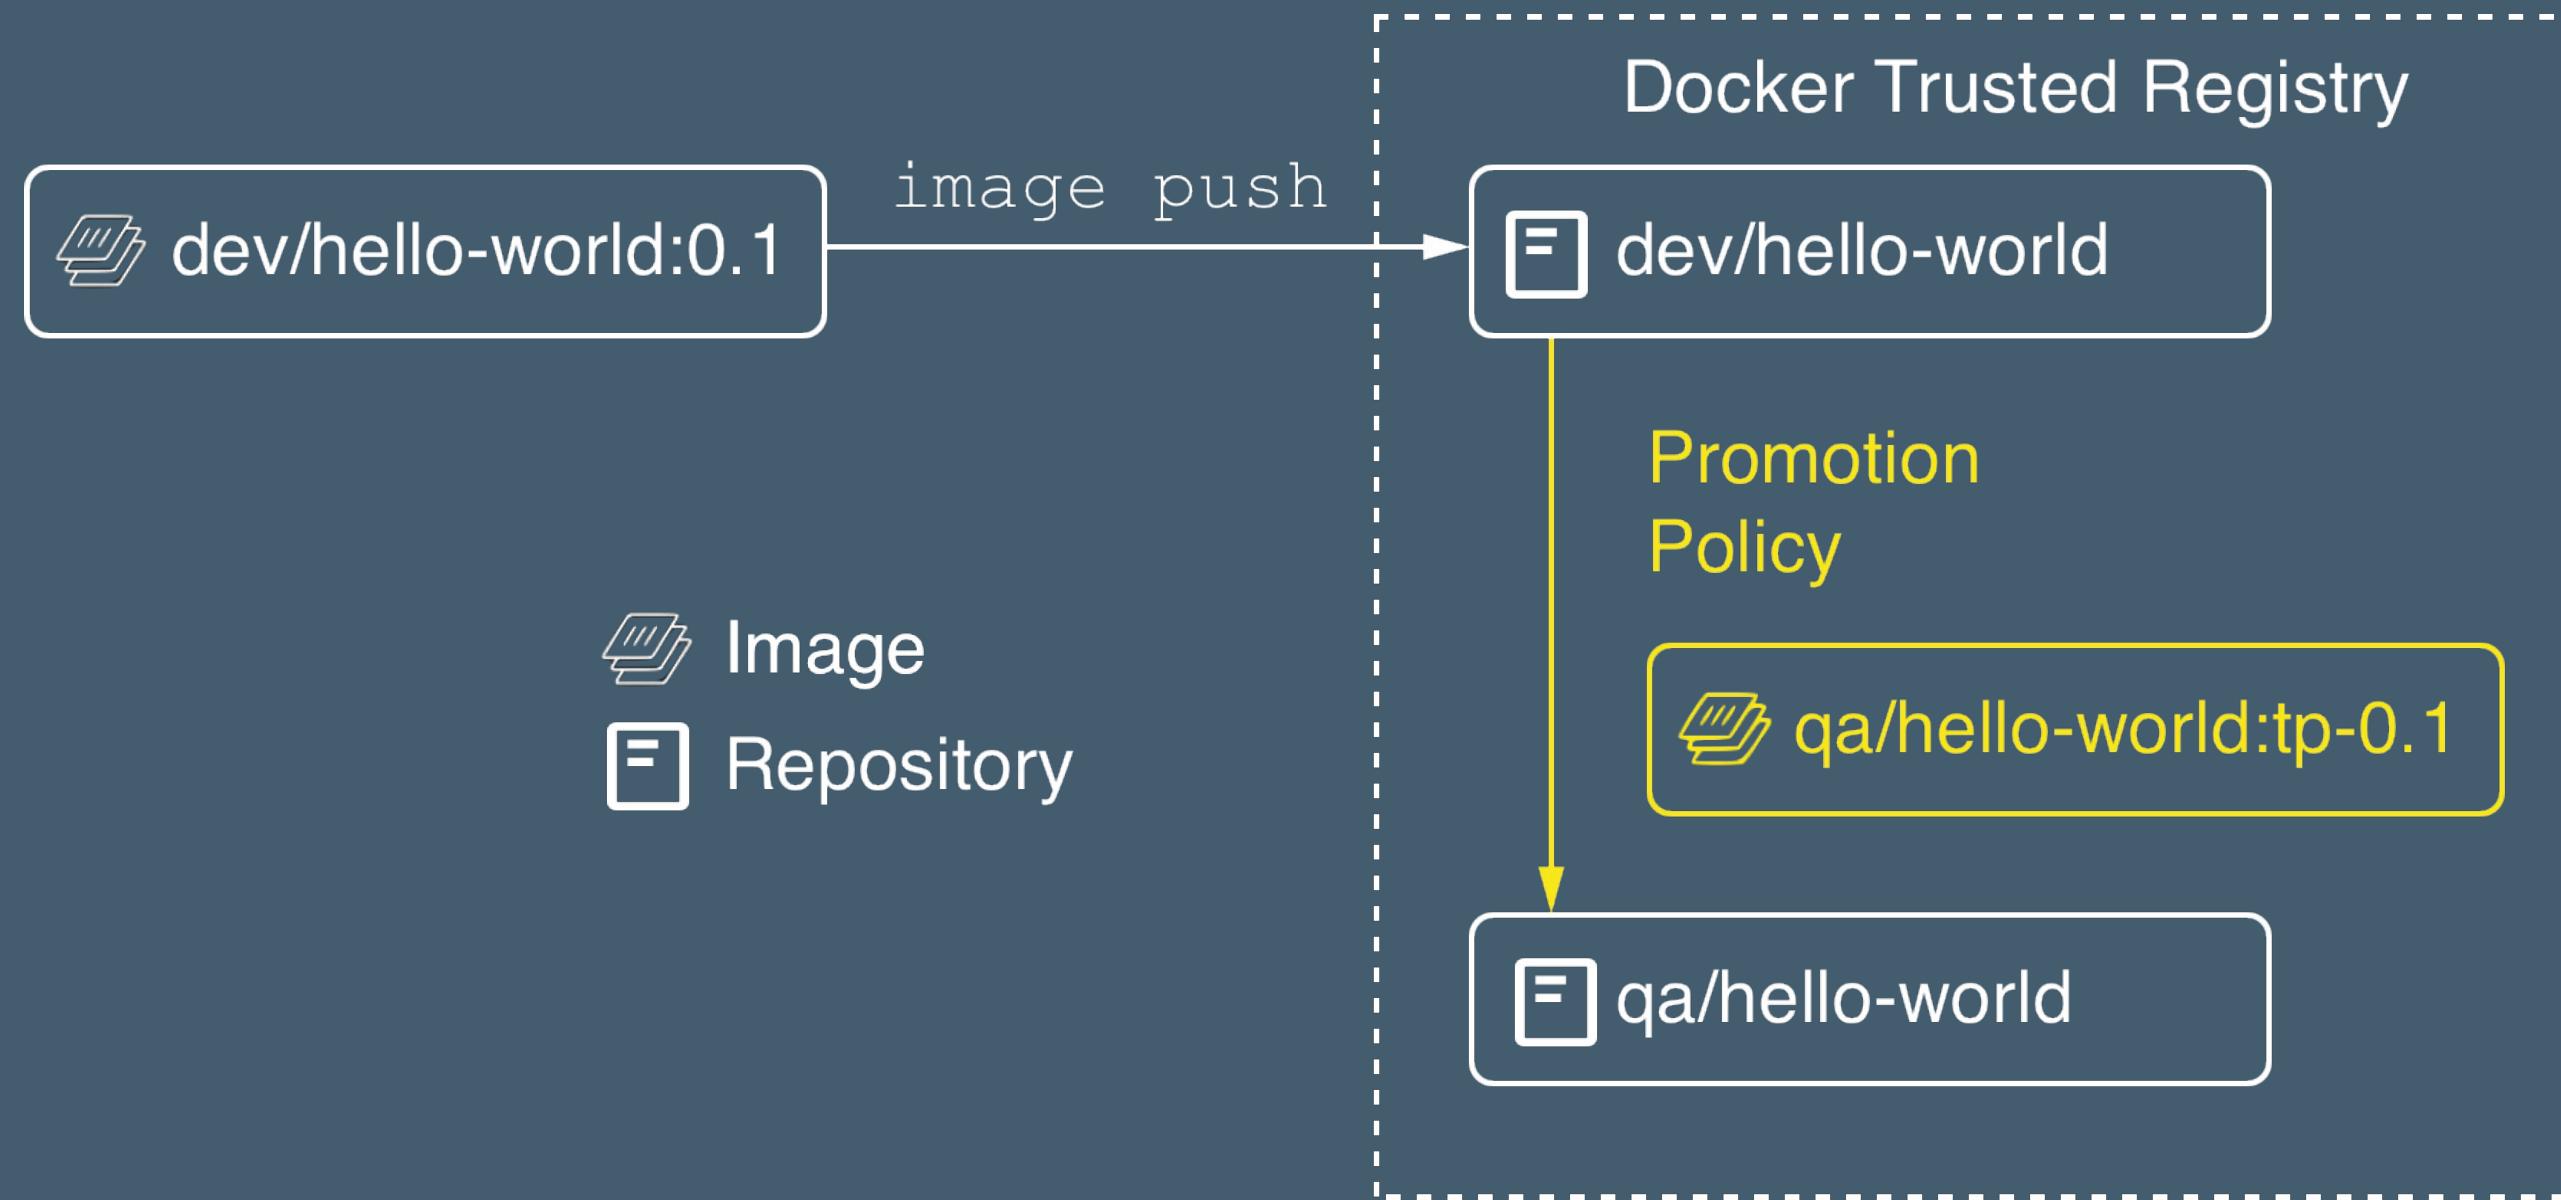


# AUTOMATION TOOLS

- Image Promotion & Mirroring
- Webhooks



# IMAGE PROMOTION

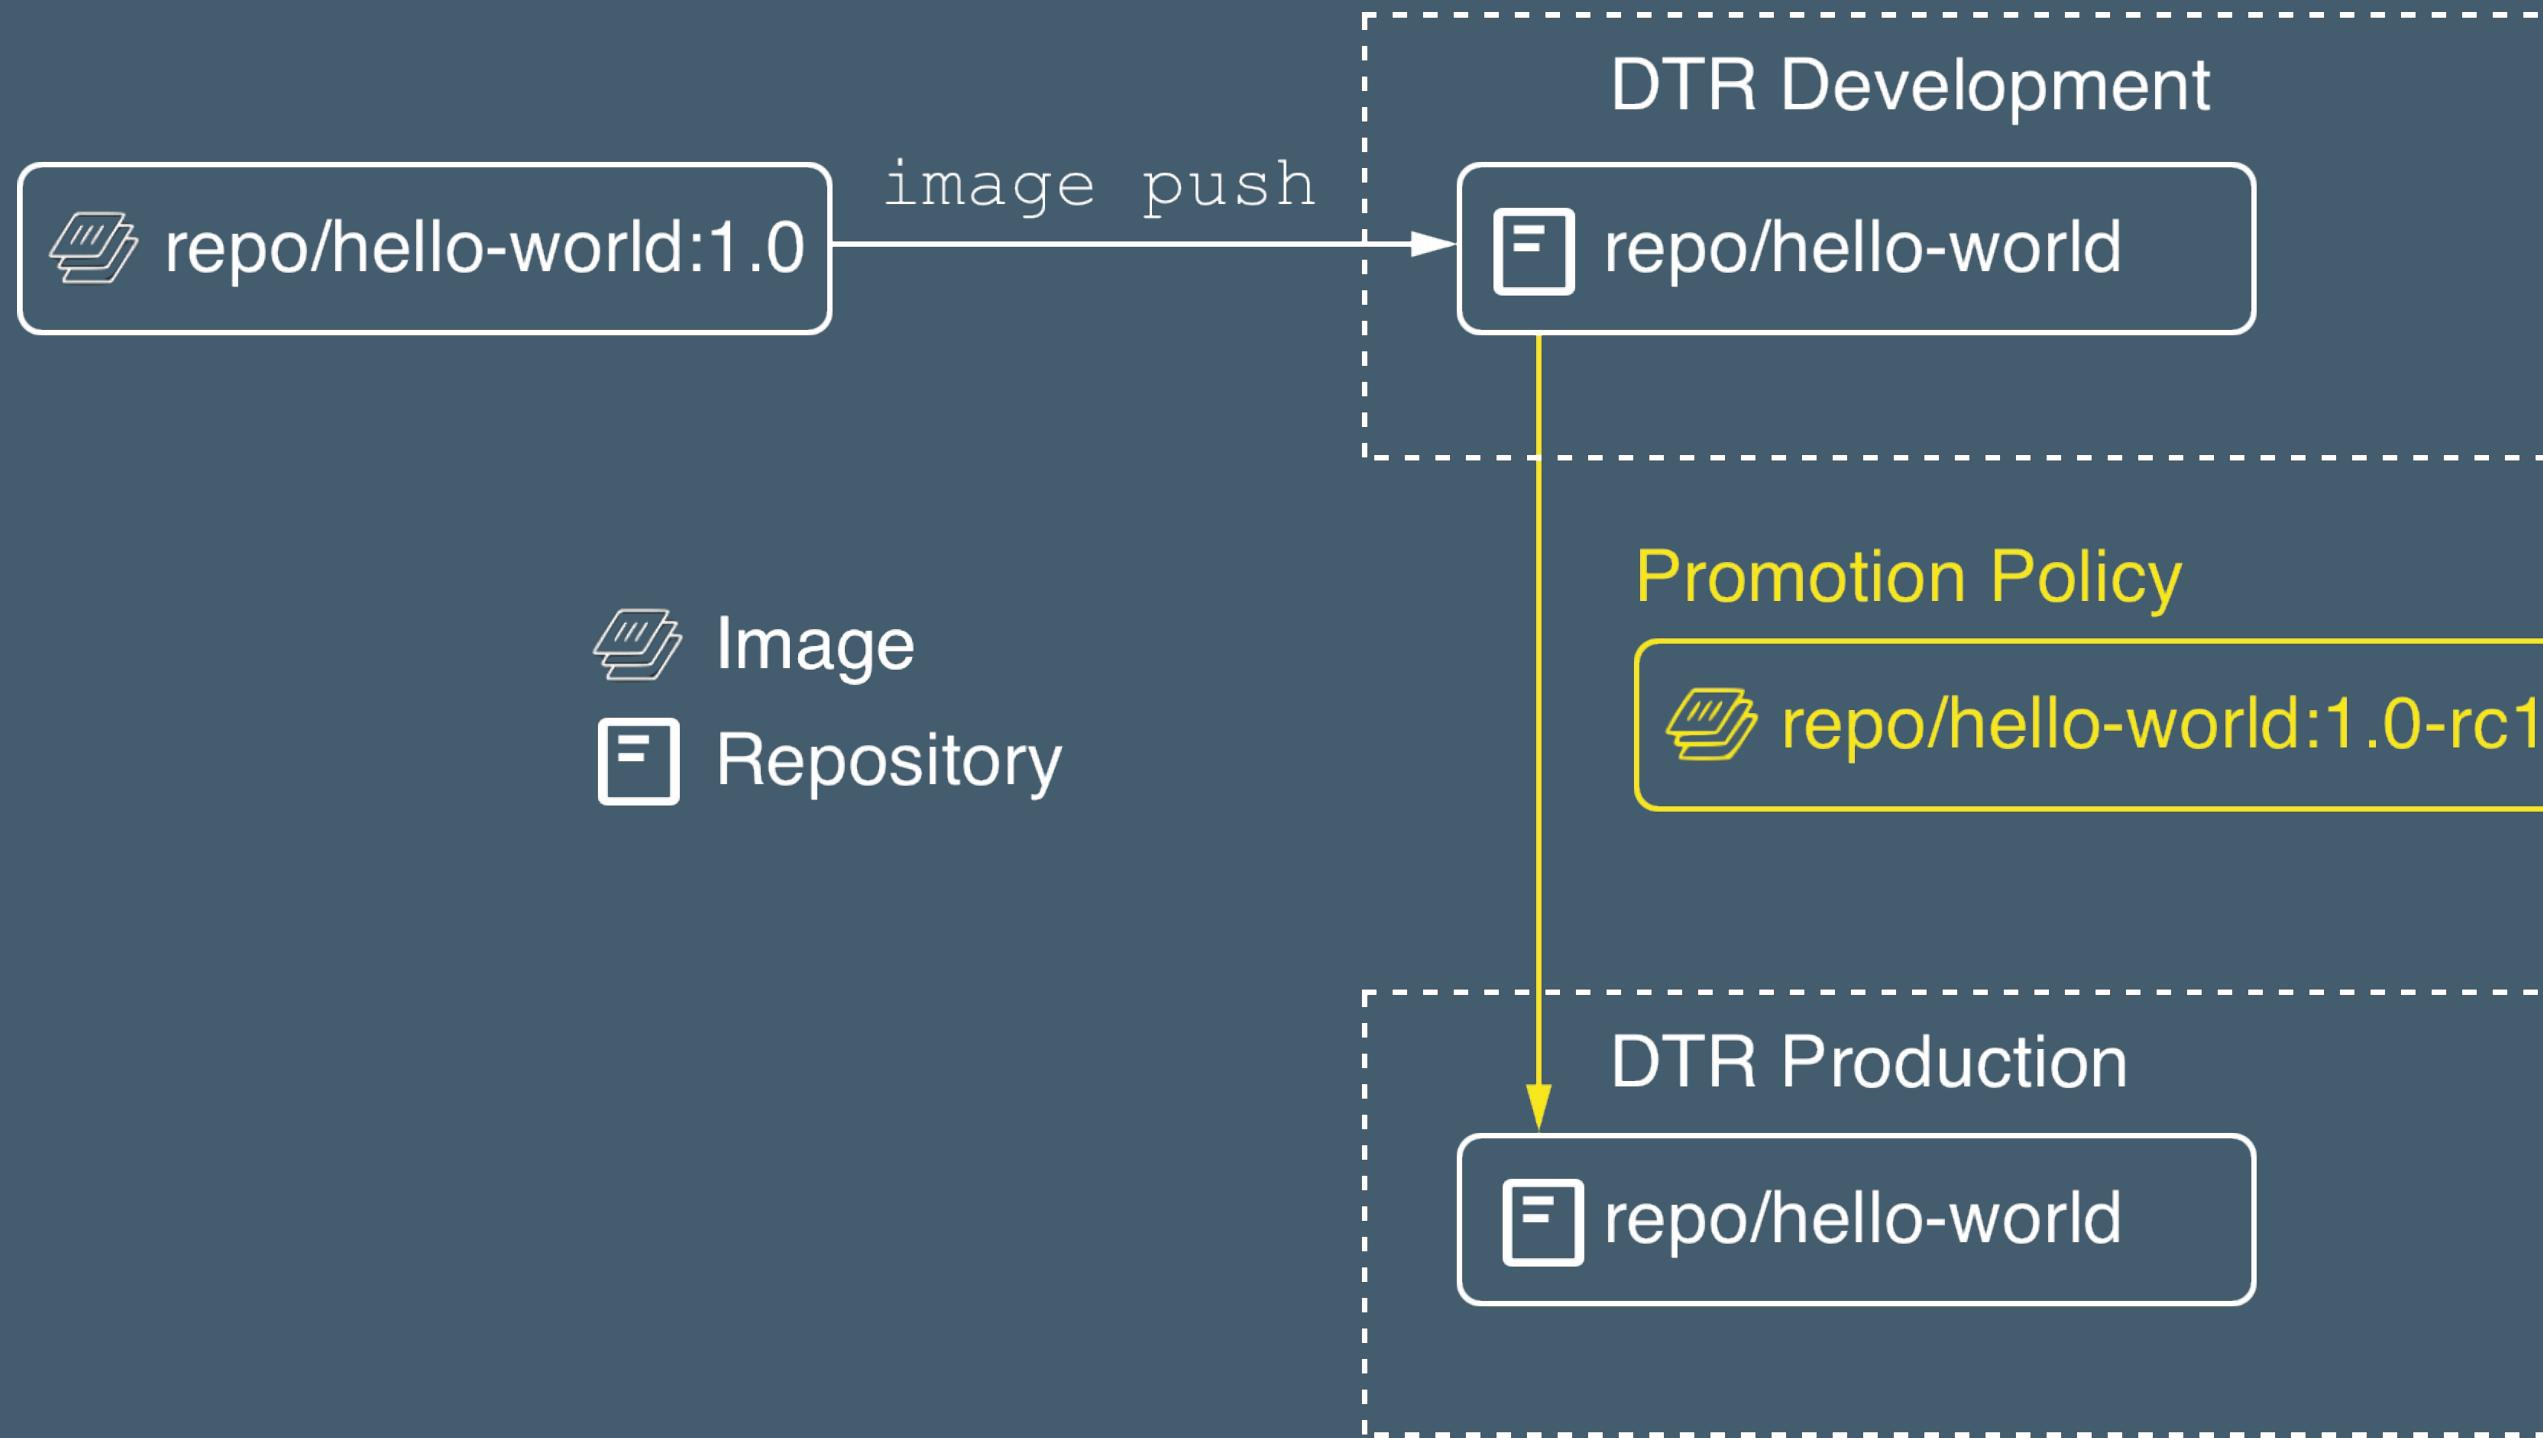


# IMAGE PROMOTION POLICIES

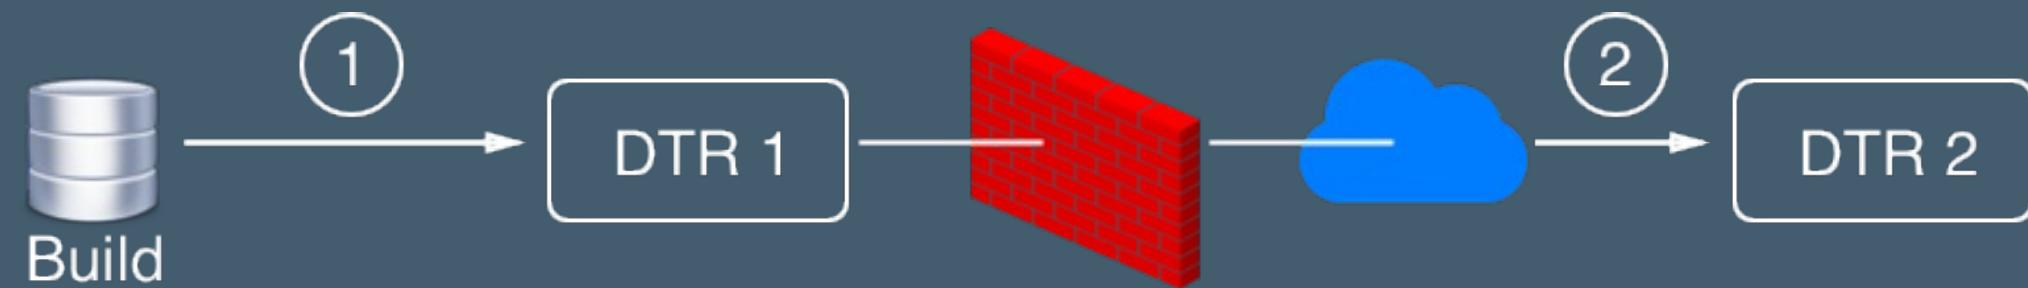
- Manual or automatic
- Automatic promotion can be triggered on:
  - Tag name
  - Package name
  - Minor / Major / Critical / All vulnerabilities
  - Component licenses
- No limit to number of policies that can be defined



# IMAGE MIRRORING



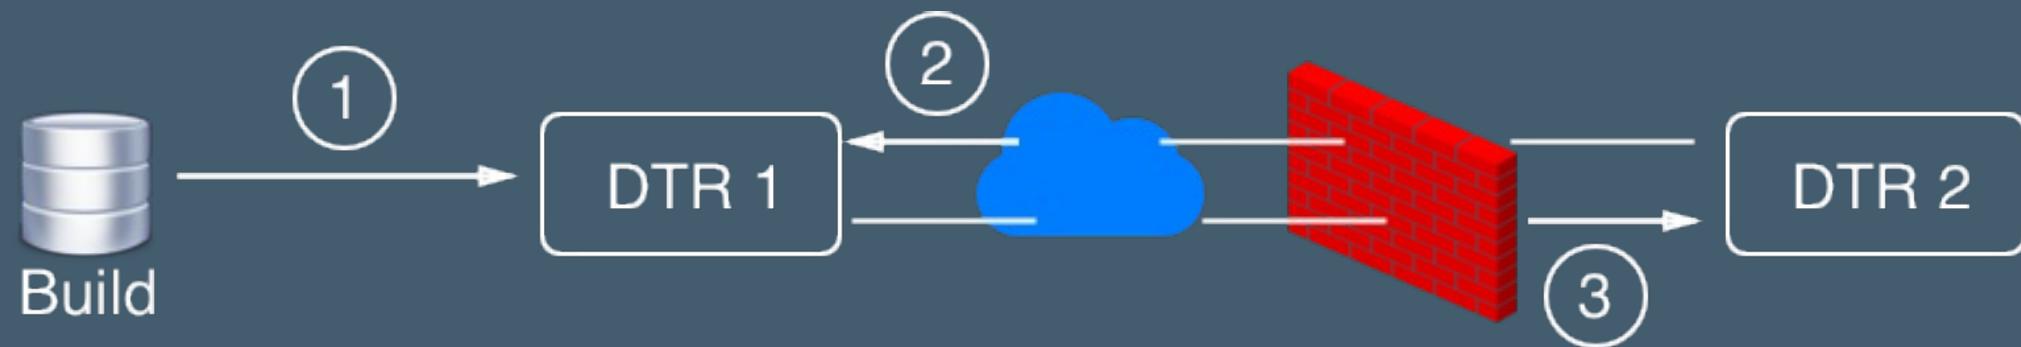
# IMAGE MIRRORING - PUSH BASED



- Image pushed to DTR 1
- If policies met => push to DTR 2
- AuthN & AuthZ managed by each DTR
- Signing & scan data not (yet) preserved



# IMAGE MIRRORING - PULL BASED



- Image pushed to DTR 1
- DTR 2 polls DTR 1 for updates
- New image found => pull to DTR 2
- Combine with Promotion Policies



# WEBHOOKS

POST message with JSON payload, triggered on:

- Tag push or delete
- Manifest push or delete
- Security scan failed
- Security scan complete

Defined per repository.



# WEBHOOK PAYLOAD

- Webhook payloads always come in a wrapper:

```
{
 "type": "...",
 "createdAt": "2012-04-23T18:25:43.511Z",
 "contents": {...}
}
```

- The **contents** key depends on the event type; see <https://dockr.ly/2JjcAW7> for the full spec.





## EXERCISE: REPOSITORY AUTOMATION

Work through the 'Image Promotion & Webhooks' exercise in your exercise book.



# DISCUSSION

- What are the pros and cons of promoting images in a single DTR, versus mirroring them across multiple DTRs?
- Questions?



## FURTHER READING

- Managing webhooks: <https://dockr.ly/2JjcAW7>
- Promotion policies overview: <https://dockr.ly/2KarnDN>
- Image Promotions and Immutable Repos: <http://bit.ly/2eEz7TH>





# TAGGING & VERSIONING STRATEGIES



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Recognize and use some common image tagging patterns.



# TAGGING IMAGES

<URL-to-registry>/<org-or-user>/<repo-name>:<tag>

https://dtr.example.com/engineering/web-ui:1.0.2

Tag is NOT unique!

Each image get SHA-256 associated == unique!



# VERSIONING STRATEGIES

- Semantic Versioning: <MAJOR.MINOR.PATCH>
- Release Date Versioning: <YY.MM.PATCH>
- Other
- Company specific versioning



# TAGGING & VERSIONING STRATEGIES

Version in Tag is not enough...

Mongo DB

- mongo:3.0.15
- mongo:3.0.15-windowsservercore

Python

- python:2.7.13
- python:2.7.13-wheezy
- python:2.7.13-alpine
- python:2.7.13-windowsservercore



# FURTHER READING

- Docker tag reference: <http://dockr.ly/2vP3vhw>
- Docker Reference Architecture: Development Pipeline Best Practices Using Docker EE: <http://dockr.ly/2tETN0V>





# BUILD SERVER



# LEARNING OBJECTIVES

By the end of this module, learners will be able to:

- Combine Jenkins, DTR and GitHub to automatically build and push images upon code updates.



# BUILD SERVER

Typical tasks of a CI server:

- Building images
- Testing images
- Pushing images
- Signing images
- Promoting images
  - to other environments
  - to other repositories
- Provisioning of infrastructure
- Notifying team members upon success or failure
- Auditing activity



# BUILD SERVERS

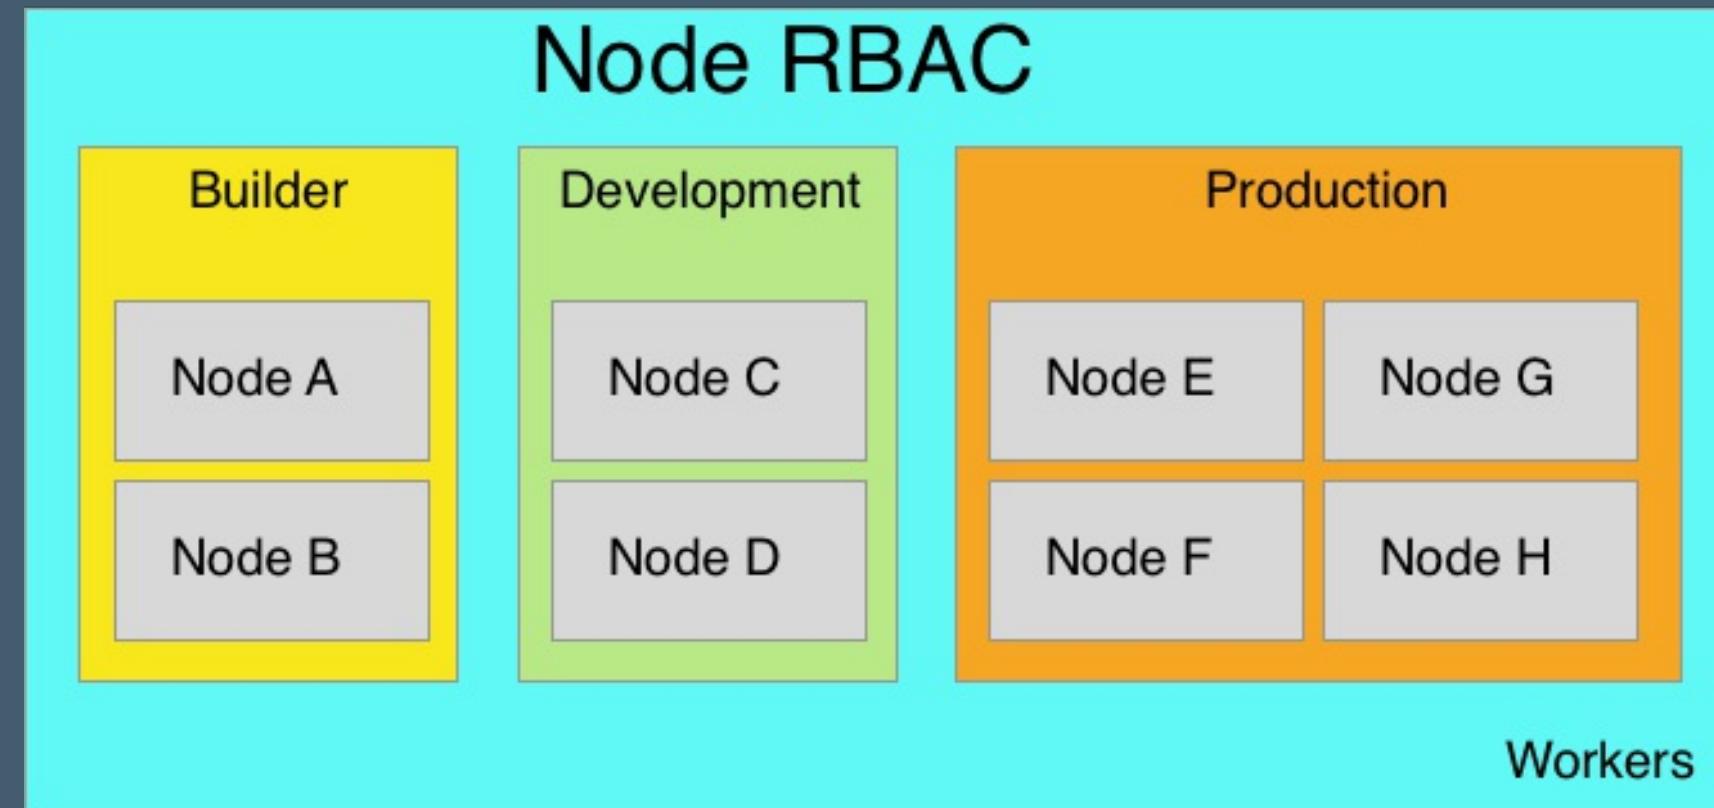
CI servers often used on premise:

- Cloud Bees: Jenkins
- Jetbrains: TeamCity
- Atlassian: Bamboo
- Microsoft: TeamFoundation server



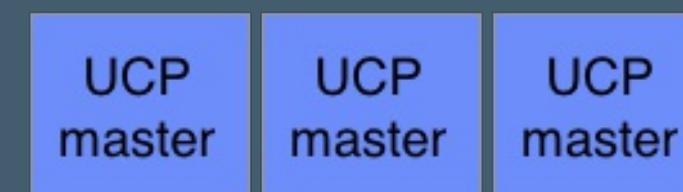
# BUILD SERVER - UCP

- Run build master/agents as containers
- Mount Docker socket to build, sign and push images
- Use node RBAC to define Builder Swarm Segment



## RUNNING TEAMCITY

```
version: '3.1'
services:
 teamcity:
 restart: unless-stopped
 image: sashgorokhov/teamcity
 ports:
 - "0.0.0.0:8111:8111"
 volumes:
 - "/mnt/teamcity:/mnt/teamcity"
 teamcity_agent:
 restart: unless-stopped
 image: sashgorokhov/teamcity-agent
 environment:
 SERVER_URL: teamcity:8111
 volumes:
 - "/var/run/docker.sock:/var/run/docker.sock"
```





## EXERCISE: BUILD SERVER

Work through the 'Build Server' exercise in the Docker for Enterprise Developers Exercises book.



# DISCUSSION

- What are some features of DTR and UCP you'll want to integrate with your CI pipeline?
- Questions?



## FURTHER READING

- Configure automated builds with Docker Hub: <http://dockr.ly/2wjwo4Z>
- Configure automated builds with Bitbucket: <http://dockr.ly/2xeTX30>
- Run only the image you trust: <https://dockr.ly/2HlaAA5>
- DockerCon Video: Delivering Ebay's CI Solution with Apache Mesos and Docker: <http://dockr.ly/2vWoBdR>



# DOCKER FOR ENTERPRISE DEVELOPERS

Thanks for coming! Please take one of our feedback surveys:

- Docker for Enterprise Developers (standalone): <http://bit.ly/2GzPQQp>
- Docker Fundamentals + Enterprise Developers (combined class):  
<https://bit.ly/2uhDfgg>

Get in touch: [training@docker.com](mailto:training@docker.com)

[training.docker.com](https://training.docker.com)



# YOU'RE ON YOUR WAY TO BECOMING DOCKER CERTIFIED!

- Study up with our Study Guides at <http://bit.ly/2yPzAdb>
- Take it online 24 hours a day
- Results delivered immediately
- Benefits include:
  - Digital certificate
  - Online verification
  - Private LinkedIn group
  - Exclusive events

**SUCCESS.DOCKER.COM/CERTIFICATION**

